

Three States, Three Donor Economies

Federal individual contributions in Washington, New York, and Texas (FEC, 2018–2026)

Companion to [the electoral-health prospectus](#). This realizes that paper's cross-state money thread, which was previously data-blocked (NY/TX had zero contributions loaded). Source: `scripts/cross_state_fec_money.py`.

Scope and method (read first)

Basis: federal individual contributions made by IN-STATE RESIDENTS, 2018–2026. NY and TX hold pure FEC bulk data (donor-residence-filtered); Washington's `individual_contributions` mixes state PDC + federal FEC, so WA is restricted to rows carrying an FEC committee id (`fec_candidate_id ~ '^[CPHS]'`) with `contributor_state = 'WA'` . All three are therefore the same thing: **how each state's residents fund federal politics**.

This is an OUTFLOW-by-donor-residence measure — it shows where each state's residents *send* money. It is **not** out-of-state money flowing *into* a state's races (the FEC ingest is filtered on the donor's state, so inflow is not observable here). The "non-constituent money" question from the parent paper still needs a recipient-keyed load.

Donor identity is a `name + zip5` proxy (over-merges common names, so concentration is, if anything, slightly understated). Figures are 2018–2026 federal cycles.

The headline

Metric	WA	NY	TX
Total federal \$ (resident donors)	\$646M	\$2.07B	\$1.94B
Contributions	5.59M	9.98M	12.56M
Distinct donors (name+zip)	361,818	671,156	836,784
Median gift	\$25	\$25	\$25
Gini (donor \$)	0.800	0.848	0.818
Top 1% of donors → share of \$	39.3%	47.5%	41.7%
Top 10% of donors → share of \$	72.3%	78.7%	74.5%
Dollars from gifts < \$200	25.0%	13.8%	20.3%
Dollars from gifts ≥ \$5,000	20.0%	34.8%	33.3%
Dollars from retired donors	24.0%	11.8%	19.5%

One line: **Washington gives the most retail, New York the most top-heavy, and each state's money carries a distinct economic fingerprint — Big Tech, Wall Street, Energy.**

Findings

1. New York is the most top-heavy; Washington the most retail

- **Defensible claim.** The top 1% of donors supply **47.5%** of New York's federal dollars versus **39.3%** in Washington (Gini 0.848 vs 0.800). Conversely, sub-\$200 gifts are a quarter (**25.0%**) of WA's dollars but only **13.8%** of NY's, and ≥\$5,000 gifts are **34.8%** of NY's money vs **20.0%** of WA's. New York's federal money is concentrated at the top; Washington's is comparatively broad-based.
- **Strongest objection.** The Gini of any voluntary-giving distribution is mechanically high everywhere (all three exceed 0.80), so the *level* is not itself pathological — only the *gap* between states is informative. And the small-dollar share is sensitive to how conduit (ActBlue/WinRed) earmarks are recorded (see Limits) — though the gift-size *amount* cut is

computed directly from the dollar value and is unaffected by that.

2. Participation is broadest, per capita, in Washington

- **Defensible claim.** New York and Texas raise $\sim 3\times$ Washington's federal dollars in absolute terms, but Washington has the widest donor *participation* relative to its size: $\sim 362\text{K}$ donors in a state of $\sim 7.9\text{M}$ ($\sim 4.6\%$) versus NY $\sim 671\text{K}/\sim 19.6\text{M}$ ($\sim 3.4\%$) and TX $\sim 837\text{K}/\sim 30.5\text{M}$ ($\sim 2.7\%$). Fewer dollars, more givers.
- **Strongest objection.** The name+zip donor proxy over-merges common names unevenly across states; population denominators are total residents, not voting-eligible adults; and WA's lower dollar total partly just reflects fewer ultra-wealthy households, not broader civic habit.

3. The retired-donor economy is largest in Washington

- **Defensible claim.** **24.0%** of Washington's federal donor dollars come from donors who list their occupation as *retired* — versus **19.5%** in Texas and just **11.8%** in New York. WA's federal money leans most heavily on people no longer in the workforce. (A looser "non-working" bucket that also folds in *not-employed* / *none* / *blank* reaches $\sim 48\%$ in WA, but that figure is soft — see objection.)
- **Strongest objection.** FEC occupation/employer strings are self-reported and noisy. The retired-only figure is the defensible one; the broader "non-working" number bundles wealthy non-earners and blank/missing fields with the genuinely jobless and carries no income signal, so it should be read as a loose upper bound, not a measurement.

4. Sector signatures: Tech (WA), Wall Street (NY), Energy/Industrial (TX)

- **Defensible claim.** The largest corporate employers of each state's federal donors are unmistakably regional:
- **WA — Big Tech:** Microsoft (\$15.2M), Amazon (\$5.0M), University of Washington (\$4.7M), Zumiez (\$4.2M), Fisher Investments (\$3.1M).
- **NY — Wall Street:** Blackstone (\$11.1M), Goldman Sachs (\$7.9M), KPMG (\$6.6M), Jane Street (\$5.8M), Soros Fund Management (\$5.4M).
- **TX — Energy / Industrial:** BNSF Railway (\$8.0M), Valero (\$7.0M), Starkey Hearing (\$5.6M), Beal Bank (\$5.1M), Lockheed Martin (\$4.6M).

The economic base that funds federal politics differs sharply by state. - **Strongest objection.** These employer sums are dominated by a handful of mega-donors at each firm, not broad rank-and-file employee giving; and the strings are free-text (Texas's single largest "employer" is the generic "*Entrepreneur*"), so firm-level totals are indicative, not audited.

5. A uniform presidential rhythm

- **Defensible claim.** All three states show presidential-cycle dollars running $\sim 2\times$ their off-year totals, in lockstep (WA \$202M/2020 vs \$79M/2018; NY \$612M vs \$299M; TX \$544M vs \$266M). Federal giving is paced by the national calendar, not state-specific dynamics.
- **Strongest objection.** This is mechanical (presidential races simply cost more) and says nothing distinctive about any state — it's a useful uniformity check, not a finding.

Follow-on tests

Computed in `scripts/cross_state_fec_tests.py`. Committee master: 44,392 committees; 100% of recipient dollars matched in all three states.

A. Is the money concentrating over time?

Top-1% donor dollar share, by cycle:

Cycle	WA	NY	TX
2018	28.2%	35.3%	29.2%
2020	34.5%	42.4%	35.6%

Cycle	WA	NY	TX
2022	30.3%	36.7%	34.9%
2024	36.2%	47.4%	41.9%
2026*	30.6%	40.3%	37.9%

- **Defensible claim.** Concentration follows a **presidential sawtooth** (peaks in 2020 and 2024) with a **mild secular upward drift**. Comparing like cycles: top-1% share rose presidential-to-presidential in all three (2020→2024: WA +1.7, NY +5.0, TX +6.3 pts) and midterm-to-midterm (2018→2022: TX +5.7, NY +1.4, WA +2.1). New York is both the most concentrated and rising fastest; **Washington is the flattest and most stable**.
- **Strongest objection.** Only three presidential and two-to-three midterm points — too few to call a secular trend confidently. The 2024 spike is entangled with record presidential joint-fundraising activity (see Test B), so "concentrating" is partly a cycle-composition artifact, not purely a structural shift. Direction = mildly rising, not a clean climb. (2026 is partial-cycle.)

B. Where does each state's money go?

Destination of residents' federal dollars (recipient committee → state/office via the FEC committee master):

Destination	WA	NY	TX
In-state Congress	13.4%	10.1%	16.0%
Out-of-state Congress	24.1%	29.5%	19.4%
Presidential	11.0%	7.0%	6.5%
PAC / party / other	51.5%	53.4%	58.1%

- **Defensible claim.** Residents fund **their own congressional delegation least of all** — just 10–16% of their federal dollars. They give roughly **2× as much to out-of-state congressional races** (WA 24%, NY 30%), and the majority (51–58%) flows to **national party committees and joint-fundraising vehicles**. Washington's single largest federal destination is the Democratic presidential JFC (~\$61M) — on par with *all* in-state congressional giving combined (~\$87M); other top destinations are the DNC, DCCC, DSCC, RNC, and the Trump JFCs. Texans are the most "local" (in/out-of-state Congress ratio 0.82); New Yorkers the most nationalized (0.34 — over 3× more to out-of-state than to their own delegation). This is the **donor-side counterpart to the nationalization-of-money literature**: even constituents' own money is overwhelmingly aimed at national / out-of- state politics rather than their representatives.
- **Strongest objection / caveat.** This remains **outflow** (where residents *send* money), not money flowing *into* each state's races. The majority "PAC/party/other" bucket is genuine national party + JFC money, but **JFCs blur the boundaries** — a presidential JFC funds the nominee *and* the party — so the "Presidential" row understates true presidential+national giving, and the splits among national vehicles are soft. The in- vs. out-of-state *Congress* split is the robust part (direct candidate-committee gifts).

C. Top donors, top recipients, and the cross-state magnets

Computed in `scripts/diag_cross_state_donors.py`.

Largest individual donors confirm the sector fingerprint at the person level: - **WA (tech/VC)**: the Cornfields (~\$3.3M), Tom Champion/Zumiez (\$3.2M), Nick Hanauer (VC), Rory & Melinda Gates. - **NY (finance/philanthropy)**: Philip Munger (\$4.1M), George Soros (\$4.0M), Stephen Schwarzman/Blackstone (\$3.2M), Agnes Gund/MoMA. - **TX (energy/industrial)**: Syed Anwar/PetroPlex (\$5.2M), Paul Foster/Western Refining (\$3.5M), the Perots, Woody Hunt.

Largest recipient committees are national vehicles, with marquee in-state contests poking through: D-state dollars concentrate in the **Harris Victory Fund** (WA \$61M, NY \$176M) and **Fight for the People PAC** (NY \$82M); TX dollars in **Trump Victory** (\$71M) and the **RNC** (\$67M). In-state contests that surface: Kim Schrier (WA-08, \$17M), Patty Murray, Ted Cruz (TX, \$35M).

Cross-state money magnets. Of 12,256 committees these donors touch, **4,602 are funded by donors in all three states**. The top are national party / JFC vehicles (Harris Victory Fund \$290M combined; RNC \$111M; DSCC / DNC / DCCC ~\$96–99M each). The single cleanest nationalization signal: **Warnock for Georgia** — a *Georgia* Senate race drawing **WA \$7.4M + NY \$16.3M + TX \$6.0M**, funded by three states whose residents cannot vote in it.

- **Caveat.** Individual-donor identity is a name+zip proxy (merges across cycles; a few employer labels are data-entry quirks). Recipient totals are donor-residence *outflow from these three states only* — not the committee's full national haul.

D. Does money chase competitive races? (money × competitiveness)

Computed in `scripts/diag_money_vs_competitiveness.py`, joining residents' U.S. House contributions to each district's competitiveness band (this project's own forecast margin: Tossup <5 / Lean 5–10 / Likely 10–20 / Solid ≥20). PARTIAL run: WA + TX donors → WA + TX House districts, 2022–2026 (post-redistricting). NY pending (DB was locked by the live inflow load); the full 3-state + inflow-side version will follow.

Band	# districts	% of districts	\$ to band	\$ / district	% of \$
Tossup (<5)	3	6.2%	\$11.3M	\$3.77M	7.9%
Lean (5–10)	2	4.2%	\$5.4M	\$2.71M	3.8%
Likely (10–20)	24	50.0%	\$76.7M	\$3.20M	53.6%
Solid (≥20)	19	39.6%	\$49.7M	\$2.62M	34.7%

- **Defensible claim.** Donor money chases competitiveness only **weakly**. Tossup districts draw the most per district (\$3.77M vs \$2.62M in Solid — a ~1.4× premium), but **88% of all dollars still flow to safe (Likely + Solid) districts**, because ~90% of districts are safe. In-state House money is dominated by support for (mostly safe-seat) candidates, not strategic targeting of the marginal race — the donor-side echo of "money follows the scoreboard." Cross-state House giving between WA and TX is negligible (\$2.4M of \$143M): residents fund their *own* state's House candidates.
- **Caveats.** WA+TX only (NY — the largest Democratic-money state — still to add). Donor-side *outflow*, not inflow. The "Lean" band has only 2 districts (noisy). Strategic/competitive targeting is more likely to surface in PAC/JFC and out-of-state money (the inflow load + the large national-vehicle bucket from Test B) than in this in-state-resident slice — which is exactly what the inflow load will let us test.

E. Inflow side — does money chase competitive races? (WA+NY+TX House & Senate)

From the recipient-anchored inflow dataset (`fec_inflow.duckdb`: 5.48M contributions / \$1.20B, all-state donors → WA/NY/TX federal candidates, built by `scripts/load_fec_inflow_bulk.py` in minutes — the API path would have taken days), joined to competitiveness. `scripts/diag_inflow_vs_competitiveness.py`.

U.S. House, 2022–2026 — \$485M across 74 districts:

Band	# dists	% dists	\$ in	\$ / district	% of \$	out-of-state share
Tossup (<5)	4	5.4%	\$47.2M	\$11.81M	9.7%	45.0%
Lean (5–10)	4	5.4%	\$42.3M	\$10.59M	8.7%	36.1%
Likely (10–20)	31	41.9%	\$190.5M	\$6.14M	39.3%	39.9%
Solid (≥20)	35	47.3%	\$205.2M	\$5.86M	42.3%	44.4%

- **Defensible claims:** 1. **The competitiveness premium is real and ~2×.** Tossup (\$11.8M/district) and Lean (\$10.6M) districts pull roughly double the per-district inflow of safe seats (~\$6M) — money *does* chase the marginal race, far more clearly than the donor-side ~1.4× (Section D). 2. **But safe seats still capture ~82% of the money** (Likely+Solid), because they're ~89% of districts. Likely and Solid draw about the *same* per district (~\$6M): once a seat is safe, *how* safe barely changes the money — the jump is **between** competitive and safe, not within safe. 3. **~40–45% of all inflow is out-of-state, roughly uniform across every band.** Nationalization is **pervasive, not battleground-specific** — roughly two-fifths of the money funding these House races comes from people who cannot vote in them, in safe and tossup seats alike.

U.S. Senate, 2018–2026 (the model does not forecast US Senate; competitiveness via actual results):

State	\$ in	out-of-state share	Senate races in window
TX	\$253.2M	45.3%	competitive — Cruz/O'Rourke 2018 (R+2.6), Cruz/Allred 2024 (R+8.8)

State	\$ in	out-of-state share	Senate races in window
NY	\$55.2M	53.5%	safe-D — Schumer / Gillibrand
WA	\$45.0M	41.1%	safe-D — Murray / Cantwell

- **Senate echoes the House, louder.** Competitive TX Senate races drew **\$253M** — **~5× safe NY (\$55M) or WA (\$45M)**: at the statewide level, competition is the single biggest money magnet. Yet out-of-state share is high *everywhere* (41–54%) and is actually **highest in safe NY (53.5%)** — national donors fund high-profile safe senators (Schumer/Gillibrand) as readily as battlegrounds. Same lesson as the House: competition lifts the total, but the out-of-state flood is profile-driven and pervasive.
- **Earmarks ARE attributed (verified — correcting an earlier caveat).** Conduit-routed (ActBlue/WinRed) money is **not** lost from these totals: FEC records each earmarked individual gift under the *candidate* committee as transaction type **15E** — **\$194M for these candidates in 2024 alone, more than the \$90M of direct 15 gifts** — and the inflow load captures it. The conduit-side 24T records (\$150M) are the *same money* seen from the conduit and are correctly excluded to avoid double-counting (`scripts/diag_earmark_inspect.py`).
- **Caveats.** House competitiveness = 2026 forecast bands on current districts; Senate banded by actual two-party result. WA contributes no congressional Tossups in the 2026 map (its competitive seats land in Lean/Likely). State-legislative money still excluded (federal only).

F. The individual layer — who votes, who gives, and whether the donor skew is real (WA, 382K match)

Computed in `scripts/diag_wa_individual_findings.py` against `data/wa_statewide.duckdb` (`voter_scores` , `voter_donor_affiliation`) with `data/wa_vrdb.duckdb` ATTACHED (`voters` , 5.51M; `voting_history` , 27.1M). WA is the one state here that needs no external voter file — it already has the registered roll, the vote history, and the person-level voter→donor match. These refresh the democracy-insight gauntlet's figures on the improved match (320K → 382,408 voters, 343,891 of them carrying a generation label).

F1 — The off-year electorate is old, and the young drop out, not down. - **Defensible claim.** Within-cohort turnout and electorate composition both swing hard by cycle type. Turnout of 18–29s **collapses 58.4% → 15.8%** (2024 presidential → off-year generals), while 65+ falls only **88.3% → 61.3%**. As composition: 65+ are **28.5%** of the 2024 presidential electorate but **~39%** of the off-year electorate, while 18–29 fall from **14.2% → 7.6%** — a senior-to-youth ratio that widens from ~2:1 to ~5:1 off-cycle. The decision in odd-year local elections is made by a substantially older electorate than the one that shows up for president. - **Strongest objection.** This is *voluntary* differential participation under WA's all-mail, postage-paid, automatic-registration system — not disenfranchisement — and the fix (move local races on-cycle) is institutional, which is itself evidence of a *responsive* system, not a failing one. The within-cohort *rate* also uses a current-roll denominator (survivorship-biased), so read the composition **shares** as the robust cut. - **Caveat.** VRDB vote history begins in 2021, so the presidential figure rests on 2024 alone and the midterm on 2022 alone; only the off-year row averages three cycles (2021/2023/2025).

F2 — The donor age skew is genuine, not a matcher artifact (the white paper's open question, answered). - **Defensible claim.** Matched donors over-represent the old and under-represent the young — raw donor-share ÷ roll-share is **Silent 1.87×**, **Boomer 1.64×**, **Gen X 1.18×**, **Millennial 0.59×**, **Gen Z 0.17×**. The white paper's strongest objection was that this is a *matcher* artifact (the (last, first-initial, zip5) uniqueness guard over-selects older, rarer-named, stable-address voters). Tested directly, **it is not**: the probability a voter is uniquely matchable on that key is nearly flat across generations — **68.9% (Gen Z) to 73.1% (Silent)**, a **~4-point spread** — so inverse-propensity re-weighting barely moves the ratios (Silent 1.87→**1.83×**, Gen Z 0.17→**0.17×**). The age skew survives the correction; it is a real property of who gives, not of who the matcher can find. (Person-level concentration on this match: top-1% of matched donors = **47.7%** of matched dollars, top-10% = **80.0%**, Gini **0.862**; geographically, **61.2%** of matched-donor dollars come from just two Seattle-metro ZIP3s — 981xx 35.9% + 980xx 25.2%.) - **Strongest objection / residual bias.** The IPW corrects the *name-commonness* channel — the matcher's actual uniqueness guard — but cannot correct the one channel it can't observe: a donor whose address changed between giving and today fails the zip match, and mobility skews young. That residual bias runs in the *same* direction as the raw skew (it deflates the young donor share), so raw should be read as an **upper bound** on the true age skew — but the mechanism the objection actually named (rare names → easier match for the old) explains essentially none of it.

F3 — Money and votes stack on the same people (association only). - **Defensible claim.** Among the 343,891 matched donors carrying a turnout score, **84.0% are super-voters versus 50.1%** of non-donors (1.68×), and their mean turnout propensity is **0.953 vs 0.748**. Financial voice and electoral voice concentrate on the same individuals rather than offsetting. - **Strongest objection.** This is cross-sectional association, not a causal "giving makes you vote" claim: donors are pre-selected for engagement (reverse causation is equally plausible), and the matcher itself favors stable-address super-voters, so

selection *and* reverse causation both inflate the gap. The benign reading — donating as a participation *gateway* — is fully live.

F4 — Robustness: concentration precision + match validation. - Bootstrap CIs on concentration

(`scripts/diag_donor_concentration_bootstrap.py` , B=1000). Because donor identity is a name+zip5 *proxy*, the concentration estimates carry sampling-style uncertainty — but it is small. Matched-donor **Gini 0.862 [95% CI 0.856–0.868]**, **top-1% 47.7% [45.6–50.0]**, **top-10% 80.0% [79.1–80.9]**; the inflow side (2024) is tighter still (Gini 0.690 [0.688–0.692]). The concentration is a precise feature of the data, not an artifact of which donors landed in the pool. (The proxy's over-merging biases concentration *down*, so true figures are, if anything, slightly higher.) - **Match-precision validation** (`scripts/diag_match_validation_sample.py`). On a seed-fixed 150-match sample, the matched donor's full first name equals the voter's in **87%** of cases, and only **13%** sit on a (last, first-initial, zip5) key shared by a same-key donor namesake — the population where a false merge or pooled-attribution is even possible (4 of 150 could not be re-found, parsing/PDC edge cases). The sample is emitted as a CSV for human adjudication (true precision = Y/(Y+N)); the structural ceiling on false merges is reassuringly low.

G. The cross-state money-flow matrix — the "Warnock" picture, made systematic

Computed in `scripts/diag_cross_state_money_matrix.py`. INFLOW from `data/fec_inflow.duckdb` (recipient-anchored, robust). OUTFLOW and the magnet list from the per-state `individual_contributions` joined to a `committee→candidate` master rebuilt from cached `cm*.zip + cn*.zip`. Recipient state = the candidate's office state (`cn.txt CAND_OFFICE_ST`), resolved `committee→candidate` — not the committee's registration state (`CMTE_ST`), which is often a DC/VA compliance-vendor address (Smiley-WA and Cornyn-TX both register in VA). Outflow and magnets are restricted to authorized candidate committees (designation P/A), so JFCs and leadership PACs — which carry one connected candidate but raise nationally — don't misattribute a recipient state. U.S. House+Senate; all cycles 2018–2026.

INFLOW — who funds each state's H+S candidates (by donor origin):

Recipient state	Total \$	In-state	In-region (other 2)	Rest-of-US
WA	\$154.6M	66.7%	6.3%	27.0%
NY	\$462.7M	55.1%	3.8%	41.1%
TX	\$582.4M	59.6%	5.9%	34.4%

OUTFLOW — where each state's residents send their H+S candidate money (by recipient region):

Donor state	Total \$	In-state	In-region (other 2)	Rest-of-US
WA	\$198.8M	43.8%	5.3%	50.9%
NY	\$672.9M	33.9%	4.0%	62.1%
TX	\$526.6M	54.1%	2.5%	43.4%

- **Defensible claims.** 1. Every one of these states sends the *majority (or near it) of its candidate money out of region. NY residents ship 62% of their House+Senate candidate dollars to candidates in other states; even Texas, the most parochial, sends 43%. This is the donor-side engine of nationalization, measured directly rather than inferred.* 2. *In-region cross-funding is negligible — WA/NY/TX barely fund each other (≤6% of either matrix in every cell).* The out-of-region money is not flowing to neighbors; it is aimed at national battlegrounds. NY is the dominant exporter (~\$418M leaves), WA the dominant net importer relative to its own giving. 3. **The magnet list is a clean battleground map.** The out-of-region candidate committees funded by donors in **all three** states are almost entirely Senate battlegrounds: **Warnock-GA (\$29.7M from these three states)**, **Ossoff-GA (\$19.2M)**, **Kelly-AZ (\$19.5M)**, **Tester-MT (\$12.8M)**, **Harrison-SC (\$11.6M)**, with Graham-SC, Perdue/Loeffler-GA, Gideon-ME, Rosen/Cortez Masto-NV and McConnell-KY close behind. **Georgia Senate races alone (Warnock + Ossoff + Perdue + Loeffler) drew ~\$67M** from residents of WA, NY, and TX — none of whom can vote in Georgia. The Section-C Warnock anecdote is the rule, not the exception.
- **Strongest objection / caveat.** The two matrices use different committee universes (inflow is the recipient-anchored bulk; outflow is per-state authorized committees), so their absolute totals are **not** directly subtractable — read the *shares within each*, not the cross-matrix difference. "Rest-of-US" inflow is robust (blank donor-state is 0.08% of inflow dollars). Out-of-region ≠ "non-constituent and therefore illegitimate": a donor may have ties, and party committees (excluded here) carry yet more of the nationalized money.

H. Sector × competitiveness — does Wall Street chase Tossups while tech/energy fund safe seats?

Computed in `scripts/diag_sector_vs_competitiveness.py`, crossing the employer/sector signatures (Sections A/C) against the inflow competitiveness bands (Section E). Each inflow contribution to a WA/NY/TX U.S. House district (2022–2026, \$485M) is classified by a keyword-on-employer sector and joined to the district's forecast band (Tossup<5 / Lean5–10 / Likely10–20 / Solid≥20). "Competitive share" = (Tossup+Lean) ÷ sector total.

Sector	\$ (House, 3 states)	Competitive share	Out-of-state share
Real estate	\$8.0M	22.0%	32.2%
Finance / Wall St	\$18.4M	21.7%	44.4%
Law	\$17.6M	20.1%	33.1%
Healthcare	\$8.4M	16.1%	35.7%
Energy	\$3.5M	14.3%	26.0%
Tech	\$7.1M	11.7%	51.8%
[all sectors baseline]	\$485M	18.5%	~43%

Sector keyword map extended after a coverage audit (`scripts/diag_sector_coverage.py`) surfaced major law firms, hedge funds, and tech names as unclassified; the additions lifted classified dollars from ~14% to ~15% and left the pattern below intact.

- **Defensible claim.** The hypothesis is **partially borne out, with a twist**. Finance money does tilt toward competition (21.7% of its dollars to Tossup/Lean vs an 18.5% baseline), and **tech is now the single least competition-seeking sector (11.7%)**, with energy next at 14.3% — both disproportionately fund safe seats. But the cleanest contrast is in *travel*, not band: **energy money is strikingly local** (only 26% out-of-state vs ~43% baseline) — it funds its own state's (Texas) incumbents — while **tech money travels the farthest (52% out-of-state) yet lands in safe seats** (national safe-D House members). Finance both travels (44%) and tilts competitive. So: Wall Street chases the marginal race somewhat; tech funds safe co-partisans at a distance; energy funds the home-state incumbent.
- **Strongest objection.** The effect is **modest and partly mechanical**. The competitive-share spread is only ~10 points (real-estate 22% to tech 12%), and tech (WA) and energy (TX) are concentrated in their home states, whose *own* House seats happen to be safe (Solid-D Seattle, Solid-R Texas) — so "tech/energy fund safe seats" is in part just "their in-state seats are safe," not a strategic preference for safety.
- **Caveat.** Even after extending the keyword map with the biggest missing law firms, hedge funds, and tech names, classified sectors are still a **thin slice — only ~15% of inflow dollars**; "retired / not-employed / blank" is 46.8% and "unclassified" ~39%, and tech (\$7.1M) and energy (\$3.5M) volumes in these three states' House races are small enough to be noisy. Employer strings are self-reported free text; the keyword map is indicative, not audited (`scripts/diag_sector_coverage.py`). House-only (Senate competitiveness isn't model-forecast).

I. Inflow concentration trend + donor retention — is the candidate-money base democratizing?

Computed in `scripts/diag_inflow_concentration_retention.py` on the inflow side (all-state donors → WA/NY/TX H+S candidates; donor = name+zip5 proxy). Section A measured concentration on the outflow; this measures it on the inflow, and adds repeat-vs-one-time retention.

Cycle	Inflow \$	Donors	Top-1% \$	Top-10% \$	Gini
2018	\$252M	244K	16.2%	52.4%	0.651
2020	\$249M	225K	17.2%	54.2%	0.672
2022	\$248M	218K	15.8%	55.9%	0.689
2024	\$284M	310K	18.4%	58.5%	0.690
2026*	\$167M	166K	16.5%	58.3%	0.686

- **Defensible claim 1 — candidate-directed money is far less concentrated than the total flow.** Inflow to candidates (gifts capped at the per-election limit) runs **top-1% ≈ 16–18%, Gini ≈ 0.69** — against the **outflow** figures of top-1%

39–48% and Gini **0.80–0.85** (Sections A & F, which include uncapped JFC/PAC/party money). The system's dollar concentration lives in the **party/JFC layer, not in direct candidate giving**, which is structurally egalitarian by the contribution cap. There is a **mild secular rise even within candidate money** (Gini 0.651→0.690, top-10% 52→59% across 2018→2024), but no presidential sawtooth — concentration here drifts up gently and plateaus.

- **Defensible claim 2 — a churning, mostly one-time base funds the candidates, but a persistent minority supplies most of the dollars.** Across 2018–2026, **78.2% of donors give in only one cycle** (supplying **41.7%** of dollars); the **21.8%** who give in ≥ 2 cycles supply **58.3%**. And cycle-over-cycle retention is **low and flat** — **~21–25%** of a cycle's donors gave in the immediately prior cycle (2020 23.4%, 2022 25.0%, 2024 21.1%), with **no rising trend**. Roughly three-quarters of each cycle's candidate donors are new-or-retained-from-the-past, not a stable subscription base. The "small-dollar democratization" picture — broad, churning participation — holds by *headcount*; the "thin layer over a concentrated core" picture holds by *dollars*.
- **Strongest objection / caveat.** The naive "returning %" *looks* like it rises (23→39%), but that is a **look-back-window artifact** — later cycles have more prior cycles to match against; the fixed one-cycle look-back (above) removes it and shows flatness. The name+zip5 donor proxy **over-merges** common names, which inflates "repeat" and deflates one-time counts, so true churn is, if anything, *higher* than reported. **2026 is a partial cycle** (its elevated retention reflects early givers skewing toward committed repeat donors before the late small-dollar surge). Capped candidate giving $\neq$ the whole money system — the uncapped layer (Section A) is where concentration concentrates.

J. Which side of a safe seat gets the money? (longshot vs favored)

New cut in `scripts/diag_loser_side_money.py`; NY + WA U.S. House inflow 2022–2026, recipient party from the committee→party map, favored party from the forecast margin. TX omitted — its committee→party map isn't built.

A safe seat for one party is a longshot for the other, so we can ask which side the money actually reaches. Almost all of it goes to the favored side, and the safer the seat, the more lopsided the split. In New York the longshot party's share of House inflow falls from 74% in the lone tossup to 27% in Likely seats to just **5.5%** in Solid (≥ 20 -point) seats; Washington runs the same way (24.6% Likely → **4.8%** Solid). Put plainly: in a truly safe district the disadvantaged party's candidate raises about a nickel on the dollar and the favored side takes the rest. (Caveats: this is money entering the race, House only; leadership PACs tied to safe incumbents can pad the favored side; and "favored" is only meaningful once a seat is actually safe — in a real tossup the challenger can out-raise the nominal favorite, which is why the tossup row looks inverted.)

Limits of inference

Full provenance + a no-AI reproduction recipe (every source, access path, and the exact scripts behind each figure): [Data Sources & Reproducibility](#).

- **Both directions now exist (Findings 1–5 + Tests A–D are outflow; Tests E–I add inflow).** The early findings measure where residents *send* money (outflow by donor residence); the recipient-anchored inflow load (Sections E–I) measures money *entering* WA/NY/TX races. The cross-state matrix (G) reports both. They use different committee universes, so read shares within each rather than subtracting one from the other.
- **Federal only.** State-legislative and local money (NYSBOE / Texas Ethics Commission) is not included — those need the adapter work flagged in the white paper.
- **Conduit-reliance metric is unreadable (but the money is captured).** ActBlue/WinRed appear as recipients of $<0.5\%$ of dollars — because, as the 2024 earmark inspection confirmed (Section E), FEC records each earmarked gift under the *candidate* committee as type `15E`, not under the conduit. So you cannot measure conduit *usage* from the recipient field — but the **money itself is fully captured** (under the candidates). **Do not read the ActBlue/WinRed-as-recipient share as "conduit reliance."** The sub-\$200 *amount* share (Finding 1) is computed from the dollar value and is unaffected.
- **Donor proxy.** name+zip5 over-merges common names, modestly understating donor counts and overstating concentration.
- **WA composition.** WA's FEC subset draws from both a donor-filtered bulk load and an earlier per-candidate load; both are filtered here to WA residents, but completeness may differ slightly from NY/TX pure-bulk.
- **No income/race.** Occupation/employer is the only socioeconomic signal, and it is self-reported.

What's done, and what's next

Tests A–I are run. Status: 1. **True inflow — DONE for WA+NY+TX, House + Senate** (Section E): recipient-anchored *bulk* load (`scripts/load_fec_inflow_bulk.py` → `fec_inflow.duckdb` , **5.48M contributions / \$1.20B** of all-state money into WA/NY/TX federal candidates) — built in minutes vs. the API path's days. 2. **Conduit/earmark attribution — DONE/verified** (Section E): earmarked ActBlue/WinRed money is attributed to candidates via FEC 15E and already counted (\$194M in 2024); conduit-side 24T duplicates correctly excluded. No fix needed. 3. **WA individual voter-file study — DONE** (Section F): turnout-by-age, donor representativeness *with* the matcher-bias re-weighting (the skew survives → genuine, not an artifact), and giving↔turnout, all on the 382K match. WA needs no external file. The **ID/NY individual** versions are still gated on acquiring those voter files. 4. **Cross-state flow matrix — DONE** (Section G): inflow provenance + outflow destination + the systematic out-of-region magnet list (Georgia Senate ~\$67M from WA/NY/TX residents). 5. **Sector × competitiveness — DONE** (Section H): finance tilts mildly competitive; tech/energy fund safe seats; energy money is distinctively local. Modest, on a thin classified slice. 6. **Inflow concentration + retention — DONE** (Section I): candidate money is far less concentrated than total flow (Gini ~0.69 vs ~0.80); base is churning/one-time by headcount, concentrated-core by dollars; cycle-over-cycle retention flat ~21–25%. 7. **State-level money** — NYSBOE (~½ day) + TEC (new adapter) to extend below the federal layer.