

Four States, Four Donor Economies

Federal individual contributions in Washington, New York, Texas, and Idaho (FEC, 2018–2026)

Stephen Kirby · Tikor Consulting · July 2026

AI-assisted drafting and analysis review. All figures are reproducible from public-record data and from the open-source scripts cited below, including `scripts/verify_cross_state_money.py`. The paper source, code, and data-acquisition recipe are public at <https://github.com/skirby359/who-decides>; the underlying voter files are not redistributed. Contact: kirby@tikorconsulting.com.

Companion to [the electoral-health prospectus](#). This realizes that paper's cross-state money thread, which was previously data-blocked (NY/TX had zero contributions loaded). Source: `scripts/cross_state_fec_money.py`. Idaho added 2026-07-19 (FEC outflow + inflow loaded to parity) as the small, deep-red pole; all sections (headline, Findings 1–5, Tests A–I, and the flow matrix G) are now four-state — every table carries WA/NY/TX/ID.

Scope and method (read first)

Basis: federal individual contributions made by IN-STATE RESIDENTS, 2018–2026. NY and TX hold pure FEC bulk data (donor-residence-filtered); Washington and Idaho mix state finance (WA=PDC, ID=Sunshine) + federal FEC in `individual_contributions`, so both are restricted to rows carrying an FEC committee id (`fec_candidate_id ~ '^([CPHS])'` , which excludes the state-finance rows) with `contributor_state = <st>` . All four are therefore the same thing: **how each state's residents fund federal politics.**

Findings 1–5 and Tests A–D are an OUTFLOW-by-donor-residence measure — they show where each state's residents *send* money (the per-state FEC ingest is donor-state-filtered). The complementary **INFLOW** side — out-of-state money flowing *into* each state's races, the "non-constituent money" question from the parent paper — is answered in **Tests E–I and the flow matrix G**, built from the recipient-anchored `fec_inflow.duckdb` . Both directions now exist; they use different committee universes, so read shares *within* each rather than subtracting one from the other.

Donor identity is a `name + zip5` proxy (over-merges common names, so concentration is, if anything, slightly understated). Figures are 2018–2026 federal cycles.

The headline

Metric	WA	NY	TX	ID
Total federal \$ (resident donors)	\$646M	\$2.07B	\$1.94B	\$76M
Contributions	5.59M	9.98M	12.56M	0.77M
Distinct donors (name+zip)	361,818	671,488	836,784	54,155
Median gift	\$25	\$25	\$25	\$25
Gini (donor \$)	0.800	0.848	0.818	0.775
Top 1% of donors → share of \$	39.3%	47.5%	41.7%	36.0%
Top 10% of donors → share of \$	72.3%	78.7%	74.5%	69.2%
Dollars from gifts < \$200	25.0%	13.8%	20.3%	29.0%
Dollars from gifts ≥ \$5,000	20.0%	34.8%	33.3%	20.1%
Dollars from retired donors	24.0%	11.8%	19.5%	31.7%

(Bold = most concentrated / top-heavy on that row; italic = least. New York holds every top-heavy extreme; Idaho now holds every retail extreme.)

One line: **New York gives the most top-heavy, Idaho the most retail; Washington and Texas sit between — and each**

Findings

1. New York is the most top-heavy; Idaho the most retail

- **Defensible claim.** The top 1% of donors supply **47.5%** of New York's federal dollars versus **36.0%** in Idaho (Gini 0.848 vs 0.775) — with WA (39.3%) and TX (41.7%) between. Conversely, sub-\$200 gifts are **29.0%** of Idaho's dollars and 25.0% of Washington's but only **13.8%** of NY's, and ≥\$5,000 gifts are **34.8%** of NY's money vs ~20% of both ID and WA. New York's federal money is concentrated at the top; **Idaho's is the most broad-based of the four** (it edges out Washington on every retail measure), and it does so despite being deep-red — retail-vs-whale structure is not a partisan property. The two big-dollar states (NY, TX) are the most top-heavy; the two small-population states (WA relative to its dollars, ID absolutely) are the most retail.
- **Strongest objection.** The Gini of any voluntary-giving distribution is mechanically high everywhere (all four exceed 0.77), so the *level* is not itself pathological — only the *gap* between states is informative. Idaho's low concentration is partly a size effect: a \$76M pool simply has fewer mega-donors to concentrate around than a \$2B one. And the small-dollar share is sensitive to how conduit (ActBlue/WinRed) earmarks are recorded (see Limits) — though the gift-size *amount* cut is computed directly from the dollar value and is unaffected by that.

2. Participation is broadest, per capita, in Washington

- **Defensible claim.** New York and Texas raise ~3× Washington's federal dollars in absolute terms, but Washington has the widest donor *participation* relative to its size: ~362K donors in a state of ~7.9M (~4.6%) versus NY ~671K/~19.6M (~3.4%), TX ~837K/~30.5M (~2.7%), and ID ~54K/~1.96M (~2.8%). Fewer dollars, more givers. Idaho sits near Texas — small-population but not disproportionately participatory; WA remains the standout.
- **Strongest objection.** The name+zip donor proxy over-merges common names unevenly across states; population denominators are total residents, not voting-eligible adults; and WA's lower dollar total partly just reflects fewer ultra-wealthy households, not broader civic habit.

3. The retired-donor economy is largest in Idaho, then Washington

- **Defensible claim.** **31.7%** of Idaho's federal donor dollars come from donors who list their occupation as *retired* — the highest of the four — followed by **24.0%** in Washington, **19.5%** in Texas, and just **11.8%** in New York. Idaho's and Washington's federal money leans most heavily on people no longer in the workforce, consistent with both states' older donor bases; New York's Wall Street money is overwhelmingly still-earning. (A looser "non-working" bucket that also folds in *not-employed* / *none* / *blank* reaches ~48% in both ID and WA, but that figure is soft — see objection.)
- **Strongest objection.** FEC occupation/employer strings are self-reported and noisy. The retired-only figure is the defensible one; the broader "non-working" number bundles wealthy non-earners and blank/missing fields with the genuinely jobless and carries no income signal, so it should be read as a loose upper bound, not a measurement. Idaho's high retired share also partly reflects its retiree-destination demographics (Coeur d'Alene, north Idaho), not only a donor-composition effect.

4. Sector signatures: Tech (WA), Wall Street (NY), Energy/Industrial (TX), MLM/timber (ID)

- **Defensible claim.** The largest corporate employers of each state's federal donors are unmistakably regional:
- **WA — Big Tech:** Microsoft (\$15.2M), Amazon (\$5.0M), University of Washington (\$4.7M), Zumiez (\$4.2M), Fisher Investments (\$3.1M).
- **NY — Wall Street:** Blackstone (\$11.1M), Goldman Sachs (\$7.9M), KPMG (\$6.6M), Jane Street (\$5.8M), Soros Fund Management (\$5.4M).
- **TX — Energy / Industrial:** BNSF Railway (\$8.0M), Valero (\$7.0M), Starkey Hearing (\$5.6M), Beal Bank (\$5.1M), Lockheed Martin (\$4.6M).
- **ID — MLM / timber / regional:** Melaleuca (~\$2.8M across name variants — the Idaho Falls direct-marketing giant), Ball Ventures (\$0.8M), Idaho Forest Group (\$0.6M), University of Idaho (\$0.3M). A distinct small-state signature: consumer-MLM, timber, and regional development capital, at one-fifth to one-tenth the dollar scale of the big three.

The economic base that funds federal politics differs sharply by state. - **Strongest objection.** These employer sums are dominated by a handful of mega-donors at each firm, not broad rank-and-file employee giving; and the strings are free-text (Texas's single largest "employer" is the generic "*Entrepreneur*"; Idaho's Melaleuca total is split across three spelling variants), so firm-level totals are indicative, not audited.

5. A uniform presidential rhythm

- **Defensible claim.** All four states show presidential-cycle dollars running ~2× their off-year totals, in lockstep (WA \$202M/2020 vs \$79M/2018; NY \$612M vs \$299M; TX \$544M vs \$266M; ID \$22.5M vs \$7.6M). Federal giving is paced by the national calendar, not state-specific dynamics — Idaho, at 1/25th the scale, keeps the identical rhythm.
- **Strongest objection.** This is mechanical (presidential races simply cost more) and says nothing distinctive about any state — it's a useful uniformity check, not a finding.

Follow-on tests

Computed in `scripts/cross_state_fec_tests.py`. Committee master: 44,606 committees; 100% of recipient dollars matched in all four states.

Idaho scope (2026-07-19, completed). All tests A–I and the flow matrix G are now four-state. The analysis scripts were made **state-agnostic**: the region is discovered by globbing `data/*_statewide.duckdb` (overridable with `CROSS_STATE_REGION`) via the shared `scripts/cross_state_common.py`, so adding a state needs no script edits — load its data and re-run. ID's federal money is small (\$76M outflow, \$11.5M inflow, vs \$0.6–2.1B and \$155–582M for the big three), so it doesn't move the *qualitative* conclusions, but it sharpens the poles (most retail, most retired, most nationalized) and its Senate money is the most out-of-state of the four (85.8%).

A. Is the money concentrating over time?

Top-1% donor dollar share, by cycle:

Cycle	WA	NY	TX	ID
2018	28.2%	35.3%	29.2%	30.9%
2020	34.5%	42.4%	35.6%	34.4%
2022	30.3%	36.8%	34.9%	31.0%
2024	36.2%	47.4%	41.9%	30.0%
2026*	30.6%	40.3%	37.9%	31.1%

- **Defensible claim.** Concentration follows a **presidential sawtooth** (peaks in 2020 and 2024) with a **mild secular upward drift** — in three of the four states. Comparing like cycles: top-1% share rose presidential-to-presidential in WA/NY/TX (2020→2024: WA +1.7, NY +5.0, TX +6.3 pts) and midterm-to-midterm (2018→2022: TX +5.7, NY +1.4, WA +2.1). New York is both the most concentrated and rising fastest. **Idaho is the exception that sharpens the rule:** it is the *flattest* of the four and its top-1% share actually *fell* into 2024 (34.4%→30.0%) rather than spiking — its money base has no whale layer thickening at the top. So the secular concentration is a big-money-state phenomenon; the small retail state doesn't show it.
- **Strongest objection.** Only three presidential and two-to-three midterm points — too few to call a secular trend confidently. The 2024 spike is entangled with record presidential joint-fundraising activity (see Test B), so "concentrating" is partly a cycle-composition artifact, not purely a structural shift. Direction = mildly rising, not a clean climb. Idaho's flatness is also partly a small-n effect (fewer donors → noisier per-cycle top-1%). (2026 is *partial-cycle*.)

B. Where does each state's money go?

Destination of residents' federal dollars (recipient committee → state/office via the FEC committee master):

Destination	WA	NY	TX	ID
In-state Congress	13.4%	10.1%	16.0%	4.8%

Destination	WA	NY	TX	ID
Out-of-state Congress	24.1%	29.5%	19.4%	25.5%
Presidential	11.0%	7.0%	6.5%	7.8%
PAC / party / other	51.5%	53.4%	58.1%	61.9%

- **Defensible claim.** Residents fund **their own congressional delegation least of all** — just 10–16% of their federal dollars in WA/NY/TX, and a startling **4.8% in Idaho**. They give **2×–5× as much to out-of-state congressional races** (WA 24%, NY 30%, ID 25.5%), and the majority (51–62%) flows to **national party committees and joint-fundraising vehicles**. Washington's single largest federal destination is the Democratic presidential JFC (~\$61M) — on par with *all* in-state congressional giving combined (~\$87M); other top destinations are the DNC, DCCC, DSCC, RNC, and the Trump JFCs. **Idaho is the extreme of nationalization:** its in/out-of-state Congress ratio is **0.19** (it sends >5× more to other states' congressional races than to its own delegation) and 62% goes to national vehicles — a safe, small state with few competitive home races gives almost entirely to the national contest. Texans are the most "local" (ratio 0.82); New Yorkers 0.34; Idaho the least local of all. This is the **donor-side counterpart to the nationalization-of-money literature**: even constituents' own money is overwhelmingly aimed at national / out-of-state politics rather than their representatives — and the smaller and safer the state, the more so.
- **Strongest objection / caveat.** This remains **outflow** (where residents *send* money), not money flowing *into* each state's races. The majority "PAC/party/other" bucket is genuine national party + JFC money, but **JFCs blur the boundaries** — a presidential JFC funds the nominee *and* the party — so the "Presidential" row understates true presidential+national giving, and the splits among national vehicles are soft. The in- vs. out-of-state *Congress* split is the robust part (direct candidate-committee gifts).

C. Top donors, top recipients, and the cross-state magnets

Computed in `scripts/diag_cross_state_donors.py`.

Largest individual donors confirm the sector fingerprint at the person level: - **WA (tech/VC):** the Cornfields (~\$3.3M), Tom Campion/Zumiez (\$3.2M), Nick Hanauer (VC), Rory & Melinda Gates. - **NY (finance/philanthropy):** Philip Munger (\$4.1M), George Soros (\$4.0M), Stephen Schwarzman/Blackstone (\$3.2M), Agnes Gund/MoMA. - **TX (energy/industrial):** Syed Anwar/PetroPlex (\$5.2M), Paul Foster/Western Refining (\$3.5M), the Perots, Woody Hunt. - **ID (MLM/one-family dominance):** the **VanderSloot family** (Frank L. & Belinda, Melaleuca founder) supply ~\$4M+ across name/zip variants — the single dominant force in Idaho's federal giving — followed by the **Ball family** (Allen/Connie, Ball Group) and William Parks. Idaho's top-donor list is more concentrated in one household than any of the big three, the individual-level counterpart to its small, safe donor economy.

Largest recipient committees are national vehicles, with marquee in-state contests poking through: D-state dollars concentrate in the **Harris Victory Fund** (WA \$61M, NY \$176M) and **Fight for the People PAC** (NY \$82M); TX dollars in **Trump Victory** (\$71M) and the **RNC** (\$67M). In-state contests that surface: Kim Schrier (WA-08, \$17M), Patty Murray, Ted Cruz (TX, \$35M).

Cross-state money magnets. Of 12,361 committees these donors touch, **4,894 are funded by donors in ≥3 of the four states**. The top are national party / JFC vehicles (Harris Victory Fund \$292M combined; Fight for the People PAC \$171M; RNC \$115M; Trump Victory \$109M; DSCC / DNC / DCCC ~\$97–100M each). The single cleanest nationalization signal: **Warnock for Georgia** — a *Georgia* Senate race drawing **WA \$7.4M + NY \$16.3M + TX \$6.0M + ID \$0.3M**, funded by four states whose residents cannot vote in it. Idaho appears on essentially every magnet at small scale, and — unlike the D-heavy WA/NY or R-heavy TX top lists — its own top recipients split cleanly both ways (Idaho State Democratic Party \$4.3M *and* RNC \$3.6M / Trump JFC \$3.0M).

- **Caveat.** Individual-donor identity is a name+zip proxy (merges across cycles; a few employer labels are data-entry quirks). Recipient totals are donor-residence *outflow from these four states only* — not the committee's full national haul.

D. Does money chase competitive races? (money × competitiveness)

Computed in `scripts/diag_money_vs_competitiveness.py`, joining residents' U.S. House contributions to each district's competitiveness band (this project's own forecast margin: Tossup <5 / Lean 5–10 / Likely 10–20 / Solid ≥20). Full four-state run: WA/NY/TX/ID donors → WA/NY/TX/ID House districts, 2022–2026 (post-redistricting), 76 districts / \$298M. (An earlier version was a partial WA+TX run; the state-agnostic refactor now covers all four.)

Band	# districts	% of districts	\$ to band	\$ / district	% of \$	cross-state \$
Tossup (<5)	4	5.3%	\$27.6M	\$6.90M	9.3%	\$2.4M
Lean (5–10)	4	5.3%	\$27.9M	\$6.97M	9.4%	\$1.3M
Likely (10–20)	32	42.1%	\$123.6M	\$3.86M	41.5%	\$9.6M
Solid (≥20)	36	47.4%	\$118.9M	\$3.30M	39.9%	\$8.4M

- **Defensible claim.** With all four states in, donor money chases competitiveness **more clearly than the earlier partial run suggested but still modestly**: Tossup/Lean districts pull ~\$6.9M each vs \$3.3M in Solid — a **~2.1× premium** (up from the WA+TX-only ~1.4×, once New York's large donor base is included). Yet **81% of all dollars still flow to safe (Likely + Solid) districts**, because ~90% of districts are safe. In-state House money is dominated by support for (mostly safe-seat) candidates, not strategic targeting of the marginal race — the donor-side echo of "money follows the scoreboard." **Cross-state House giving is small (~\$21.7M of \$298M, ~7%)**: residents overwhelmingly fund their own state's House candidates. Idaho contributes only to the safe bands (cd02 R+16 → Likely, cd01 R+35 → Solid; no ID tossup), reinforcing that safe-state in-district money is a safe-seat phenomenon.
- **Caveats.** Donor-side *outflow*, not inflow (Section E is the inflow counterpart). The competitiveness bands are this project's 2026 forecast on current districts. Strategic targeting shows up more in PAC/JFC and out-of-state money (Test B's large national-vehicle bucket + the inflow side) than in this in-state-resident slice.

E. Inflow side — does money chase competitive races? (WA+NY+TX+ID House & Senate)

From the recipient-anchored inflow dataset (*fec_inflow.duckdb* : 5.50M contributions / \$1.21B, all-state donors → WA/NY/TX/ID federal candidates — ID added 2026-07-19 — built by *scripts/load_fec_inflow_bulk.py* in minutes; the API path would have taken days), joined to competitiveness. *scripts/diag_inflow_vs_competitiveness.py*, now four-state.

U.S. House, 2022–2026 — \$488M across 76 districts:

Band	# dists	% dists	\$ in	\$ / district	% of \$	out-of-state share
Tossup (<5)	4	5.3%	\$47.2M	\$11.81M	9.7%	45.0%
Lean (5–10)	4	5.3%	\$42.3M	\$10.59M	8.7%	36.1%
Likely (10–20)	32	42.1%	\$192.3M	\$6.01M	39.4%	40.0%
Solid (≥20)	36	47.4%	\$206.1M	\$5.73M	42.2%	44.2%

- **Defensible claims:** 1. **The competitiveness premium is real and ~2×**. Tossup (\$11.8M/district) and Lean (\$10.6M) districts pull roughly double the per-district inflow of safe seats (~\$6M) — money *does* chase the marginal race, far more clearly than the donor-side ~2.1× (Section D). 2. **But safe seats still capture ~82% of the money** (Likely+Solid), because they're ~89% of districts. Likely and Solid draw about the *same* per district (~\$6M): once a seat is safe, *how* safe barely changes the money — the jump is **between** competitive and safe, not within safe. 3. **~36–45% of all inflow is out-of-state, roughly uniform across every band**. Nationalization is **pervasive, not battleground-specific** — roughly two-fifths of the money funding these House races comes from people who cannot vote in them, in safe and tossup seats alike.

U.S. Senate, 2018–2026 (the model does not forecast US Senate; competitiveness via actual results):

State	\$ in	out-of-state share	Senate races in window
TX	\$253.2M	45.3%	competitive — Cruz/O'Rourke 2018 (R+2.6), Cruz/Allred 2024 (R+8.8)
NY	\$55.2M	53.5%	safe-D — Schumer / Gillibrand
WA	\$45.0M	41.1%	safe-D — Murray / Cantwell
ID	\$6.6M	85.8%	safe-R — Crapo / Risch

- **Senate echoes the House, louder.** Competitive **TX** Senate races drew **\$253M — ~5× safe NY (\$55M) or WA (\$45M)**: at the statewide level, competition is the single biggest money magnet. Yet out-of-state share is high

everywhere (41–54%) and is actually **highest in safe NY (53.5%)** — national donors fund high-profile safe senators (Schumer/Gillibrand) as readily as battlegrounds. Same lesson as the House: competition lifts the total, but the out-of-state flood is profile-driven and pervasive.

- **Idaho extends the pattern to the bottom of the size distribution — most sharply of all.** ID's federal candidates drew **\$11.5M** total inflow (House \$4.9M, all in the safe bands; Senate \$6.6M) — ~1/20th of Texas — yet its **Senate money is 85.8% out-of-state, the highest of the four** (WA 41%, TX 45%, NY 54%), and its House inflow lands only in Likely/Solid (no ID tossup exists). A small, safe, deep-red state's candidate money is *overwhelmingly* non-constituent: the "profile/incumbency pulls national money" mechanism operates at \$6.6M even harder than at \$250M — Crapo and Risch are safe, so almost none of their money needs to come from Idahoans. Nationalization is most extreme, not least, at the safe bottom.
- **Earmarks ARE attributed (verified — correcting an earlier caveat).** Conduit-routed (ActBlue/WinRed) money is **not** lost from these totals: FEC records each earmarked individual gift under the *candidate* committee as transaction type 15E — **\$194M for these candidates in 2024 alone, more than the \$90M of direct 15 gifts** — and the inflow load captures it. The conduit-side 24T records (\$150M) are the *same money* seen from the conduit and are correctly excluded to avoid double-counting (`scripts/diag_earmark_inspect.py`).
- **Caveats.** House competitiveness = 2026 forecast bands on current districts; Senate banded by actual two-party result. WA contributes no congressional Tossups in the 2026 map (its competitive seats land in Lean/Likely). State-legislative money still excluded (federal only).

F. The individual layer — who votes, who gives, and whether the donor skew is real (WA/NY/ID: donor skew F5 + giving→turnout F6)

Computed in `scripts/diag_wa_individual_findings.py` against `data/wa_statewide.duckdb` (`voter_scores` , `voter_donor_affiliation`) with `data/wa_vrdb.duckdb` ATTACHED (`voters` , 5.51M; `voting_history` , 27.1M). WA is the one state here that needs no external voter file — it already has the registered roll, the vote history, and the person-level voter→donor match. These refresh the democracy-insight gauntlet's figures on the current match (314,974 voters, 312,245 of them carrying a turnout score).

□ Panel and specification note (revised 2026-07-27). Two corrections apply to the F2–F4 figures below, and both are folded in.

First, these are computed on a pooled voter→donor match — WA's `individual_contributions` holds federal FEC money (\$646.2M) and state PDC filings (\$394.6M), so one person's federal and state giving stacks into a single donor total. Pooling inflates measured concentration. `docs/donor-class-and-the-electorate.md` reports WA as two separate panels — federal 147,745 donors / \$346.3M / top-1% 41.2% / Gini 0.815; state (PDC) 217,114 / \$122.5M / 43.5% / 0.821 — and those supersede the pooled numbers here for any donor-level claim.

Second, the match specification changed on 2026-07-27 to the full-first-name key alone, after a stratified blinded rating of 480 records found precision is a property of the match key: 100% on that key (120/120, Wilson [96.9–100]) against 47.9–71.7% on the three initial-based keys, which carried every household false merge. Every figure below has been recomputed on the new specification. WA pooled moved 382,408 → 314,974 voters (\$574.21M → \$468.85M). The direction of every F2–F4 finding is unchanged and most strengthen. Rebuild with `scripts/match_wa_voters_to_donors.py --source {fec,state,all} --tiers full`.

F1 — The off-year electorate is old, and the young drop out, not down. - Defensible claim. Within-cohort turnout and electorate composition both swing hard by cycle type. Turnout of 18–29s **collapses 58.4% → 15.8%** (2024 presidential → off-year generals), while 65+ falls only **88.3% → 61.3%**. As composition: 65+ are **28.5%** of the 2024 presidential electorate but **~39%** of the off-year electorate, while 18–29 fall from **14.2% → 7.6%** — a senior-to-youth ratio that widens from ~2:1 to ~5:1 off-cycle. The decision in odd-year local elections is made by a substantially older electorate than the one that shows up for president. - **Strongest objection.** This is *voluntary* differential participation under WA's all-mail, postage-paid, automatic-registration system — not disenfranchisement — and the fix (move local races on-cycle) is institutional, which is itself evidence of a *responsive* system, not a failing one. The within-cohort *rate* also uses a current-roll denominator (survivorship-biased), so read the composition **shares** as the robust cut. - **Caveat.** VRDB vote history begins in 2021, so the presidential figure rests on 2024 alone and the midterm on 2022 alone; only the off-year row averages three cycles (2021/2023/2025).

F2 — The donor age skew is genuine, not a matcher artifact (the white paper's open question, answered). - Defensible claim. Matched donors over-represent the old and under-represent the young — raw donor-share ÷ roll-share is **Silent 1.96×**, **Boomer 1.71×**, **Gen X 1.22×**, **Millennial 0.59×**, **Gen Z 0.17×**. The white paper's strongest objection was that this is a *matcher* artifact (the (last, first-initial, zip5) uniqueness guard over-selects older, rarer-named, stable-address voters). Tested directly, **it is not**: the probability a voter is uniquely matchable on that key is nearly flat across generations — **68.9% (Gen Z) to 73.1% (Silent)**, a **~4-point spread** — so inverse-propensity re-weighting barely moves the ratios (Silent

1.96→**1.91×**, Gen Z 0.09→**0.09×**). The age skew survives the correction; it is a real property of who gives, not of who the matcher can find. (Person-level concentration on this match: top-1% of matched donors = **46.6%** of matched dollars, top-10% = **79.3%**, Gini **0.857**; geographically, **61.4%** of matched-donor dollars come from just two Seattle-metro ZIP3s — 981xx 36.2% + 980xx 25.1%. All pooled-panel figures — see the note above; on the federal panel the multipliers run Silent **2.67×** to Gen Z **0.04×**, concentration top-1% **41.2%** / Gini **0.815**, and the two-ZIP3 share **63.5%**.) - **Strongest objection / residual bias**. The IPW corrects the *name-commonness* channel — the matcher's actual uniqueness guard — but cannot correct the one channel it can't observe: a donor whose address changed between giving and today fails the zip match, and mobility skews young. That residual bias runs in the *same* direction as the raw skew (it deflates the young donor share), so raw should be read as an **upper bound** on the true age skew — but the mechanism the objection actually named (rare names → easier match for the old) explains essentially none of it.

F3 — Money and votes stack on the same people (association only). - Defensible claim. Among the 312,245 matched donors carrying a turnout score, **87.6% are super-voters versus 50.9%** of non-donors (1.72×), and their mean turnout propensity is **0.967 vs 0.749**. Financial voice and electoral voice concentrate on the same individuals rather than offsetting. (Pooled panel; the split is essentially identical either way — federal **88.0% vs 52.0%**, state **88.9% vs 51.5%**.) - **Strongest objection**. This is cross-sectional association, not a causal "giving makes you vote" claim: donors are pre-selected for engagement (reverse causation is equally plausible), and the matcher itself favors stable-address super-voters, so selection *and* reverse causation both inflate the gap. The benign reading — donating as a participation *gateway* — is fully live.

F4 — Robustness: concentration precision + match validation. - Bootstrap CIs on concentration
(`scripts/diag_donor_concentration_bootstrap.py` , B=1000). Because donor identity is a name+zip5 *proxy*, the concentration estimates carry sampling-style uncertainty — but it is small. Matched-donor **Gini 0.857, top-1% 46.6%, top-10% 79.3%**; the inflow side (2024) is tighter still (Gini 0.690 [0.688–0.692]). Re-bootstrapped per panel on the current specification: federal **Gini 0.815 [0.806–0.822]**, top-1% **41.2% [38.6–43.4]**; state **Gini 0.821 [0.806–0.838]**, top-1% **43.5% [38.7–48.9]**. The concentration is a precise feature of the data, not an artifact of which donors landed in the pool. (The proxy's over-merging biases concentration *down*, so true figures are, if anything, slightly higher.) - **Match-precision validation** — superseded and now much stronger (`scripts/diag_match_validation_stratified.py` + `scripts/score_match_validation.py`). A **stratified blinded rating of 480 matched records** (20 per state × panel × tier, evidence file carrying no stratum labels, scorer written before the verdicts) found that precision is **a property of the match key**, not a single pooled number: the full-first-name key `STRICT_ZIP5_FULL` measured **100.0%** (120/120, Wilson 95% CI 96.9–100) while the three initial-based keys ran **47.9–71.7%**. Population-weighted across the panels, donor-weighted precision is **93.0%**. Of 152 confirmed-false matches, **129 were household/relative merges and every one of them sat on an initial-based key — none on the full-name key**. An independent human rater re-rated 150 of the 480 blind and **agreed on all 75 full-name-key records**. This is why the papers' primary specification is the full-name key alone; the pooled figures in F1–F3 above are the all-tier match and therefore carry the 93.0% weighted precision, not 100%. (*An earlier version of this bullet reported an adjudicated 150-match sample — "87% full-first-name agreement, 13% namesake risk, CSV emitted for human review". Those verdicts were never retained, so that pass cannot be re-derived; it is withdrawn in favour of the above.*)

F5 — The donor age skew is cross-state and cross-partisan (NY + ID confirm WA), and the donor class is less unaffiliated than the electorate.

Computed in `scripts/diag_cross_state_donor_representativeness.py` — the money-linked individual layer extended from WA to the two states that now have a voter file + a voter→donor match (NY 558K, ID 41K; WA 315K). Generation is derived uniformly from birth year (WA/NY `birthdate` , ID `age_snapshot`) so the states are comparable; `voter_scores` is unused (empty for NY/ID). Matched-donor share ÷ roll share per generation, RAW and inverse-propensity re-weighted by $P(\text{matchable} | \text{generation})$. TX excluded (no voter file). (When first written, `voter_scores` was empty for NY/ID and the turnout-linked cuts were WA-only; that gap is now closed — see F6.)*

Generation	WA raw / rwt	NY raw / rwt	ID raw / rwt
Silent	1.96× / 1.91×	1.50× / 1.45×	2.08× / 2.03×
Boomer	1.71× / 1.70×	1.65× / 1.63×	1.71× / 1.68×
Gen X	1.22× / 1.24×	1.21× / 1.23×	0.96× / 1.00×
Millennial	0.54× / 0.54×	0.58× / 0.59×	0.46× / 0.47×
Gen Z	0.09× / 0.09×	0.14× / 0.15×	0.10× / 0.10×

Recomputed 2026-07-27. This table and F6's had gone stale: the primary-specification switch rebuilt all three pooled matches (NY 308,032 → 558,017, ID 47,762 → 41,136, WA 382,408 → 314,974) and only the WA cells were refreshed. The

figures above are what `diag_cross_state_donor_representativeness.py` prints today. The correction *widens* the cross-state spread, so the claim below is weaker than the one it replaces — see the revised claim 1.

- **Defensible claim 1 — the age skew is a property of who gives in every state, but its steepness is not constant.** In all three states the oldest cohort (Silent) is over-represented among matched donors (**1.50–2.08×**) and the youngest (Gen Z) is sharply under-represented (**0.09–0.14×**). The *direction* is universal and the *magnitude* is not: the old-to-young gradient is **21.6×** in WA and **21.8×** in ID but only **10.5×** in NY, so New York's donor pool is materially less age-skewed than the other two. An earlier version of this section claimed the gradient was "essentially identical" across the three states; **that is withdrawn** — it was an artifact of the stale table, in which NY's Silent ratio read 1.87× rather than 1.50×. The skew does **survive the matcher-bias correction everywhere** (the IPW reweight moves every ratio by ≤ 0.05), which is the part that answers the WA white paper's open question: the skew is genuine, not an artifact of the uniqueness guard over-selecting older, rarer-named voters. Donor concentration is high in all three (Gini **WA 0.857 / NY 0.884 / ID 0.821**; ID lowest, matching its retail economy from Finding 1 — ID read 0.822 until 2026-07-27, a double-rounding of the diagnostic's 4-decimal 0.8215; the value is 0.821450). **ID here is the POOLED voter→donor match** (FEC + Sunshine, 41,136 donors) — an earlier version called it the FEC match, which it is not; the FEC-only panel has 23,303. The *state*-Sunshine-only match in the donor-class companion gives a slightly lower Gini, 0.798.
- **Defensible claim 2 — the donor class is less unaffiliated and more Democratic-tilted than the electorate** (the party-of-record cut WA cannot do). Matched-donor registered-party vs the full roll:
- **NY:** electorate **D 48% / R 22% / unaffiliated-or-other 30%** → donors **D 59% / R 25% / O 17%**. Donors are +11 pts more Democratic and little more than half as unaffiliated — the "donor class older *and* more Dem" finding, confirmed at the person level.
- **ID:** electorate **D 12% / R 63% / O 25%** → donors **D 20% / R 67% / O 13%**. Even in a deep-red state the donor pool is *more partisan than the electorate on both sides* (the unaffiliated quarter collapses to an eighth), and the Democratic share rises proportionally more (12→20%) than the Republican (63→67%).
- **Strongest objection.** The registered-party comparison mixes states with different party registration regimes (NY closed-primary registration vs ID's open system), so the *levels* aren't directly comparable across states — only the within-state donor-vs-roll *gap* is. And the match itself (Section F4) favors stable-address voters, who skew older and more partisan; the age IPW corrects the name-uniqueness channel but not residential mobility, so the age skew is an **upper bound** (as in WA). The direction and cross-state consistency are the robust part.

F6 — Giving→turnout replicates cross-state under identical definitions: donors are super-voters at ~1.6–1.8× the non-donor rate in all three voter-file states.

Computed in `scripts/diag_cross_state_giving_turnout.py` on `voter_scores` populated for NY/ID (2026-07-19, `scripts/populate_ny_id_voter_scores.py`) with WA-identical definitions — `is_super_voter = voted 2022+ AND registered ≥8 years`; `turnout_propensity` = the WA step/history blend. This closes the F1/F3 gap flagged in F5: the earlier NY giving→turnout cut used a different metric ("3 of last 4 generals"), so the ratios weren't comparable. Now they are.

State	Donors super-voter %	Non-donors	Ratio	Avg propensity (donor / non)
WA	87.6% (312.3K)	50.9% (5.14M)	1.72×	0.967 / 0.749
NY	83.1% (558.0K)	45.0% (12.98M)	1.85×	0.892 / 0.653
ID	50.0% (41.1K)	30.1% (988.8K)	1.66×	0.892 / 0.851

Recomputed 2026-07-27, same cause as F5: the NY and ID rows were still on the pre-switch pooled matches (308.0K and 47.8K) while WA's had been refreshed.

- **Defensible claim.** The engagement gap between the donor class and the rest of the roll is a **feature of every state measured, not a state peculiarity**: with identical definitions the donor/non-donor super-voter ratio runs **1.66–1.85×** across a blue mega-state, a Pacific vote-by-mail state, and a deep-red closed-primary state. (An earlier version called this a "tight 1.62–1.76× band"; on the corrected figures the band is wider and NY sits at its top, so "tight" is withdrawn — the robust claim is that every state shows a substantial gap in the same direction, not that they share a magnitude.) Within the donor class the gradient is generational everywhere (donor super-voter rates: Silent/Boomer ~0.55–0.93 vs Gen Z ~0.03–0.17) — the "money-linked electorate" is the same doubly-selected (old + habitual) slice in all three states.
- **Strongest objection.** Association only — donors are pre-selected for civic engagement, and the matcher favors

stable-address voters (who are also habitual voters), so the ratio blends behavior with selection in every state (consistently, which is why the *cross-state comparison* is informative even though the *level* is inflated). ID's absolute rates aren't comparable to WA/NY: its roll is a survivorship-biased current extract (who-decides-idaho §III) — the propensity floor is inflated (non-donor avg 0.851) and its 8-year-registration super-voter clause bites harder on a churned roll, which is why ID's levels sit lower while its *ratio* matches.

G. The cross-state money-flow matrix — the "Warnock" picture, made systematic

Computed in `scripts/diag_cross_state_money_matrix.py`. INFLOW from `data/fec_inflow.duckdb` (recipient-anchored, robust). OUTFLOW and the magnet list from the per-state `individual_contributions` joined to a committee→candidate master rebuilt from cached `cm*.zip` + `cn*.zip`. Recipient state = the candidate's office state (`cn.txt CAND_OFFICE_ST`), resolved committee→candidate — not the committee's registration state (`CMTE_ST`), which is often a DC/VA compliance-vendor address (Smiley-WA and Cornyn-TX both register in VA). Outflow and magnets are restricted to authorized candidate committees (designation P/A), so JFCs and leadership PACs — which carry one connected candidate but raise nationally — don't misattribute a recipient state. U.S. House+Senate; all cycles 2018–2026.

INFLOW — who funds each state's H+S candidates (by donor origin):

Recipient state	Total \$	In-state	In-region (other 3)	Rest-of-US
WA	\$154.6M	66.7%	6.6%	26.7%
NY	\$462.7M	55.1%	3.9%	41.0%
TX	\$582.4M	59.6%	6.1%	34.2%
ID	\$11.5M	31.9%	14.8%	53.3%

OUTFLOW — where each state's residents send their H+S candidate money (by recipient region):

Donor state	Total \$	In-state	In-region (other 3)	Rest-of-US
WA	\$198.7M	43.8%	5.4%	50.8%
NY	\$672.8M	33.9%	4.1%	62.0%
TX	\$526.5M	54.1%	2.6%	43.3%
ID	\$17.1M	21.4%	10.6%	68.0%

- Defensible claims.** 1. **Every one of these states sends the majority (or near it) of its candidate money out of region.** NY residents ship **62%** of their House+Senate candidate dollars to candidates in other states; **Idaho ships the most of all — 68%**; even Texas, the most parochial, sends **43%**. This is the donor-side engine of nationalization, measured directly rather than inferred. 2. **Idaho is the only net money-importer for its own candidates.** ID candidates raise just **31.9%** of their money in-state (vs 55–67% for the big three) and **53.3% from the rest of the US** — the highest external dependence of the four. But ID residents also send **68% of their money out** (\$17.1M out vs \$11.5M in): a small, safe, deep-red state whose handful of federal races draw national R money while its own donors chase the national contest. **In-region cross-funding stays negligible** — the four states barely fund *each other* (≤6–7% of any big-state cell; ID's 14.8% in-region inflow is just \$1.7M of tiny absolute dollars). NY is the dominant exporter (~\$418M leaves); WA and ID are net importers relative to their own giving. 3. **The magnet list is a clean battleground map — and Idaho rides along.** The out-of-region candidate committees funded broadly across the region are almost entirely Senate battlegrounds: **Warnock-GA (\$30.0M, incl. ID \$0.33M)**, **Kelly-AZ (\$19.9M, ID \$0.42M)**, **Ossoff-GA (\$19.5M)**, **Tester-MT (\$13.2M, ID \$0.40M)**, **Harrison-SC (\$11.8M)**, with Graham-SC, Perdue/Loeffler-GA, Gideon-ME, Rosen/Cortez Masto-NV and McConnell-KY close behind. **Georgia Senate races alone (Warnock + Ossoff + Perdue + Loeffler) drew ~\$68M** from residents of the four states — none of whom can vote in Georgia. Idaho's contributions are small (\$0.1–0.4M per race) but present across the same marquee list, and they **split both ways** — ID money shows up on the Democratic magnets (Warnock/Kelly/Ossoff/Tester) *and* the Republican ones (Graham, Perdue, Loeffler, McConnell), the only state here whose out-of-region giving visibly funds both parties' Senate battlegrounds. The Section-C Warnock anecdote is the rule, not the exception.
- Strongest objection / caveat.** The two matrices use different committee universes (inflow is the recipient-anchored bulk; outflow is per-state authorized committees), so their absolute totals are **not** directly subtractable — read the *shares within each*, not the cross-matrix difference. "Rest-of-US" inflow is robust (blank donor-state is 0.08% of inflow

dollars). Out-of-region ≠ "non-constituent and therefore illegitimate": a donor may have ties, and party committees (excluded here) carry yet more of the nationalized money.

H. Sector × competitiveness — does Wall Street chase Tossups while tech/energy fund safe seats?

Computed in `scripts/diag_sector_vs_competitiveness.py`, crossing the employer/sector signatures (Sections A/C) against the inflow competitiveness bands (Section E). Each inflow contribution to a WA/NY/TX/ID U.S. House district (2022–2026, \$488M) is classified by a keyword-on-employer sector and joined to the district's forecast band (Tossup<5 / Lean5–10 / Likely10–20 / Solid≥20). "Competitive share" = (Tossup+Lean) ÷ sector total.

Sector	\$ (House, 4 states)	Competitive share	Out-of-state share
Real estate	\$8.0M	21.9%	32.1%
Finance / Wall St	\$18.4M	21.6%	44.4%
Law	\$17.7M	20.0%	33.3%
Academia / public	\$8.1M	17.6%	48.3%
Healthcare	\$8.4M	16.1%	35.7%
Energy	\$3.6M	14.1%	26.6%
Tech	\$7.1M	11.6%	51.9%
[all sectors baseline]	\$488M	18.4%	~43%

Sector keyword map extended after a coverage audit (`scripts/diag_sector_coverage.py`) surfaced major law firms, hedge funds, and tech names as unclassified; the additions lifted classified dollars from ~14% to ~15% and left the pattern below intact.

- **Defensible claim.** The hypothesis is **partially borne out, with a twist**. Finance money does tilt toward competition (21.6% of its dollars to Tossup/Lean vs an 18.4% baseline), and **tech is the least competition-seeking sector (11.6%)**, with energy next (14.1%) — they do disproportionately fund safe seats. But the cleanest contrast is in *travel*, not band: **energy money is strikingly local** (only 26.6% out-of-state vs ~43% baseline) — it funds its own state's (Texas) incumbents — while **tech money travels farthest (51.9% out-of-state) yet lands in safe seats** (national safe-D House members). Finance both travels (44%) and tilts competitive. So: Wall Street chases the marginal race somewhat; tech funds safe co-partisans at a distance; energy funds the home-state incumbent.
- **Strongest objection.** The effect is **modest and partly mechanical**. The competitive-share spread is only ~10 points (real-estate 22% to tech 12%), and tech (WA) and energy (TX) are concentrated in their home states, whose *own* House seats happen to be safe (Solid-D Seattle, Solid-R Texas) — so "tech/energy fund safe seats" is in part just "their in-state seats are safe," not a strategic preference for safety.
- **Caveat.** Even after extending the keyword map with the biggest missing law firms, hedge funds, and tech names (`scripts/diag_sector_coverage.py`), classified sectors are still a **thin slice — only ~14% of inflow dollars**; "retired / not-employed / blank" is 46.8% and "unclassified" 38.6%, and tech (\$7.1M) and energy (\$3.6M) volumes in these four states' House races are small enough to be noisy. Employer strings are self-reported free text; the keyword map is indicative, not audited. House-only (Senate competitiveness isn't model-forecast).

I. Inflow concentration trend + donor retention — is the candidate-money base democratizing?

Computed in `scripts/diag_inflow_concentration_retention.py` on the inflow side (all-state donors → WA/NY/TX/ID H+S candidates; donor = name+zip5 proxy). Section A measured concentration on the outflow; this measures it on the inflow, and adds repeat-vs-one-time retention. State-agnostic: it simply reflects whatever states are in `fec_inflow.duckdb` (ID folded in 2026-07-19 — its ~\$12M barely moves these all-state figures).

Cycle	Inflow \$	Donors	Top-1% \$	Top-10% \$	Gini
2018	\$254M	246K	16.2%	52.4%	0.651
2020	\$251M	226K	17.2%	54.2%	0.672
2022	\$250M	219K	15.8%	55.8%	0.688

Cycle	Inflow \$	Donors	Top-1% \$	Top-10% \$	Gini
2024	\$285M	311K	18.4%	58.5%	0.690
2026*	\$170M	168K	16.5%	58.2%	0.686

- **Defensible claim 1 — candidate-directed money is far less concentrated than the total flow.** Inflow to candidates (gifts capped at the per-election limit) runs **top-1% $\approx$ 16–18%, Gini $\approx$ 0.69** — against the **outflow** figures of top-1% **39–48%** and Gini **0.80–0.85** (Sections A & F, which include uncapped JFC/PAC/party money). The system's dollar concentration lives in the **party/JFC layer, not in direct candidate giving**, which is structurally egalitarian by the contribution cap. There is a **mild secular rise even within candidate money** (Gini 0.651→0.690, top-10% 52→59% across 2018→2024), but no presidential sawtooth — concentration here drifts up gently and plateaus.
- **Defensible claim 2 — a churning, mostly one-time base funds the candidates, but a persistent minority supplies most of the dollars.** Across 2018–2026, **78.2% of donors give in only one cycle** (supplying **41.6%** of dollars); the **21.8%** who give in ≥ 2 cycles supply **58.4%**. And cycle-over-cycle retention is **low and flat** — **~ 21 –25%** of a cycle's donors gave in the immediately prior cycle (2020 23.4%, 2022 25.0%, 2024 21.2%), with **no rising trend**. Roughly three-quarters of each cycle's candidate donors are new-or-returned-from-the-past, not a stable subscription base. The "small-dollar democratization" picture — broad, churning participation — holds by *headcount*; the "thin layer over a concentrated core" picture holds by *dollars*.
- **Strongest objection / caveat.** The naive "returning %" *looks* like it rises (23→39%), but that is a **look-back-window artifact** — later cycles have more prior cycles to match against; the fixed one-cycle look-back (above) removes it and shows flatness. The name+zip5 donor proxy **over-merges** common names, which inflates "repeat" and deflates one-time counts, so true churn is, if anything, *higher* than reported. **2026 is a partial cycle** (its elevated retention reflects early givers skewing toward committed repeat donors before the late small-dollar surge). Capped candidate giving $\neq$ the whole money system — the uncapped layer (Section A) is where concentration concentrates.

J. Which side of a safe seat gets the money? (longshot vs favored)

New cut in `scripts/diag_loser_side_money.py`; WA/NY/TX/ID U.S. House inflow 2022–2026, recipient party from the committee→party map (97.9–100% of dollars resolvable per state), favored party from the forecast margin. Four-state since 2026-07-19 — TX needed only its FEC committee master loaded (data parity, no code change); the script is state-agnostic via `cross_state_common.py`.

A safe seat for one party is a longshot for the other, so we can ask which side the money actually reaches. Almost all of it goes to the favored side, and the safer the seat, the more lopsided the split. In New York the longshot party's share of House inflow falls from 74% in the lone tossup to 27% in Likely seats to just **5.5%** in Solid (≥ 20 -point) seats; Texas runs the same staircase from the other side (57.3% Tossup → 24.2% Lean → 14.0% Likely → **2.8%** Solid, on \$159.9M), and Washington matches (24.6% Likely → **4.8%** Solid). Put plainly: in a truly safe district the disadvantaged party's candidate raises about a nickel on the dollar and the favored side takes the rest. (Caveats: this is money entering the race, House only; leadership PACs tied to safe incumbents can pad the favored side; "favored" is only meaningful once a seat is actually safe — in a real tossup the challenger can out-raise the nominal favorite, which is why the tossup rows look inverted; and Idaho's two districts carry too little money — \$2.8M total — to read, its Solid row being a single small district where one funded longshot flips the percentage.)

K. State-level money — the first cut over the state-disclosure layer

Computed in `scripts/cross_state_state_money.py` (state-agnostic via `cross_state_common.py`) on the STATE-disclosure rows now in each state's `candidate_finance`: WA PDC (PDC: prefix), NY BOE (NY:), TX TEC (Tx:), ID Sunshine (SUNSHINE:). Every section above is federal; this is the first analysis of the state layer. The apples-to-apples core is state-HOUSE candidate money (every lower chamber is fully up each cycle; senates are staggered heterogeneously), pooled over the 2022+2024 cycles (post-redistricting maps; 2026 is in-cycle and excluded). Gap-filling mirrors the donor-side backfills: TX legislative party and ALL of ID's office/district/party are resolved by joining filer names to the SoS candidates roster (unique-party guard); ID resolves 65% of person-filer dollars, TX 59% of legislative dollars.

What each state's system captures (not the same universe — read before comparing):

	WA (PDC)	NY (BOE)	TX (TEC)	ID (Sunshine)
Rows / \$	16,604 / \$1.52B	20,349 / \$1.73B	19,416 / \$5.17B	2,067 / \$55M

	WA (PDC)	NY (BOE)	TX (TEC)	ID (Sunshine)
Cycles	2014–2026	2014–2026	2014–2026	2020–2025
Universe	full filer universe	full filer universe	full filer universe	full filer universe
Office labels	native (legislative rows)	native (partial)	native (partial)	none → roster join
Party labels	native (legislative rows)	native	mostly Unknown → roster join	none → roster join

(WA was originally legislative-candidates-only; the full filer universe — statewide execs, party/caucus committees, PACs, ballot committees — was bulk-loaded 2026-07-19 from the PDC Campaign Finance Summary via `load-pdc-filer-universe`: one CSV export request, 43K filer-years, ~13 seconds. The legislative rows keep their richer office/party/individual-PAC detail; K1–K4 are computed on those rows and are unchanged by the expansion.)

K1 — The price of a house seat varies 26× across the four states. Per-seat-per-cycle candidate receipts (2022+2024): **TX \$978K » WA \$253K > NY \$138K > ID \$37K**. Median candidate raise: TX \$85.5K ≈ WA \$83.6K > NY \$45.0K > ID \$13.7K (TX's mean, \$353K, is 4× its median — Texas sets **no dollar limit** on individual contributions to non-judicial state candidates, barring only corporate and labor-organization money (Tex. Elec. Code § 253.094; the Judicial Campaign Fairness Act, §§ 253.151–253.176, is the one capped exception), and that produces a whale-candidate layer WA's capped system lacks — though note that caps do *not* flatten donor-level concentration, which is tested in [donor-class-and-the-electorate.md](#) Appendix G). A competitive TX house seat costs more than most **congressional** seats' inflow; an ID house seat costs less than a WA school-board race.

K2 — Does state money chase competitiveness the way federal money does (~2×, Section D)? Only in the red states. Per-district house-candidate \$ by forecast band, with the (Tossup+Lean)/Solid premium:

Band (\$/district)	WA	NY	TX	ID
Tossup	\$0.99M	\$0.30M	\$3.52M	\$0.37M
Lean	\$1.28M	\$0.23M	\$2.13M	\$0.23M
Likely	\$1.16M	\$0.18M	\$1.65M	\$0.17M
Solid	\$0.93M	\$0.31M	\$1.68M	\$0.13M
Premium	1.12×	0.85×	1.72×	2.49×

- **Defensible claim.** TX and ID look federal-like (competitive seats draw ~1.7–2.5× the candidate money of safe ones), but **WA is nearly flat (1.12×**) and **NY inverts (0.85×** — **safe Assembly seats out-raise the battlegrounds per district**). The mechanism is visible in K5: in NY the battleground money runs through the **party campaign committees** (DACC/DSCC/NYSSRCC/RACC are 4 of NY's top-10 filers, \$33–43M each) rather than through candidate committees, and safe-seat incumbents (leadership, chairs) raise heavily anyway; the candidate-committee lens undercounts exactly the competitive spending NY routes around it. WA sits between the poles: its eight **caucus committees** (HDCC/HROC, the Senate Truman/Kennedy funds, Leadership Council/WA WINS, WSDC/SRCC) raised **\$36.2M in 2022+2024 — about half the \$72.9M all legislative candidates raised combined** — a real routing-around-candidates layer, milder than New York's, which (together with WA's contribution caps compressing what any single hot race can absorb directly) is consistent with its flat candidate-level premium. That is a claim about what a *candidate committee* can take in, and it does not extend to the shape of the donor pool: caps demonstrably fail to flatten donor-level concentration ([donor-class-and-the-electorate.md](#) Appendix G).
- **Strongest objection.** The forecast bands are 2026 labels applied to 2022+2024 money (bands are stable, but not identical, across cycles); WA's Lean band holds a single district; and NY's office classification is partial — if BOE office-coding is biased toward incumbents, safe-seat money is inflated. Read the WA/NY numbers as "no large candidate-level premium," not as precise ratios.

K3 — State-legislative money mirrors chamber control, hard. Democratic share of party-resolved legislative receipts (2022+2024): **NY 74.8% D, WA 60.9% D, TX 27.4% D, ID 23.5% D**. Unlike federal inflow (which both parties export to the same national battlegrounds), state money stays home and tracks who holds/contests the chamber. (TX carries \$130M party-unresolved after the roster join; ID resolves 65% of person-\$ — directional reads only for those two.)

K4 — State legislatures are PAC-funded; Congress is individual-funded. Individual share of legislative candidate receipts (2022+2024): **WA 31.4% (PACs 62.4%), TX 35.8% (entities 64.2%), NY 38.2% (PACs 29.5% + party/other transfers**

32.3%), ID 49.9%. The federal benchmark from the same DBs: House candidate committees run ~**65% individual / 16–25% PAC** (WA H 64.6%/25.3% on \$104M; NY H 67.2%/15.7% on \$369M). The individual small-donor era documented in Sections A–I is a *federal* phenomenon — statehouse campaigns are still funded chiefly by organized money. (TEC's split is individual-vs-entity by construction; "PAC" there means all non-individual money.)

K5 — The panorama: what the full filer universe shows (all four states). - **TX:** Texans for Greg Abbott has raised **\$424.5M across 6 cycles** — 1.4× the entire TX House candidate universe over 2022+2024 combined. ActBlue Texas (\$168.6M) and Dan Patrick (\$106.9M) follow; Texans for Lawsuit Reform (\$105.5M) is the biggest non-candidate PAC. Org-name filers hold 2/3 of all TX state money (\$2.9B vs \$1.5B). - **NY:** Hochul \$68.5M and Cuomo \$48.7M top the list, but the signature is the **committee layer**: the state Democratic Committee, both Senate campaign committees, and three union PACs (NYSUT, SEIU, UNITE HERE) occupy most of the top 10. - **ID:** the top filer is a **ballot-measure committee** (Idahoans for Open Primaries, \$5.6M — more than any Idaho candidate; the sitting governor raised \$1.2M over 3 cycles), and orgs hold 2/3 of ID state money. Consistent with Section F5's picture of a small retail donor base: even Idaho's biggest money fights are initiative fights, not candidate fights. - **WA:** the top filer is a **union PAC** — SEIU's Political Education and Action Fund (\$39.2M across 12 years) — followed by New Direction PAC (\$32.2M) and two nine-figure-fight ballot committees (No on 1631, the oil-industry carbon-fee opposition, \$31.6M in a single year; Yes! To Affordable Groceries, \$22.4M). No candidate cracks the top ten: WA's biggest single candidate haul ever is Bob Ferguson's ~\$19M 2024 governor run, roughly half of SEIU's cumulative total. Committees hold **69% of all WA state money** (\$1.05B of \$1.52B) — the same ~2/3 org share as TX and ID. The caucus committees sit just below the mega-PACs (HDCC \$19.2M and the Truman Fund \$18.7M across 12 years) — the "half the candidate layer" routing channel quantified in K2.

Limits of inference

Full provenance + a no-AI reproduction recipe (every source, access path, and the exact scripts behind each figure): [Data Sources & Reproducibility](#).

- **Both directions now exist (Findings 1–5 + Tests A–D are outflow; Tests E–I add inflow).** The early findings measure where residents *send* money (outflow by donor residence); the recipient-anchored inflow load (Sections E–I) measures money *entering* WA/NY/TX/ID races. The cross-state matrix (G) reports both. They use different committee universes, so read shares within each rather than subtracting one from the other.
- **Sections A–J are federal; K is the state layer.** State-legislative and local money is loaded for all four states (WA PDC, NY BOE, ID Sunshine, TX TEC — item 7) and Section K is the first cut over it. The two layers use different disclosure regimes and filer universes (WA's is legislative-candidates-only) — compare *within* K, not K against A–J dollar figures.
- **Conduit-reliance metric is unreadable (but the money is captured).** ActBlue/WinRed appear as recipients of <0.5% of dollars — because, as the 2024 earmark inspection confirmed (Section E), FEC records each earmarked gift under the *candidate* committee as type 15E, not under the conduit. So you cannot measure conduit *usage* from the recipient field — but the **money itself is fully captured** (under the candidates). **Do not read the ActBlue/WinRed-as-recipient share as "conduit reliance."** The sub-\$200 *amount* share (Finding 1) is computed from the dollar value and is unaffected.
- **Donor proxy.** name+zip5 over-merges common names, modestly understating donor counts and overstating concentration.
- **WA & ID composition.** WA's FEC subset draws from both a donor-filtered bulk load and an earlier per-candidate load; ID's is bulk-only (donor-filtered, loaded 2026-07-19) and shares its `individual_contributions` table with state Sunshine rows, which the FEC-committee-id regex filters out. Both are restricted here to in-state residents, but completeness may differ slightly from NY/TX pure-bulk.
- **No income/race.** Occupation/employer is the only socioeconomic signal, and it is self-reported.

What's done, and what's next

Tests A–I are run. Status: 0. **Idaho fully integrated — DONE 2026-07-19** (all sections A–I + matrix G): ID FEC outflow + inflow loaded to parity; it is the small, deep-red, most-retail, most-retired, most-nationalized pole (safe-R Senate money 85.8% out-of-state). The analysis scripts were made **state-agnostic** in the same pass (item 8). 1. **True inflow — DONE for WA+NY+TX+ID, House + Senate** (Section E): recipient-anchored *bulk* load (`scripts/load_fec_inflow_bulk.py` → `fec_inflow.duckdb` , **5.50M contributions / \$1.21B** of all-state money into WA/NY/TX/ID federal candidates) — built in minutes vs. the API path's days. 2. **Conduit/earmark attribution — DONE/verified** (Section E): earmarked ActBlue/WinRed

money is attributed to candidates via FEC 15E and already counted (\$194M in 2024); conduit-side 24T duplicates correctly excluded. No fix needed. 3. **Individual voter-file study — DONE, WA + cross-state donor layer** (Section F): WA has the full stack (turnout-by-age F1, donor representativeness F2 with matcher-bias re-weighting, giving↔turnout F3). **F5 (2026-07-19) extends the donor-representativeness cut to NY + ID** (voter files + matches now loaded): the donor age skew is near-identical and survives the IPW correction in all three, and the party-of-record cut shows donors less unaffiliated / more Dem-tilted in NY and ID. **F6 (2026-07-19) closes the turnout half:** NY/ID `voter_scores` populated with WA-identical definitions (`scripts/populate_ny_id_voter_scores.py` — NY 13.54M rows from the parsed history, ID 1.03M from the melted participation table), and the giving→turnout cut replicates at a substantial donor/non-donor super-voter ratio in every state (WA **1.72×** / NY **1.85×** / ID **1.66×** — see F6; "near-constant" overstated it and is withdrawn). The individual layer is now complete for all three voter-file states. 4. **Cross-state flow matrix — DONE, now four-state** (Section G): inflow provenance + outflow destination + the systematic out-of-region magnet list (Georgia Senate ~\$68M from WA/NY/TX/ID residents). ID is the biggest out-of-region exporter (68%) and the only net candidate-money importer. 5. **Sector × competitiveness — DONE** (Section H): finance tilts mildly competitive; tech/energy fund safe seats; energy money is distinctively local. Modest, on a thin classified slice. 6. **Inflow concentration + retention — DONE** (Section I): candidate money is far less concentrated than total flow (Gini ~0.69 vs ~0.80); base is churning/one-time by headcount, concentrated-core by dollars; cycle-over-cycle retention flat ~21–25%. 7. **State-level money — DONE for all four states** (2026-07-19). WA (PDC), NY (NYSBOE `ny_finance.py`, 20,349 rows), ID (Sunshine), and now **TX** are all in `candidate_finance`. The new **TEC adapter** (`etl/adapters/tx_tec.py`) streams the Texas Ethics Commission bulk zip (`TEC_CF_CSV.zip` → `contribs_###.csv` / `expend_###.csv` / `filers.csv` / `cover.csv`) via DuckDB into the `StateFinanceAdapter` framework with a TX: id prefix: **19,416 candidate-cycle rows, 2014–2026, ~\$5.2B receipts** (top: Texans for Greg Abbott \$120M/2022, Beto for Texas \$98.6M, ActBlue Texas \$100M). Individual/PAC split from `contributorPersentTypeCd`; party is mostly Unknown (TX has no party registration) and resolved downstream. Tests in `tests/test_etl/test_tx_tec_adapter.py`. 8. **State-agnostic scripts — DONE 2026-07-19:** all 8 cross-state scripts now enumerate the region via `scripts/cross_state_common.py` (glob `data/*_statewide.duckdb`, or a `CROSS_STATE_REGION` override) instead of hardcoded state lists; the competitiveness read and the reports/ output path are shared, and the magnet "broadly funded" threshold scales with the region. **Adding a state N needs no script edits** — load `data/N_statewide.duckdb` (+ VRDB) via the existing loaders, run `FEC_INFLOW_STATES=N python scripts/load_fec_inflow_bulk.py`, then re-run the 8 scripts. 9. **State-level money analysis — DONE 2026-07-19** (Section K, `scripts/cross_state_state_money.py`): first cut over the state-disclosure layer loaded in item 7. Headlines: house seats cost 26× more in TX than ID; the federal ~2× competitiveness premium holds only in TX/ID (WA flat, NY inverted — party committees route around candidate committees); statehouse money is PAC-funded (~31–38% individual) vs Congress's ~65% individual; TX/ID's biggest state-money filers are a governor's committee (\$424.5M) and a ballot-measure committee respectively. **Follow-on completed same day:** the full WA PDC filer universe (statewide execs, party/caucus cmtes, PACs, ballot cmtes) bulk-loaded via the new `load-pdc-filer-universe` CLI (one Socrata CSV export request — 43K filer-years, ~13s; WA 1,559→16,604 rows / \$1.52B), giving WA its K5 panorama row (top filer = SEIU's PAC; committees hold 69% of WA state money) and quantifying the caucus-committee routing layer (\$36.2M ≈ half the candidate layer, 2022+2024).

Related work

The building blocks — donor concentration, out-of-district money, whether money chases competitive races — are established individually; the contribution here is the recipient-anchored, four-state, federal-and-state comparison built from bulk FEC and state-disclosure records on one harmonized frame. It sits in these literatures:

- **Mapping the donor universe.** Bonica, "Mapping the Ideological Marketplace" (2014) and the DIME database — the individual-contribution infrastructure this paper rebuilds from bulk filings and joins to voter files (Section F).
- **How much money, and why so little.** Ansolabehere, de Figueiredo & Snyder, "Why Is There So Little Money in U.S. Politics?" *Journal of Economic Perspectives* (2003) — the framing for the money-and-competitiveness cuts (Sections D, E, H): contributions behave more like consumption/participation than investment.
- **Cross-district and out-of-state money.** Gimpel, Lee & Pearson-Merkowitz, "The Check Is in the Mail: Interdistrict Funding Flows in Congressional Elections," *American Journal of Political Science* (2008) — the direct antecedent to the inflow/outflow matrix and the out-of-region magnet list (Sections B, E, G).
- **Donors, polarization, and influence.** Barber, "Ideological Donors, Contribution Limits, and the Polarization of American Legislatures," *Journal of Politics* (2016); Schlozman, Verba & Brady, *The Unheavenly Chorus* (2012); Bonica, McCarty, Poole & Rosenthal, "Why Hasn't Democracy Slowed Rising Inequality?" *JEP* (2013) — the concentration and

Gini results (Sections F, I).

- **State campaign finance.** La Raja & Schaffner, *Campaign Finance and Political Polarization: When Purists Prevail* (2015) — context for the state-disclosure layer and the PAC-vs-individual split across statehouses (Section K).
- **Party and the donate.** Grumbach & Sahn, "Race and Representation in Campaign Finance," *American Political Science Review* (2020) — the donor-composition frame for the individual layer (Section F).