

The Donor Class Is Not the Electorate

Who funds elections in Washington, New York, and Idaho — old, concentrated, and (where party is observable) skewed toward Democrats in a blue state *and* a red one — measured from individual voter-to-donor matches

Stephen Kirby · Tikor Consulting · July 2026

AI-assisted drafting and analysis review. All figures are reproducible from public-record data available through lawful request and from the open-source scripts cited below, including `scripts/verify_donor_class.py` and `scripts/diag_contribution_limits.py`. The paper source, code, and data-acquisition recipe are public at <https://github.com/skirby359/who-decides>; the underlying voter files are not redistributed. Contact: kirby@tikorconsulting.com.

*Paper #3 of the electoral-health series (companion to [who-decides-washington.md](#), [who-decides-new-york.md](#), [who-decides-idaho.md](#), [safe-seat-washington.md](#), and [cross-state-fec-money.md](#)). DRAFT — pending the independent-verification gate in [publication-checklist.md](#).***

Provenance. Each panel is built by a per-state match script run once per money source (`--source fec / --source state`), writing to `voter_donor_affiliation_fec` and `voter_donor_affiliation_state`; `scripts/verify_donor_class.py` re-derives both from scratch SQL. Washington figures: `scripts/match_wa_voters_to_donors.py` and `scripts/diag_wa_individual_findings.py` — WA's registered roll (5.51M) + 27.1M vote records + birthdates (`data/wa_vrdb.duckdb`), matched to 172,998 federal and 269,204 state donors. New York figures: `scripts/match_ny_voters_to_donors.py`, `scripts/backfill_ny_committee_party.py`, `scripts/diag_ny_match_bias.py`, `scripts/diag_ny_primary_participation.py`, `scripts/diag_ny_donor_extras.py`, `scripts/diag_ny_electorate_extras.py` — NY's NYSVOTER roll (13.54M; individual party enrollment + DOB; `data/ny_vrdb.duckdb`) matched to 10.02M FEC itemized contributions and, via `scripts/load_ny_contributions.py`, to 3.95M NYSBOE state contributions (`data/ny_statewide.duckdb`, 307,841 federal / 424,020 state matched voters). Idaho figures: `scripts/match_id_voters_to_donors.py`, `scripts/backfill_id_recipient_party.py`, `scripts/diag_id_electorate_extras.py` — ID's statewide roll (1.03M; individual party affiliation + age; `data/id_vrdb.duckdb`) matched to both Idaho Sunshine state filings (27,250 voters) and federal FEC contributions (27,196 voters), in `data/id_statewide.duckdb`. Cross-state dollar concentration: `scripts/cross_state_fec_money.py`. Contribution-limit tests (Appendix G): `scripts/diag_contribution_limits.py`. Each figure below traces to one of these scripts.

Two panels, read this first. American donors give into two separate money systems governed by different rules, and a voter roll can be matched to either. Every finding below is therefore computed twice, once per system, and never pooled:

panel	what it is	matched donors
Federal (primary)	FEC itemized individual contributions	WA 172,998 · NY 307,841 · ID 27,196
State (secondary)	WA Public Disclosure Commission · NY State Board of Elections · Idaho Sunshine	WA 269,204 · NY 424,020 · ID 27,250

Pooling the two systems — which earlier drafts did for Washington without intending to — inflates measured concentration, because one person's federal and state giving stacks into a single donor total while a one-system donor's does not; on Washington's data, pooling reads 46.6% top-1% against 42.4% federal and 43.8% state. The panels are kept separate for exactly that reason.

This is a strength, not a limitation. Federal and state money are capped differently, solicited by different campaigns, and reported by different agencies. Every finding in this paper holds in both — the donor class is old, top-heavy, metro-concentrated, partisan-skewed, and stacked with high turnout whether the money is federal or state. Where the panels differ, they differ informatively. One difference holds in all three states: federal money is older money — New York's federal donors are 47.9% over 65 against 38.4% of its state donors, Idaho's 64.7% against 51.1%, and Washington's Silent Generation multiplier runs 2.56× federal against 1.64× state. Two more hold in the two large states and invert in Idaho: state money reaches more donors (New York 424,020 state against 307,841 federal; Washington 269,204 against 172,998; Idaho essentially level at 27,250 and 27,196), and federal money concentrates harder in the primary metro (Manhattan supplies 50.3% of New York's federal dollars but 20.6% of its state dollars — while in Idaho it is Boise's grip that loosens, from 49.2% to 36.4%). Where the comparison is available, state giving is the broader, younger, more geographically dispersed layer — and still unrepresentative on every axis measured here.

Abstract

Campaign money is usually described by how much is raised. This paper asks whose money it is, and whether the people who fund elections resemble the people who vote in them. Three state voter-registration files — Washington (5.51M registrants), New York (13.54M, with individual party enrollment and date of birth), and Idaho (1.03M, with party affiliation and age) — are linked person by person to itemized campaign contributions under a conservative uniqueness rule. Because American donors give into two separately regulated money systems, every result is computed twice and never pooled: a federal panel (172,998 matched donors in Washington, 307,841 in New York, 27,196 in Idaho) and a state panel (269,204, 424,020, and 27,250). Every finding holds in both. The matched donor class is markedly older than the electorate that elects the officials it funds: 47.9% of New York's federal donors and 64.7% of Idaho's are 65 or older, against a quarter and a third of their rolls, and Washington's Silent Generation gives at 2.6 times its share of the roll while Generation Z gives at one-tenth of its own. Money concentrates sharply, with the top 1% of matched donors supplying 42.4% of federal dollars in Washington, 51.2% in New York, and 35.8% in Idaho, and a single metropolitan area supplying a fifth to a half of each state's total. Where party of record is observable, the donor class over-represents registered Democrats and under-represents the unaffiliated by double digits — in deep-blue New York (+15.0 points Democratic federally, +8.9 in state money) and in deep-red Idaho (+8.0) alike, so the tilt is not an artifact of a state's majority party. The same individuals also vote at far higher rates than non-donors, stacking financial and electoral voice rather than offsetting them. Comparing the two panels isolates a difference present in every state — federal money is older money — and two more that hold in the two large states: state money reaches more donors, and federal money concentrates far harder in the primary metropolis. The age skew survives inverse-propensity re-weighting for match bias, and hand rating of a 150-record sample puts match precision near 90%. Contribution caps do not explain the cross-state differences in concentration. The paper measures itemized giving and registration records, not policy influence, and its causal claims are limited to association.

Keywords: campaign finance; political donors; voter files; record linkage; party of record; contribution limits; donor concentration; Washington; New York; Idaho.

The question

Campaign money is usually described by *how much* is raised. The more consequential question for representation is *whose* money it is — and whether the people who fund elections look anything like the people who vote in them. If the donor class is a representative cross-section of the electorate, money is just amplified participation. If it is a narrow, unrepresentative slice, then a distinct and self-selected population is setting the financial terms of every race before the first ballot is cast.

This question can be answered at the individual level in three states, by matching the registered-voter roll to itemized donors, person by person (a conservative name + ZIP match; see Appendix C). Washington supplies the demographic and behavioral cut; **New York and Idaho — which, unlike Washington, publish each voter's party — supply the dimension WA cannot: who the donor class is *partisan-ly*, and where its money goes.** And they bracket the political spectrum: New York is ~48% registered Democratic, Idaho ~63% registered Republican.

The short answer, consistent across all three states: **the donor class is not the electorate.** It is markedly older, geographically and financially concentrated in a small top tier, and — where party is observable — skewed toward Democrats while nearly excluding the largest non-partisan bloc. The striking part is that the Democratic tilt of the donor class relative to the electorate appears in deep-blue New York **and in deep-red Idaho**: it is not an artifact of a state's majority party.

Finding 1 — The donor class is old, and the skew is real

In all three states matched donors are far older than the voters they fund.

New York (`match_ny_voters_to_donors.py`) — age-band share, age as of 2024-11-05:

age band	federal donors	state donors	all active voters	2024 GE voters
18–29	3.0%	4.9%	18.0%	14.1%
30–44	14.1%	17.8%	25.6%	23.1%
45–64	34.9%	38.9%	31.2%	34.6%

age band	federal donors	state donors	all active voters	2024 GE voters
65+	47.9%	38.4%	25.2%	28.2%

Nearly **half of NY's federal donors are 65 or older**, versus a quarter of the active roll; its state donors are younger but still tilted, at 38.4%. **Washington** shows the same shape as generation multipliers (donor share ÷ roll share), and the skew is sharper in federal money than in state:

generation	federal panel	state panel
Silent	2.56×	1.64×
Boomer	1.97×	1.51×
Gen X	0.98×	1.26×
Millennial	0.39×	0.68×
Gen Z	0.10×	0.20×

Idaho replicates it in a third state, in both money systems (age here is current-roll age, not election-time DOB, so bands are read against the current roll):

age band	federal donors	state donors	all voters	2024 GE voters
18–29	1.4%	2.6%	15.2%	13.1%
30–44	7.1%	13.1%	22.8%	22.4%
45–64	26.8%	33.2%	30.9%	31.9%
65+	64.7%	51.1%	31.0%	32.6%

Nearly **two-thirds of Idaho's federal donors are 65+**, and more than half of its state donors — against a third of the roll, with the under-30 share reduced to 1.4% and 2.6% respectively. The donor class is the grayest slice of the electorate in blue and red states alike.

Note the consistent gap between the panels: in **all three states, federal money is older money**. New York's 65+ donor share falls from 47.9% federal to 38.4% state, Idaho's from 64.7% to 51.1%, and Washington's state-level giving reaches Gen X and Millennials at roughly twice the rate federal giving does (Gen X 1.26× vs 0.98×, Millennial 0.68× vs 0.39×). The generational narrowing is sharpest in exactly the layer where the most money moves. It is also the layer most visible to national donors: state races are contested in every district, federal races in a handful, and the age gradient tracks that difference.

The skew is not a matching artifact. The obvious objection is that the match key (last name + first name + ZIP5, required to be *unique* on the roll) selects older, rarer-named, stable-address voters. Tested directly in the two states where the re-weighting was run (WA and NY), it does not: the probability a voter is uniquely matchable is **nearly flat across age** — NY **94.5–95.4%** across the four bands (0.9-pt spread, `diag_ny_match_bias.py`); WA **96.3–97.6%** across generations (1.4-pt spread). Inverse-propensity re-weighting therefore **does not move the distribution**: NY's 65+ donor share goes 47.9% → **47.9%** (0.0-pt shift), and every Washington generation multiplier is **unchanged to two decimal places in both panels**. The age skew is a property of *who gives*, not of *who the matcher can find*. (Idaho uses the identical conservative matcher but was not separately re-weighted; its 65+ donor shares are reported without that adjustment.) Appendix F carries the full validation.

Finding 2 — The donor class is whale-dominated

Money concentrates at the very top of the matched-donor distribution, in every state and in both money systems:

panel	matched donors	matched \$	top 1% → share of \$	top 10%	Gini
Federal					
Washington	172,998	\$420.3M	42.4%	74.8%	0.820

panel	matched donors	matched \$	top 1% → share of \$	top 10%	Gini
New York	307,841	\$1,196.1M	51.2%	81.4%	0.867
Idaho	27,196	\$49.6M	35.8%	70.4%	0.786
State					
Washington (PDC)	269,204	\$153.9M	43.8%	76.0%	0.827
New York (NYSBOE)	424,020	\$379.5M	48.5%	78.3%	0.846
Idaho (Sunshine)	27,250	\$15.9M	39.3%	70.8%	0.798

Estimator: donors are ranked by total matched dollars and split into 100 equal-count buckets ($NTILE(100)$); "top 1% / 10%" is the top 1 / 10 buckets' dollars ÷ all matched dollars. Appendix C gives the full definition and the reason equal-count buckets are used; Appendix E gives the bootstrap confidence intervals.

Two readings. First, the ordering **New York > Washington > Idaho** holds in *both* panels, and matches the statewide all-donor figures — so it is a property of these states' donor economies, not of who the matcher can find or which money system is examined.

Second, the state-versus-federal comparison does **not** run one way. State money is slightly *more* concentrated in Washington (43.8% vs 42.4%) and Idaho (39.3% vs 35.8%), but *less* so in New York (48.5% vs 51.2%) — even though state contributions are capped far lower than federal ones in all three. Whatever caps do, they do not produce a consistent ordering between the two layers. Appendix G takes that up directly.

A geographic corollary everywhere — and the sharpest panel difference in the paper. In WA, **63.4%** of federal matched-donor dollars come from two Seattle-metro ZIP3s (981xx 37.3% + 980xx 26.1%); the state panel is broader at 55.1% (32.3% + 22.8%).

New York is the extreme case. **Manhattan (New York County) supplies 50.3%** of the state's *federal* matched-donor dollars — top three ZIP3s **63.4%** — but only **20.6%** of its *state* dollars, where the leaders are suburban **Nassau (15.1%)** and **Suffolk (11.1%)** and the top three ZIP3s take just **38.1%**. New York's federal money is a Manhattan phenomenon; its state money is a Long Island one. Any account of "the New York donor class" that rests on the federal file alone is describing one island.

Idaho shows single-metro dominance from a state with no large city: **Ada County (Boise) supplies 49.2%** of state matched-donor dollars from 10,037 donors, and 36.4% of federal dollars from 10,338. Idaho's federal money is the one case where a second center appears: resort-county **Blaine (Sun Valley) supplies 11.7%** of federal dollars from 1,097 donors — 4.0% of Idaho's federal donors carrying nearly three times their weight.

(Idaho is the least concentrated of the three in both panels, and its statutory per-election caps do visibly bind — 6,797 itemized state gifts land on exactly \$1,000, against 448 at \$750. The caps are nonetheless not what makes Idaho less concentrated: Idaho is the least top-heavy inside the *federal* layer too, under the same federal limits as Washington and New York, and its far more tightly capped state layer is *more* top-heavy than its federal one, not less. Appendix G runs the test; Appendix D carries the statutes and the literature.)

At the statewide level (all itemized donors, not only those matched to a voter), the same top-heaviness appears across all four states loaded — the top 1% of donors supply **39.3%** of federal dollars in WA, **47.5%** in NY, **41.7%** in TX, and **36.0%** in ID (`cross_state_fec_money.py`). Texas enters here and nowhere else in this paper: no Texas voter file has been obtained, so Texas money can be described in aggregate but cannot be matched to individual voters. The "small-dollar democratization" narrative coexists with a money system whose itemized dollars are dominated by a thin top stratum and a single metro.

Finding 3 — The donor class is partisan-skewed toward Democrats (New York and Idaho)

This is the cut Washington cannot supply. Using each donor's **own** NY party enrollment (100% present), the donor class over-represents registered Democrats and **nearly excludes the unaffiliated** — in both money systems:

party	registration	federal donor share	skew	federal \$ share	state donor share	skew	state \$ share
-------	--------------	---------------------	------	------------------	-------------------	------	----------------

party	registration	federal donor share	skew	federal \$ share	state donor share	skew	state \$ share
DEM	47.8%	62.8%	+15.0	71.0%	56.7%	+8.9	55.0%
REP	22.3%	21.4%	-0.9	16.5%	25.4%	+3.2	27.3%
NOPARTY (blank)	25.5%	12.5%	-13.0	10.6%	13.5%	-12.0	14.1%
OTHER (minor)	4.5%	3.3%	-1.2	1.9%	4.4%	-0.1	3.6%

Registered Democrats are +15 points over their share of the roll in federal money and supply **71% of federal matched dollars**; Republicans give roughly in proportion. The constant across both panels is the **exclusion of the unaffiliated**: NY's "blank" enrollees are a quarter of all registrants and only an eighth of federal donors, barely better at a seventh of state donors.

The Democratic tilt itself is substantially a *federal* phenomenon. It halves in state money (+15.0 → +8.9), where Republicans move from slightly under-represented to slightly over (-0.9 → +3.2). New York's donor class leans Democratic most sharply where the money is nationalized; its state-level donor pool is closer to — though still not — the electorate. The unaffiliated bloc is shut out either way.

Where the money goes — crossover. Recipient party comes from the bulk FEC committee and candidate masters on the federal panel (`backfill_ny_committee_party.py` , 87.8% of matched donors resolved) and, on the state panel, from `backfill_ny_recipient_party.py` — NYSBOE publishes no party on the filer, so party is reconstructed from explicit party words in committee names, the finance rows that already carry a party, and the committee→candidate→roster chain. That reaches **37.7%**, the thinnest coverage in this paper; the state row below should be read as indicative, and the reasoning behind it is in Appendix C.

own party	federal: donors	→ D	→ R	state: donors	→ D	→ R
DEM	174,156	94.4%	3.9%	89,010	88.3%	8.9%
REP	57,330	14.2%	82.6%	48,708	12.2%	84.7%
NOPARTY	30,568	65.6%	31.0%	15,878	54.8%	41.8%
OTHER	8,268	40.8%	57.0%	6,114	36.9%	58.8%

Two patterns, both present in both panels: registered **Republicans fund Democrats at ~3.6× the rate Democrats fund Republicans** federally (14.2% vs 3.9%) — a deep-blue donor ecosystem; and the unaffiliated bloc, invisible to registration-based analysis, **leans Democratic in its actual giving** (65.6% → D federally, 54.8% state), so NY's independents are not centrist by behavior. Both patterns attenuate in state money, where partisans of both stripes are more loyal to their own side and the unaffiliated split closer to even.

A stability check on the thin state coverage: extending the backfill lifted resolution from 25.9% to 37.7% of matched state donors and moved the reported rates by 1.0 point (DEM→D), 4.5 (REP→R) and -2.1 (NOPARTY→D). The direction of every cell survived; the magnitudes shifted enough that the state column should be treated as approximate.

Idaho — the same skew, in the reddest state, in both money systems. The decisive test of whether the Democratic tilt of the donor class is real or just a blue-state artifact is to run it in a state where Republicans hold a 5:1 registration edge. Using each donor's own Idaho affiliation, in each panel separately:

party	registration	federal donor share	skew	state donor share	skew
REP	62.9%	66.7%	+3.8	66.5%	+3.6
DEM	11.8%	19.9%	+8.0	20.9%	+9.1
UNAFF (unaffiliated)	23.9%	12.7%	-11.2	12.0%	-11.8
OTHER (minor)	1.4%	0.8%	-0.6	0.6%	-0.8

Dollar shares track the same way — Republicans supply 68.1% of federal and 71.1% of state matched dollars, Democrats 21.1% and 20.6%, the unaffiliated 10.3% and 8.0% (federal \$49.6M across 27,196 donors; state \$15.9M across 27,250).

Republicans supply the plurality of Idaho's money, as a 63%-Republican state must — but relative to their numbers the **most over-represented donors are registered Democrats** (+8.0 federal, +9.1 state, well over half again their share of the roll), and the unaffiliated quarter is the most *under*-represented in both (−11.2, −11.8). The same directional finding as New York, from the opposite end of the spectrum, and it does not depend on which money system is examined: the donor class leans Democratic-of-the-electorate and runs against the unaffiliated, whether the electorate around it is blue or red and whether the money is federal or state.

Crossover (Idaho). The two panels differ sharply in how well recipient party can be resolved, and this is where the federal panel earns its place as the spine. Idaho Sunshine carries no party on the recipient, so for the state panel recipient party must be reconstructed from the Secretary of State candidate roster plus party/committee name patterns (`backfill_id_recipient_party.py`), resolving only ~52% of matched donors. The federal panel needs no reconstruction — the FEC committee and candidate masters carry party directly — and resolves **86.7%**, comparable to New York's 79%.

own party	federal: donors	→ D	→ R	state: donors	→ D	→ R
DEM	5,037	96.5%	2.8%	3,365	93.5%	4.0%
REP	15,579	17.6%	81.5%	9,420	18.8%	79.3%
UNAFF	2,819	74.0%	24.5%	1,303	72.8%	24.5%
OTHER	139	19.4%	79.1%	59	18.6%	81.4%

The two direction-safe patterns from New York replicate in both panels: registered **Democrats are near-monolithic donors** (96.5% → D federal, 93.5% state, against NY's 94.4%), and **unaffiliated donors lean Democratic** in their actual giving (roughly 3:1 in Idaho, 2:1 in New York) — independents are not centrist by behavior in either state.

The Republican→Democratic rate is the one figure that needs a caveat, and only for the state panel: there the unresolved recipient pool (local Republican candidates and R-aligned PACs absent from the roster) skews Republican, so Republican donors' Republican-side giving is disproportionately the untraced part, making 18.8% an **upper bound**. On the federal panel, with 86.7% resolved from authoritative party labels, the 17.6% figure carries no such hedge — and it lands within a point of the state panel's upper bound, which suggests the state reconstruction was not badly biased after all.

In-state vs out-of-state, by party. Money flowing *into* NY's federal races (`fec_inflow.duckdb` , all-state donors → NY candidates; `diag_ny_donor_extras.py`) is **44.8% out-of-state for both parties** — nationalization is party-symmetric at the aggregate, consistent with the cross-state finding that out-of-state share is uniform across competitiveness (`cross-state-fec-money.md` §G). The one asymmetry is by office: NY's **Senate Democrats draw 54.1% of their money from out-of-state** (Schumer/Gillibrand as national magnets) versus ~43–45% for House candidates of both parties. So the donor class is partisan-skewed in *who it is* (§3 above) but not in *how far its money travels* — except at the marquee Senate tier.

The skew holds in every kind of district. Mapping matched donors to their congressional district's competitiveness (`diag_ny_electorate_extras.py`), the donor pool's Democratic share **exceeds the registered Democratic share in every band** — Tossup 57.7% donor vs 40% registrant, Solid 71.6% vs 56% — so the donor class is more Democratic than the electorate not just statewide but locally, regardless of how contested the seat is. And **two-thirds of donors (205K of 308K) live in Solid districts** (mostly Solid-D Manhattan): the money originates in safe seats, consistent with the cross-state finding that safe seats supply most of it. Idaho shows the same safe-seat origin from the red side: the great bulk of matched Idaho donors sit in Solid-R legislative districts (where the donor pool is 71% Republican), while the handful of more competitive districts carry a far more balanced donor mix (~46% Republican / ~34% Democratic) (`diag_id_electorate_extras.py`) — the money, in both states, originates overwhelmingly in the seats that are not in doubt.

Finding 4 — Financial voice and electoral voice stack on the same people

In both states, the people who give are the people who reliably vote, and it holds in both money systems — this is the finding least sensitive to the panel split. In Washington, federal matched donors are **85.4% super-voters versus 51.4%** of non-donors (mean turnout propensity 0.966 vs 0.754); the state panel is nearly identical at **84.9% versus 50.8%** (0.950 vs 0.751). In New York, federal matched donors voted in **3.03 of the last four federal generals on average versus 1.78** for non-donors, and **73.0% are super-voters (≥3 of 4) versus 37.1%**; the state panel lands within a tenth of a point of that on every measure (3.02 vs 1.77; **73.1% versus 36.8%**), despite drawing on 116,000 more people. The same individuals concentrate *both* forms of influence rather than one offsetting the other. (Association only — donors are pre-selected for engagement, so reverse causation is equally plausible; the benign "donating as a gateway to participation" reading is fully live, and is treated as objection 3 in Appendix A.)

A closed-primary corollary at the nominating stage, in both party-of-record states. NY's **closed** primaries restrict each party's primary to its enrollees, so the **25.3% enrolled "blank" are excluded by law** ($\approx 0.1\text{--}0.6\%$ primary participation), and in blue NY the Democratic primary is frequently the decisive contest (2021 odd-year DEM 16.9% vs REP 5.0%). Idaho is the mirror image: its **closed Republican primary is the decisive contest** in nearly every seat (80–86% of all primary ballots are Republican; see [who-decides-idaho.md](#)), and the unaffiliated quarter is excluded in practice — **6.6% primary participation vs 83% in the November general**. The population that nominates is small and party-gated in both states — and, per Finding 3, funded by a donor class narrower and more skewed still.

What this paper does not claim, and limits

- **The match is a proxy, and a floor.** Voter $\leftrightarrow$ donor identity rests on (last, first, ZIP5) uniqueness, not a shared identifier. It is conservative by design (ambiguous keys are dropped, not guessed), so the matched set is a **floor**, not a census, of the donor population. Hand rating of 150 matched records puts apparent precision at **$\approx 90\%$** (Appendix F).
- **Itemized giving only, and two panels never pooled.** Sub-\$200 unitemized giving is invisible, so the *small-dollar* end is undercounted — which, if anything, **understates** the concentration in Findings 1–2. Federal and state money are reported as separate panels (Appendix C); dollar magnitudes are comparable within a panel, not across them, and New York has no state panel to compare.
- **Composition, not rates.** All three matches use the current roll, so turnout *rates* for older cycles are biased by survivorship. Share-of-population figures, which need no denominator, carry the findings. Idaho's age is a current-roll integer, so its bands are current-age, not election-time.
- **Recipient party is partial, and least complete on the state panels.** It resolves for **87.8%** of NY federal and **86.7%** of ID federal matched donors, where the FEC masters carry party outright, but only $\sim 52\%$ of ID state and **37.7%** of NY state donors, where it must be reconstructed from committee names and candidate rosters. Both state crossover columns are approximate, and Idaho's majority-party rate specifically is an upper bound. Own-party and age cuts use the 100%-present party of record throughout and are unaffected.
- **No policy-influence claim.** This paper measures who gives and who votes. It does not measure whether money changes votes, wins elections, or moves policy, and the giving $\leftrightarrow$ turnout relationship in Finding 4 is reported as association only.
- **Contribution limits are not the mechanism** behind Idaho's lower concentration (Appendix G), and no causal claim is made in either direction about what caps do to a donor pool.

Appendix A states each objection in full, with the bound on it.

What it means

Across three states that differ in size, partisanship, and election administration — and across two separately regulated money systems within them — the population that finances campaigns is the same kind of population: **old, top-heavy, geographically concentrated, and — where party is observable — skewed toward Democrats relative to the electorate and away from the unaffiliated**. That every finding survives the federal/state split is the strongest evidence here that it describes donors rather than a quirk of one disclosure regime: the two systems cap contributions differently, are solicited by different campaigns, and are reported to different agencies, yet produce the same donor class. New York and Idaho's party-of-record turns the Washington finding from "the donor class is demographically unrepresentative" into the sharper, falsifiable claim that it is *also* partisan-unrepresentative in a specific direction — and, critically, in the **same** direction in a deep-blue and a deep-red state, so the Democratic tilt is a property of who donates, not of a state's majority party. Combined with the turnout and safe-seat papers, the picture is a series of narrowing filters between the registered population and actual influence — who votes, who votes in the decisive primary, and who pays — each one older and less representative than the last. This is the evidentiary core of the electoral-health series' "donor class $\neq$ electorate" finding, now resolved by party across the spectrum.

Appendices

Appendix A — The objections, in full

The strongest objections to this paper are about the *match*, not the findings, and three of the four are testable. Each is stated at full strength, then bounded.

1. The matcher selects old people, so the age skew is manufactured. The match key requires a (last name, first name, ZIP5) triple that is *unique* on the roll. Rare names, stable addresses, and long tenure all correlate with age, so the objection is that Finding 1 measures matchability rather than giving. This is the one objection that can be tested head-on, and it fails: the probability that a voter is uniquely matchable is nearly flat across age (NY 94.5–95.4% across four bands; WA 96.3–97.6% across five generations), so inverse-propensity re-weighting does not move the distribution — NY's 65+ donor share is unchanged at 47.9%, and every Washington generation multiplier is unchanged to two decimal places in both panels. A selection mechanism that flat cannot produce a 2.5× senior over-representation. Full tables in Appendix F.

2. Household false-merges inflate individual donors. Because the key is surname plus ZIP, a married couple sharing both can collapse into one matched voter, attributing a spouse's giving to their partner. Hand rating of a 150-record sample (2026-07-10, two review rounds) found this to be the dominant error mode, at ≈90% apparent precision (15/150 flagged on the second, more thorough pass; 9 on the first). Its effect on the findings is small by construction: the mis-attributed partner shares household, ZIP, and typically similar age, so the age, geography, and concentration cuts move very little. Some flags were unverifiable for missing donor detail, so true precision may be slightly higher.

3. Donors vote more because donating is a gateway to participation, not because the same elite holds both forms of voice. Finding 4 is an association, and the causal arrow is genuinely ambiguous: donors are pre-selected for engagement, and a first contribution plausibly *increases* subsequent turnout. Nothing here distinguishes the two readings, and the benign one is fully live. What survives either reading is the descriptive point, which is all Finding 4 claims: the two forms of influence sit on the same people rather than on complementary populations.

4. Idaho's crossover rate overstates Republican defection. Recipient party in Idaho is reconstructed rather than published, and resolves only ~52% of matched donors. The unresolved pool — local Republican candidates and R-aligned PACs absent from the Secretary of State roster — skews Republican, so Republican donors' Republican-side giving is disproportionately the untraced part. The 18.8% Republican→Democratic figure is therefore an upper bound, and no cross-state claim is made about majority-party crossover. The two patterns that *are* claimed (near-monolithic Democratic loyalty; unaffiliated donors leaning Democratic) are robust to the unresolved pool because it cannot plausibly reverse them.

5. Itemization hides the small-dollar end, so concentration is overstated. Sub-\$200 giving is unitemized and therefore invisible here. This objection runs the wrong way for the paper's conclusion: adding the missing small-dollar mass to the denominator would *lower* every top-1% share, but it would also add the population the findings describe as absent. As reported, the concentration figures describe the *itemized* universe, and the direction of the bias is toward understating how thin the top stratum is relative to all givers.

6. State contribution caps, not donor behavior, explain Idaho's flatter distribution. This was the paper's own earlier explanation, and Appendix G shows it does not survive testing. It is retained here as an objection because it is the intuitive reading.

Appendix B — Data access and privacy

- **Washington.** The standard statewide **VRDB extract**, the single public extract the Secretary of State publishes. By statute the public file carries name, address, political jurisdiction, gender, **year of birth**, voting record, registration date, and registration number, and no other registration information is available for public inspection (RCW 29A.08.710) — the statutory reason this series uses year of birth rather than full date of birth. Use is restricted to elections and political purposes and **may not be commercial** (RCW 29A.08.720); the file is **not redistributable**, with penalties at RCW 29A.08.740. Washington further strengthened voter-data protections in 2026 (SB 5892 / Ch. 213, Laws of 2026, eff. Mar. 25, 2026).
- **New York.** The NYSVOTER statewide file, obtained by public-records request under New York's Freedom of Information Law (N.Y. Pub. Off. Law art. 6), subject to the elections-purpose certification the State Board of Elections requires. NY publishes **individual party enrollment** and **full date of birth**; only aggregate cohort counts are released here.
- **Idaho.** The statewide voter file from the Secretary of State. Idaho Code **§ 74-120** releases registrant **age** while withholding date of birth, driver's-license number, and (for the list product) address detail — the statutory reason Idaho age is a current-roll integer rather than an election-time DOB, and the reason Idaho's age bands are read against the current roll.

- **Contribution data are public records in every layer used.** FEC bulk individual contribution files; Idaho Sunshine state contribution filings; Washington PDC filings. No access restriction applies to them, and none are voter-file derived.
- **What is released.** Aggregate counts and shares only, with cell sizes in the thousands to millions. No individual-level records, names, or addresses appear in this paper, in the verification scripts' output, or in the repository. The one artifact that contains individual rows — the 150-record hand-rating sample used in Appendix F — was written to a scratch directory and is not committed. Only citations and code are published, not data. Full provenance and access dates: [Data Sources & Reproducibility](#).

Appendix C — Methods

- **Source and unit.** Each state's registration roll is joined to itemized individual contributions. The unit of analysis is a **matched voter-donor**: one row per registered voter for whom a unique contribution identity could be established.
- **Panels, and why the money systems are never pooled.** A state's contribution table can hold more than one money system, distinguished by the source prefix on each contribution's identifier: Washington carries federal FEC rows (\$646.2M) alongside state PDC filings (\$394.6M), and Idaho carries FEC rows (\$76.2M) alongside Sunshine state filings (\$53.3M). The match is therefore run **once per source**, writing to a separate panel table each time, and no figure in this paper mixes them. This matters quantitatively: pooling lets one person's federal and state giving stack into a single donor total while a one-system donor's does not, which mechanically raises measured concentration — Washington reads top-1% 46.6% pooled against 42.4% federal and 43.8% state. It also matters conceptually, since the two systems are capped and administered differently (Appendix G).
- **The New York state panel.** NYSBOE publishes contributions as a transaction-level feed (`data.ny.gov 4j2b-6a2j`, 12.6M rows back to 1999) carrying contributor last name, first name, and ZIP — everything the match key needs. Earlier drafts reported that New York had no state panel; that was a tooling gap, not a data gap. The repo's NY adapter read the same feed but kept only roll-up columns for `candidate_finance` and discarded contributor identity. `scripts/load_ny_contributions.py` now loads the per-contribution rows. Scope: individual contributors (`CNTRBR_TYPE_DESC = 'Individual'`) on Schedule A (monetary receipts, the direct analog of the FEC itemized individual file), cycles 2018–2026 to align with the WA PDC window — 3,954,090 contributions totalling \$880.3M, of which \$379.5M matches to a registered New York voter. Out-of-state donors are retained, as they are in the WA and ID state panels; the voter-roll join drops them.
- **The match key and its uniqueness rule.** Matching proceeds in tiers, the first being `STRICT_ZIP5_FULL` — (last name, full first name, ZIP5) — followed by first-initial variants. Every tier carries the same guard: the key must resolve to **exactly one** voter on the roll and one donor identity, or the record is dropped rather than guessed. Adding the full-first-name tier raised Washington's matched count from 320K to 382K (+19%) and it is now the dominant tier. Because ambiguity is dropped, the matched set is a floor.
- **Concentration estimator.** Donors are ranked by total matched dollars and split into 100 **equal-count** buckets (`NTILE(100)`) over donors with `total_donated > 0`; the top-1% and top-10% figures are the top 1 and top 10 buckets' dollars divided by all matched dollars. Equal-count buckets are used deliberately: capped and round-number giving produces heavy ties, and an earlier draft using `PERCENT_RANK` drifted from an exact decile at small N, reading Idaho's top-10% as 69.0% rather than 70.8%. Gini is computed on the same donor totals by the rank-weighted formula. Appendix G reuses this identical estimator so its layer comparisons are commensurable with Finding 2.
- **Reconstructing recipient party on the state panels.** Neither NYSBOE nor Idaho Sunshine publishes the recipient's party, so the crossover cut needs it inferred. `backfill_ny_recipient_party.py` works in four uniqueness-guarded tiers — an explicit party word in the committee name (dropped if a name claims both), a party already present on the finance row, a committee→candidate→roster chain, and containment of a roster candidate's full name — reaching 37.7% of matched state donors. A fifth, bare-surname tier of the kind Idaho can use was built and **rejected**: Idaho's recipient strings are "LAST, FIRST" candidate names, whereas New York's are free-text committee names, so searching for a surname inside a phrase misfires — it read "FRIENDS OF DAVID KNAPP" as Republican via the surname *David* and "SARATOGA COUNTY GREEN PARTY" as Republican via *Green*. It would have added roughly \$67M of apparently-resolved money at the cost of silent misassignment. Corporate, labor and trade PACs are left unresolved by design in both states: they are genuinely non-partisan recipients and a large share of state money, and giving them a party would manufacture a crossover result rather than measure one.
- **Match-bias re-weighting.** $P(\text{uniquely matchable})$ is estimated per age band or generation directly on the roll, and donor shares are re-weighted by its inverse. This is the test in Finding 1 and Appendix F; it was run for WA and NY, not for ID.
- **Age conventions differ by state, by statute.** WA and NY supply birth date (WA year-only, per RCW 29A.08.710), so

ages are election-time. Idaho supplies a current-roll integer age, so Idaho bands are current-age and are compared against the current roll, not an election-time cohort.

- **Rates versus shares.** Turnout *rates* computed from a current roll are inflated by survivorship wherever the roll has shrunk — acutely in Idaho, whose 2026 roll (1.03M) is smaller than the 1.18M registered at the 2024 election. Rate cuts are therefore not reported for Idaho, and all headline figures in this paper are denominator-free composition shares.
- **Reproduction.** `scripts/verify_donor_class.py` re-derives Findings 1–4 for all three states from scratch SQL, without importing the analysis code; `scripts/diag_contribution_limits.py` produces Appendix G. Neither imports the other. The verifier explicitly does not cover the crossover tables, the re-weighting, or the hand-rated sample — those are separate scripts and, in the last case, a human step.

Appendix D — Related work

That the donor class is small, wealthy, and unrepresentative is well established; the contribution here is the *voter-file-matched* view — linking individual donors to the registration roll across three states to show the donor class's age, concentration, partisan tilt, and turnout overlap on the same records. It sits in these literatures:

- **The shape of the donor class.** Bonica, "Mapping the Ideological Marketplace," *American Journal of Political Science* 58(2) (2014): 367–386, doi:10.1111/ajps.12062, and the DIME database; Schlozman, Verba & Brady, *The Unheavenly Chorus: Unequal Political Voice and the Broken Promise of American Democracy* (Princeton, 2012); Hill & Huber, "Representativeness and Motivations of the Contemporary Donorate: Results from Merged Survey and Administrative Records," *Political Behavior* 39(1) (2017): 3–29, doi:10.1007/s11109-016-9343-y — the donors-skew-old finding (Finding 1), and the closest methodological analog to the match used here. Bonica, McCarty, Poole & Rosenthal, "Why Hasn't Democracy Slowed Rising Inequality?" *Journal of Economic Perspectives* 27(3) (2013): 103–124, doi:10.1257/jep.27.3.103 — the concentration/whale result (Finding 2).
- **Unequal voice and its consequences.** Verba, Schlozman & Brady, *Voice and Equality: Civic Voluntarism in American Politics* (Harvard, 1995); Gilens, *Affluence and Influence: Economic Inequality and Political Power in America* (Princeton, 2012); Gilens & Page, "Testing Theories of American Politics: Elites, Interest Groups, and Average Citizens," *Perspectives on Politics* 12(3) (2014): 564–581, doi:10.1017/S1537592714001595 — the normative stakes of a participation-and-money elite.
- **Party and the donorate.** Grumbach & Sahn, "Race and Representation in Campaign Finance," *American Political Science Review* 114(1) (2020): 206–221, doi:10.1017/S0003055419000637; Grumbach, Sahn & Staszak, "Gender, Race, and Intersectionality in Campaign Finance," *Political Behavior* 44 (2022): 319–340, doi:10.1007/s11109-020-09619-0 — donor-pool composition by party and group, the frame for the Democratic tilt (Finding 3, New York and Idaho).
- **Giving and voting as stacked participation.** Verba, Schlozman & Brady (1995) again, on the co-occurrence of participatory acts; the giving→turnout overlap (Finding 4) is the individual-record instance, framed strictly as association.
- **Voter-file / individual-level method.** Ansolabehere & Hersh, "Validation: What Big Data Reveal About Survey Misreporting and the Real Electorate," *Political Analysis* 20(4) (2012): 437–459, doi:10.1093/pan/mps023; Hersh, *Hacking the Electorate: How Campaigns Perceive Voters* (Cambridge, 2015). On match bias and the current-roll caveat, the hand-rated ≈90% match precision and the inverse-propensity re-weighting in Appendix F address the older/stable-address/uncommon-name skew directly.
- **Contribution limits and the shape of the donor pool.** The statutory regimes tested in Appendix G: Idaho Code § 67-6610A (individual contributions capped at \$1,000 per election to legislative, judicial, and local candidates and \$5,000 per election to statewide candidates, with primary and general counted separately and self-funding exempt; 2026 S.B. 1422 proposed raising these to \$1,500 and \$6,000 but was retained on the calendar); RCW 42.17A.405, recodified as RCW 29B.40.020 effective Jan. 1, 2026 (Washington's caps, indexed and administered by the Public Disclosure Commission); N.Y. Elec. Law § 14-114 (New York's comparatively high caps); Tex. Elec. Code § 253.094 (Texas bars corporate and labor-organization contributions but imposes **no dollar limit** on individual gifts to non-judicial state candidates; the Judicial Campaign Fairness Act, §§ 253.151–253.176, is the exception); 52 U.S.C. § 30116 (the federal per-election individual limit, \$3,500 for the 2025–26 cycle, indexed for inflation); and 52 U.S.C. § 30118 (federal prohibition on corporate contributions, which the Idaho and Texas state systems do not share). On the constitutional architecture that leaves per-gift caps standing while removing any ceiling on a donor's total giving: *Buckley v. Valeo*, 424 U.S. 1 (1976) (contribution limits upheld, expenditure limits struck), and *McCutcheon v. FEC*, 572 U.S. 185 (2014) (invalidating the biennial aggregate limit). On the empirical consequence — that limits redistribute large-donor influence across vehicles rather than removing it: Barber, "Ideological Donors, Contribution Limits, and the Polarization

of American Legislatures," *Journal of Politics* 78(1) (2016): 296–310, doi:10.1086/683453; La Raja & Schaffner, *Campaign Finance and Political Polarization: When Purists Prevail* (Michigan, 2015). Appendix G's result is consistent with that literature.

Appendix E — Full distribution tables

Matched-donor concentration, with bootstrap intervals. Both Washington panels were bootstrapped at $B=1,000$ resamples (`diag_donor_concentration_bootstrap.py`). The intervals are tight enough that the concentration finding does not rest on a handful of top donors, and the two panels' intervals overlap — the state layer's higher point estimate is not separable from the federal layer's at this precision:

Washington, matched donors	federal panel	95% interval	state panel	95% interval
top 1% → share of matched \$	42.4%	[40.2–44.9]	43.8%	[39.6–48.9]
top 10% → share of matched \$	74.8%	[73.7–76.0]	76.0%	[74.2–78.2]
Gini	0.820	[0.812–0.828]	0.827	[0.814–0.843]

Statewide (all itemized donors, not only matched), four states. From `cross_state_fec_money.py` over each state's FEC individual layer, donor-residence filtered. This is the one table in the paper where Texas appears, since aggregate donor figures need no voter file:

statewide concentration	WA	NY	TX	ID
top 1% → share of \$	39.3%	47.5%	41.7%	36.0%
top 10% → share of \$	72.3%	78.7%	74.5%	69.2%
Gini	0.800	0.848	0.818	0.775

Idaho is the least concentrated of the four on every measure, and the ordering matches the matched federal panel ($NY > WA > ID$) — so the panel result is not an artifact of who the matcher can find. It also cannot be attributed to state contribution caps, since this table is entirely federal money under identical federal limits; Appendix G develops that point.

Candidate money versus total flow. Concentration is a property of the *uncapped* vehicles more than of candidate committees. Money reaching candidates (each gift bounded by the per-election limit) runs top-1% $\approx$ **16–18%** and Gini $\approx$ **0.69**, against **39–48%** and **0.80–0.85** for total outflow including party committees, joint fundraising committees, and PACs (`cross-state-fec-money.md` §I).

Geographic concentration of matched dollars, by panel.

state / panel	leading geography	share of matched \$	next
WA federal	ZIP3 981xx (Seattle)	37.3%	980xx 26.1% → two ZIP3s = 63.4%
WA state	ZIP3 981xx (Seattle)	32.3%	980xx 22.8% → two ZIP3s = 55.1%
NY federal	New York County (Manhattan)	50.3%	Westchester 12.4%, Kings 6.6%; top-3 ZIP3 = 63.4%
NY state	New York County (Manhattan)	20.6%	Nassau 15.1%, Suffolk 11.1%; top-3 ZIP3 = 38.1%
ID federal	Ada County (Boise), 10,338 donors	36.4%	Blaine 11.7%, Bonneville 11.2%
ID state	Ada County (Boise), 10,037 donors	49.2%	Kootenai 9.0%, Canyon 5.5%

Two of the three states put federal money in a tighter geographic box than state money, and in New York the gap is enormous — Manhattan's share of matched dollars falls by **30 points**, from half of the federal layer to a fifth of the state one, with suburban Nassau and Suffolk taking its place. The mechanism is plausible and the same in Washington: state

legislative races are contested in every district, so state money is raised everywhere, while federal money pools where national donors live. Idaho inverts it — Ada County's grip *loosens* from 49.2% of state dollars to 36.4% of federal — not because Idaho's federal money is broadly spread, but because it relocates to wealthy enclaves outside the capital, above all resort-county Blaine at 11.7% from 1,097 donors. Concentration is the constant; which geography does the concentrating depends on the money system.

Donor pool versus registration, by district competitiveness (NY). The Democratic share of the donor pool exceeds the Democratic share of registrants in every band, and two-thirds of all matched donors (205K of 308K) live in Solid districts:

band	donor pool, D share	registrants, D share
Tossup	57.7%	~40%
Solid	71.6%	~56%

Idaho shows the mirror image from the red side: the bulk of matched donors sit in Solid-R legislative districts, where the donor pool is ~71% Republican, while the few more competitive districts carry a far more balanced mix (~46% R / ~34% D) (`diag_id_electorate_extras.py`).

Giving and turnout, side by side.

	donors	non-donors
WA super-voter share, federal panel	85.4%	51.4%
WA mean turnout propensity, federal panel	0.966	0.754
WA super-voter share, state panel	84.9%	50.8%
WA mean turnout propensity, state panel	0.950	0.751
NY generals voted, of last 4, federal panel	3.03	1.78
NY super-voter share (≥3 of 4), federal panel	73.0%	37.1%
NY generals voted, of last 4, state panel	3.02	1.77
NY super-voter share (≥3 of 4), state panel	73.1%	36.8%

The give↔vote overlap is the finding least sensitive to which money system is examined: Washington's two panels differ by half a point, and New York's by a tenth of one — even though New York's state panel draws on 116,000 more people than its federal panel.

Appendix F — Match validation and robustness

Is matchability age-dependent? This is objection 1, tested directly. P(a voter is uniquely matchable) is computed on the roll itself, then donor shares are re-weighted by its inverse:

state	spread of P(matchable)	re-weighted result
NY, four age bands	94.5%–95.4% (0.9-pt spread)	65+ donor share 47.9% → 47.9%
WA, five generations	96.3%–97.6% (1.4-pt spread)	every multiplier unchanged to 2 d.p., both panels (federal Silent 2.48 → 2.48×, Gen Z 0.10 → 0.10×; state Silent 1.59 → 1.59×)

A selection gradient that flat cannot generate the observed senior over-representation. Two notes on the Washington row. First, P(uniquely matchable) is a property of the *roll* and the match key, not of which contributions are used, so the same propensities re-weight both panels; only the multipliers they act on are panel-specific. Second, the WA test uses a slightly narrower denominator than Finding 1 (active registrants carrying a ZIP), so its raw multipliers sit a few hundredths below Finding 1's — the quantity of interest is the raw-to-re-weighted *difference*, which is zero to two decimals either way. An earlier draft, computed on the pooled match and a stricter matchability definition, reported a 68.9–73.1% spread and reached the same null result (Silent 1.87 → 1.83×). Idaho was not separately re-weighted, so its 64.7% and 51.1% figures carry

no such adjustment.

Hand-rated match precision. A 150-record sample was rated by hand on 2026-07-10 in two rounds (`diag_match_validation_sample.py`): 9 records flagged on the first pass and **15 on the second, more thorough pass**, i.e. **≈90% apparent precision**. The dominant error mode is spousal or household false-merge (shared surname and ZIP). Several flags were unverifiable for want of donor detail, so true precision may be slightly higher. The sample contains individual-level rows and is deliberately not committed to the repository.

How much of the donor side is reachable at all? Ceiling analysis on the donor files (`diag_donor_match_ceiling.py`) — the share of donors whose key could in principle resolve to a unique voter, on the first-initial key and the full-name key:

	WA	NY	TX	ID
initial key	65%	69%	61%	86%
full-name key	72%	76%	68%	96%

Idaho's ceiling is far higher (smaller roll, fewer name collisions), which is worth keeping in view when comparing matched counts across states: the three matched sets are not equally complete samples of their donor populations.

One residual bias, and its direction. A donor who moved between giving and the extract date fails the ZIP match, and mobility skews young — deflating the young donor share. The raw age skew in Finding 1 is therefore an **upper bound**. What that leaves untouched is the specific mechanism objection 1 names (rare names make the old easier to match), which the flat $P(\text{matchable})$ above rules out.

Appendix G — Contribution limits and the top of the distribution

Earlier drafts of Finding 2 explained Idaho's lower matched top-1% share by asserting that state contribution limits "compress the very top of the state-money distribution." The explanation is intuitive, and it is wrong. This appendix tests it (`diag_contribution_limits.py`) and reports what replaces it.

The test does not require new data. Idaho residents appear in **two money systems at once**, under two different rule sets, so the regime can be varied while holding the donor population nearly fixed. Washington supplies the same pair.

G1 — the statutory regimes. Individual contributions to a candidate, per election:

layer	legislative	statewide	ceiling on a donor's <i>total</i> giving
Idaho state (Sunshine)	\$1,000	\$5,000	none
Washington state (PDC)	capped, indexed	capped, indexed	none
New York state (BOE)	capped, high	capped	none
Texas state (TEC)	no dollar limit	no dollar limit	none
Federal (FEC)	\$3,500 (2025–26)	\$3,500	none since 2014

Note the direction of the caps: Idaho's legislative limit is *lower* than the federal per-election limit, but **no** system caps a donor's total, the federal aggregate ceiling having been struck in *McCutcheon v. FEC*, 572 U.S. 185 (2014). Idaho and Texas also permit direct corporate and PAC contributions to candidates, which federal law forbids, and Idaho caps nothing at all on ballot-measure committees. Statutes in Appendix D.

G2 — the cap binds. Bunching on the round statutory value is unmistakable in Idaho's state filings:

gift amount	itemized gifts
\$750	448
\$900	190
\$999	94
\$1,000 (legislative cap)	6,797
\$1,001	1

gift amount	itemized gifts
\$1,100	18
\$5,000 (statewide cap)	734
\$5,001	0

A 15× spike on the cap value and a cliff immediately above it is what a binding constraint looks like. So the premise of the original explanation holds: the cap is real and donors hit it.

G3 — and yet capped layers are not less concentrated. Applying Finding 2's estimator to each layer separately (donor = name + ZIP5):

layer	rows	\$	donors	top 1%	top 10%	Gini
ID state (capped), all filers	216,700	\$53.3M	54,019	56.4%	81.6%	0.872
ID state (capped), persons only	181,539	\$23.9M	47,356	39.7%	71.2%	0.800
ID federal, all filers	770,765	\$76.2M	54,155	36.1%	69.2%	0.775
ID federal, persons only	770,128	\$76.2M	54,088	36.1%	69.2%	0.775
WA state (capped), all filers	2,816,398	\$348.3M	728,255	44.4%	75.5%	0.823
WA federal, persons only	5,578,905	\$645.6M	361,184	39.3%	72.3%	0.800

Two readings, both against the original explanation. First, in **both** states the capped *state* layer is **more** top-heavy than the uncapped-aggregate *federal* layer — Idaho 39.7% vs 36.1% among persons, and 56.4% vs 36.1% with committees included; Washington 44.4% vs 39.3%. Second, the persons-only Idaho state figure (39.7%, Gini 0.800) closely reproduces Finding 2's matched 39.3% / 0.798, which is an independent check that these layer definitions match the paper's.

G4 — what a cap would do, mechanically. Trimming every gift in Washington's federal layer to Idaho's caps isolates pure truncation, holding the donor population fixed:

per-gift cap	top 1%	top 10%	\$ retained
uncapped (actual)	39.3%	72.3%	\$646M
\$5,000 (ID statewide cap)	32.2%	68.7%	\$571M
\$3,500 (federal limit)	31.2%	67.7%	\$554M
\$1,000 (ID legislative cap)	26.4%	62.9%	\$454M

So truncation is *not* a small effect: an Idaho-style legislative cap applied mechanically would cut the top-1% share by 12.9 points. That sharpens the puzzle rather than resolving it. Idaho's actually-capped state layer sits at 39.7%, **13.3 points above** what pure truncation predicts.

What replaces the original explanation. The cap binds at the moment of the gift, and then its compression is undone downstream. Nothing limits a donor's *total*, so the same people reach the cap again and again across many recipients; and the tail displaces into vehicles the state does not cap — 41.9% of Idaho's state dollars sit in gifts above \$5,000, the largest single Sunshine contribution is \$1,245,000, and with committees included the state layer reaches 56.4%. This is the displacement the limits literature describes (Barber 2016; La Raja & Schaffner 2015), and it matches the rest of the series: Idaho's state legislative money is ~50% PAC-funded and its single largest filer is a ballot-measure committee (`cross-state-fec-money.md` K4, K5), while capped candidate-side inflow runs top-1% ≈ 16–18% against 39–48% for total outflow (\$!).

Independently, the state-versus-federal distinction could not have explained the gap it was invoked for. Idaho is the least concentrated of the four states **inside the federal layer too** (36.0% statewide, against WA 39.3% and NY 47.5%), under identical federal rules (Appendix E). Idaho's flatter distribution is a property of its small, retail donor base, not of its state caps.

Caveats. The Sunshine layer covers cycles 2022–2025 (odd years included) while the federal layers cover 2018–2026, so the windows are not identical. Separating people from organizations relies on a name heuristic, and the two state files differ:

Idaho Sunshine files people as "LAST, FIRST", so a comma test works, whereas Washington PDC files them as "LAST FIRST", so no reliable persons-only cut is available for the WA state layer and its all-filer row is the one to read. PDC's "SMALL CONTRIBUTIONS" unitemized pseudo-contributor is excluded from all Washington state cuts; left in, it keys as a single enormous donor and inflates every figure. The FEC files are persons by construction, which is why their two cuts nearly coincide — itself a check that the heuristic is not driving the result. No causal claim is made about what caps do to a donor pool; the claim is the narrower one that caps do not account for the cross-state differences in Finding 2.

End note — data, reproduction, and series

Reproduction, in dependency order:

Each match script builds ONE panel per run, selected with `--source` ; the two panels are written to `voter_donor_affiliation_fec` and `voter_donor_affiliation_state` and are never pooled (Appendix C).

```
# Washington - both panels:
python scripts/match_wa_voters_to_donors.py --source fec
python scripts/match_wa_voters_to_donors.py --source state
python scripts/diag_wa_individual_findings.py

# New York - both panels:
python scripts/load_ny_voters.py                # NYSVOTER FOIL -> ny_vrdb.duckdb
python scripts/diag_ny_turnout_party.py --rebuild # voter_participation table
python scripts/backfill_ny_committee_party.py    # bulk FEC committee/candidate party -> 79%
python scripts/load_ny_contributions.py         # NYSBOE per-contribution rows (state layer)
python scripts/backfill_ny_recipient_party.py    # state recipient party -> 37.7%
STATE=NY python scripts/match_ny_voters_to_donors.py --source fec
STATE=NY python scripts/match_ny_voters_to_donors.py --source state
STATE=NY python scripts/diag_ny_match_bias.py    # age-skew validation
STATE=NY python scripts/diag_ny_primary_participation.py
STATE=NY python scripts/diag_ny_donor_extras.py # geography, giving<->turnout, in/out-of-state

# Idaho - both panels (party of record + current-roll age):
python scripts/load_id_voters.py                # ID SoS voter file -> id_vrdb.duckdb
STATE=ID python scripts/backfill_id_recipient_party.py # state recipient party from SoS roster + patte
STATE=ID python scripts/match_id_voters_to_donors.py --source fec
STATE=ID python scripts/match_id_voters_to_donors.py --source state
python scripts/diag_id_electorate_extras.py      # donor-mix x LD competitiveness

# Cross-state dollar concentration, and the appendices:
python scripts/cross_state_fec_money.py         # Appendix E statewide table
python scripts/diag_donor_concentration_bootstrap.py # Appendix E intervals (both WA panels)
python scripts/diag_donor_match_ceiling.py      # Appendix F ceilings
python scripts/diag_contribution_limits.py      # Appendix G

# Independent re-derivation of Findings 1-4 across both panels
# (from-scratch SQL, imports no analysis code):
python scripts/verify_donor_class.py
```

All inputs are public records (FEC bulk files; Idaho Sunshine and Washington PDC filings; state voter files obtained under each state's lawful-use terms — NY NYSVOTER FOIL, WA VRDB, ID SoS). See [data-sources-and-reproducibility.md](#) for the full source ledger and the matcher/competitiveness method notes, and [publication-checklist.md](#) for the verification ledger of expected values.

This is Paper #3 of the electoral-health series: [who-decides-washington.md](#) (the gray off-year electorate), [who-decides-new-york.md](#) and [who-decides-idaho.md](#) (party-resolved electorates), [safe-seat-washington.md](#) (observed competitiveness), and [cross-state-fec-money.md](#) (the four-state money layer).