

The State of Electoral Health: WA / NY / TX / ID

A research prospectus built from precinct results, campaign finance, and the individual voter record

Draft outline — 2026-06-27. Derived from the democracy-insight gauntlet (*idea-gauntlet/RESULTS-democracy-insight-2026-06-27.md*): 10 research questions scored on 11 dimensions, V-Dem-anchored, against a pre-registered null.

Scope and integrity statement (read first)

This program studies the **electoral** dimension of democracy — participation, representation, political equality, and contestation. It is **silent** on rule of law, civil liberties, press freedom, and executive constraint, and it makes no claim about them.

Three commitments govern every section below:

1. **Pre-registered null.** The working hypothesis is that "Washington's (and the comparison states') electoral democracy is functioning." A finding moves against that null only when the evidence earns it.
2. **Accuracy gate.** Each finding contributes to any larger conclusion *only in proportion to its data sufficiency × inferential strength*. A striking number built on data the warehouse cannot actually carry is reported as a limit, not a result.
3. **No verdict.** This document assembles the *components* of a democratic-health assessment. It deliberately does **not** declare a "failure of democracy." The point is to be the instrument capable of detecting one honestly — not to reach for the headline.

Bottom line of the scan that produced this prospectus: of 10 questions, **nine scored "weak-against-null" and one "supports-null"; none reached "moderate" or "strong" against the null.** The accuracy-weighted failure signal was **22/100**. The honest reading of the present evidence is **measurable electoral stress that is real, largely known to the literature, reformable, and admits strong benign readings — not, on this data, systemic failure.** Critically, the findings that could most strengthen a failure case are precisely the ones the data is currently blind to (see *Boundary of Inference*).

Status update — 2026-06-28 (first analyses now executed). This began as a prospectus of analyses *to run*. Several have since been run on data already on disk, and the realized write-ups exist: - **Finding 1** → [who-decides-washington.md](#) (gray off-year electorate, from 27.1M VRDB vote records). - **Finding 2** → [safe-seat-washington.md](#) — now on **observed** margins, extended to a **complete four-state lower-chamber map** (WA 87.8 / NY 88.6 / TX 94.0 / ID 92.9% not close), and the observed counts **validate** the model's projection used below. (WA read 88.8% until the seat universe was rebuilt from certified statewide summaries; the old figure came from a results-table universe that silently dropped 24 King County House seats per cycle in 2016 and 2018.) - **Findings 4 & 5** → [cross-state-fec-money.md](#) §F — and the matcher-bias objection to Finding 5 has been **tested and rejected** (see that finding below); the donor match is **314,974** on the full-name-key specification adopted 2026-07-27 (it was 382K all-tier, and 320K before tier 0). - **The "non-constituent money" boundary item is no longer blocked** — NY/TX contributions were loaded and the recipient-anchored inflow built, so the cross-state nationalization test *did* run ([cross-state-fec-money.md](#) §E-I). None of this changes the **"stress, not failure"** verdict, but it moves the lead findings from *literature-borrowed* to *established in-data*, and resolves one of the three boundary blockers. The two that remain (individual party-of-record; party-resolved turnout/crossover) still gate a stronger verdict and await the NY/ID voter files.

Method

Each research question was scored 0–10 on eleven dimensions by an independent, web-searching reviewer instructed to be adversarial toward its own conclusions:

- **Insight layer (research quality):** data sufficiency, inferential strength, novelty (vs. the existing political-science literature), systemic significance, usefulness, robustness.
- **Diagnostic layer (anatomy of a failure claim):** function impairment (severity vs. the democratic ideal), trajectory (is it worsening?), entrenchment (does it resist self-correction?), counter-thesis strength (how good is the "this is actually healthy" reading — *reverse-scored*), convergence (does it triangulate with the other findings?).

Function impairment is anchored to the **V-Dem component indices** (electoral, liberal, participatory, deliberative,

egalitarian) so severity is legible and comparable across states and time.

Data provenance: ~5.1M precinct-result rows across 19 elections (WA) plus NY/TX/ID; ~8.6M individual contributions (FEC + state PDC, ~\$1.3B, 2018–2026) with employer/occupation/ZIP; ~9.2M voter-score rows; ~27.1M individual VRDB vote records; and the rare asset — **314,974 voters matched to their donations** at the person level, on the full-name-key specification adopted 2026-07-27; see §F of the cross-state money paper).

Findings — the publishable core (write these up first)

Ordered by insight composite. Each carries its single most defensible claim, the strongest objection to it, the V-Dem function and diagnostic read, and the first concrete analysis to run.

1. Who actually decides? The gray off-year electorate

Insight 62 · failure-contribution 22 · null: weak-against · V-Dem: participatory, egalitarian, electoral · trajectory: stable

- **Defensible claim.** The electorate that decides odd-year general elections in WA is roughly half the size of the presidential electorate and dramatically older. From 27.1M VRDB vote records (~100% birthdate coverage): voters 65+ were **~37–40%** of off-year ballots (2021/2023/2025) vs **28.5%** in 2024, while 18–29 were **7–8%** off-year vs **14.2%** presidential — a **~5:1** senior-to-youth ratio off-year vs ~2:1 presidential. Individually, 18–29 turnout collapses **58.4%** → **15.8%** (presidential → off-year) while 65+ falls only 88.3% → 61.3%.
- **Strongest objection.** This is *voluntary* differential participation, not disenfranchisement. WA has all-mail, postage-paid, automatic/same-day registration — the gray electorate reflects who *chose* to vote, and the gap is mechanically expected from low salience. The fix is purely institutional (move local races on-cycle to even Novembers), which is itself evidence the system is responsive, not failing.
- **First analysis — DONE** ([who-decides-washington.md](#)): Join `voting_history` to `voters.birthdate`; classify age-as-of-each-election into cohorts; produce a cycle-type × cohort table of both (a) within-cohort turnout and (b) cohort *share* of the actual electorate, for every general 2021–2025. Report composition shares (denominator-free) alongside rates to neutralize the registration-churn objection. **Stop short of any partisan-consequence claim** — that needs statewide party-of-record (see §2 of the follow-on).
- **Key literature.** Hajnal & Trounstein (off-cycle timing skews the electorate older/whiter); Anzia (*Timing and Turnout*); Lucero et al. 2025 (on-cycle consolidation cut California's over-55 local share from ~half to 28%).

2. Safe-seat democracy: the collapse of general-election contestation

Insight 52 · failure-contribution 34 (highest) · null: weak-against · V-Dem: electoral, participatory, deliberative · trajectory: worsening

- **Defensible claim.** The large majority of seats are non-competitive. On the model's own symmetric Cook-style bands, seats at ≥10-pt margin run **WA 90% (53/59), NY 86% (206/240), TX 81% (167/205), ID 92% (34/37)**, with genuine Tossups a small minority. When the general is foregone, the operative decision moves to lower-turnout primaries. This is a structural counting result that does **not** depend on any blocked/weak signal.
- **Strongest objection.** Those are *model-projected* 2026 margins, sensitive to the 10-pt threshold; and "non-competitive" conflates a closed-shop seat with one that is lopsided because voters *genuinely* lean that way (self-sorting). A Solid-D Seattle seat and a Solid-R rural-ID seat each *represent* their electorates — the null's "outcomes reflect voter choice."
- **First analysis — DONE** ([safe-seat-washington.md](#)): Ran **observed** margins, not the projection: per state, the most recent general for partisan legislative + congressional offices → margin per district → band counts. **The companion has since split this into two dimensions that must not be merged:** *candidate competition* (was the race close, on the top-two margin regardless of party) and *partisan availability* (did the ballot offer both a Democrat and a Republican). They overlap but are different questions — in WA 2024, 83.5% of seats were not close and 35.3% offered no D-v-R option, and of the fifteen same-party generals fourteen were also lopsided but **one was decided by six points**. Treating "no major-party choice" as automatically non-competitive, as this bullet originally did, is the conflation the rewrite removed. For WA, built the 2016–2024 seat-level trend (not-close share runs **78.9–88.1%** across the five cycles, dipping in the 2018 wave) and the primary-to-general turnout ratio in safe seats (~0.5) — the one place this data *adds* to Cook/Ballotpedia/Unite America by attaching real turnout (and the voter-donor join) to the "primary decides" claim.
- **Key literature.** Cook PVI (swing districts 164 in 1997 → ~72 post-2016); Ballotpedia Competitiveness Index (38% of

state-leg seats uncontested in 2024); Unite America (*The Primary Problem*).

3. Whale-dominated money behind a small-dollar facade

Insight 48 · failure-contribution 22 · null: weak-against · V-Dem: egalitarian, participatory · trajectory: indeterminate

- **Defensible claim.** WA money is broad by headcount but concentrated by dollar. For recipient-cycles with ≥ 100 distinct donors ($n=2,821$), the **median itemized gift is \$25** yet single gifts reach **\$2.5M**, with a per-recipient Gini ~ 0.61 . The data can quantify, per race and cycle, the sub-\$200 retail share vs. the whale layer.
- **Strongest objection.** The recipient key is `fec_candidate_id`, and the highest-volume "recipients" are conduits (ActBlue, JFCs), so a naive per-recipient Gini measures a *conduit's pass-through book*, not a candidate's race. A true per-race figure needs an earmark-attribution layer that does not yet exist and must be validated against double-counting. And a Gini that mixes \$25 and \$2.5M is mechanically high regardless of democratic health.
- **First analysis.** Build the conduit-attribution layer *first* (partition direct-to-candidate vs. conduit-routed; verify whether ActBlue/WinRed rows carry attributable earmark memos). Then compute, per ultimate-recipient $\times$ cycle: Gini, top-1%/top-10% dollar share, and sub-\$200 vs. $\geq \$200$ split — reported **with and without** conduit pass-throughs, since that toggle determines the answer.
- **Key literature.** Bonica & Rosenthal (wealth-elasticity of giving); Bouton & Cagé (*Small Campaign Donors*); Brennan Center (small-dollar grew, megadonors grew faster — "thin layer over a concentrated core").

4. Money and votes stack on the same people

Insight 47 · failure-contribution 22 · null: weak-against · V-Dem: participatory, egalitarian · trajectory: indeterminate · (most novel: 6)

- **Defensible claim (cross-sectional only).** Among the 314,974 matched voters, donors are a participation elite: **87.6% are super-voters vs 50.9%** of non-donors (average turnout propensity **0.967 vs 0.749**; donor-class verifier F4). Financial voice and electoral voice concentrate on the same individuals rather than offsetting.
- **Strongest objection.** The *causal/longitudinal* version ("giving *makes* people vote more; inequality *deepens*") is unsupportable here: donors are pre-selected for engagement (reverse causation is equally plausible), the match is biased toward older/stable-address/super-voting/uncommon-named people (inflating the gap), and `voter_donor_affiliation` collapses each voter to one row, so first-gift-then-vote sequencing can't be reconstructed on a shallow 2021–2026 history. The benign reading — donating as a *gateway* that broadens participation — is live.
- **First analysis — DONE** (`cross-state-fec-money.md` §F3): the cross-sectional benchmark — matched donors are **87.6% super-voters vs 50.9%** of non-donors (**1.72 $\times$**), mean turnout propensity **0.967 vs 0.749** — framed strictly as association, with the match-bias diagnostic from §F2 (the giving $\leftrightarrow$ voting overlap is real; the *causal/longitudinal* version remains out of reach).
- **Key literature.** Verba/Schlozman/Brady (*Voice and Equality*; *Unheavenly Chorus*) — money is the most income-skewed form of participation; the donor pool is a structural elite (the *constant* against which any worsening must be measured).

5. The donor class is not the electorate

Insight 42 · failure-contribution 22 · null: weak-against · V-Dem: egalitarian, participatory · trajectory: indeterminate

□ **Panel note (2026-07-26).** The WA figures in this finding were computed on a **pooled** voter $\leftrightarrow$ donor match, before it was discovered that WA's `individual_contributions` holds federal (FEC: , \$646.2M) and state (PDC: , \$394.6M) money and the matcher had no source filter. Pooling stacks one person's federal and state giving into a single donor total and inflates concentration. `donor-class-and-the-electorate.md` now reports two panels — **federal** 147,745 donors / top-1% **41.2%** / Gini **0.815**, **state** 217,114 / **43.5%** / **0.821** — which supersede the pooled numbers below. Every finding survives, and the age skew *strengthens* on the federal panel (Silent **2.67 $\times$** , Gen Z **0.04 $\times$**).

- **Defensible claim.** At the person level, donors are a narrow slice that does not mirror the electorate: **~ 3.5 – 6% of voters**, skewed old (pooled WA: Silent **1.96 $\times$** , Boomer **1.71 $\times$** over-represented; Gen Z **0.09 $\times$** , Millennial **0.54 $\times$** under-represented — $2.67 \times / 2.04 \times / 0.04 \times / 0.35 \times$ on the federal panel), overwhelmingly super-voters, geographically concentrated (**61.4% of WA donor dollars from two Seattle-metro ZIP3s**; 63.5% federal-only), and internally top-heavy (top 1% of matched donors supply **46.6%** of matched dollars pooled, **41.2%** federal; top 10% **79.3% / 74.2%**). RETIRED (\$154.0M, 23.8%) and NOT EMPLOYED (\$128.8M, 19.9%) are the two largest occupation blocs. (*These are contribution-level, not matched-donor, figures — occupation is on the gift, not the voter — and they are stated on the*

series' documented outflow basis: FEC rows, Washington-resident donors, \$646.2M. Until 2026-07-27 they read \$221.7M / 21.3% and \$147.4M / 14.1%, computed on the unfiltered pooled table — FEC plus state PDC plus non-resident donors, \$1,050.8M — i.e. the same two-money-system pooling this series corrected everywhere else, sitting inside a finding whose every other number is panel-scoped. On the federal basis both blocs are a LARGER share, so the point sharpens.) The voter→donor join makes this *person-level*, not ecological — the genuine value-add over the standard FEC-aggregate literature.

- **Strongest objection — now tested and rejected.** The objection was that the skew is a **matcher artifact**: the (last name, first initial, ZIP) unique-key over-selects the older, rarer-named, stable-address residents it then "finds" over-represented. **Tested directly** (cross-state §F2): the probability a voter is uniquely matchable on that key is **nearly flat across generations (69.1%–73.3%, a ~4-pt spread)**, so inverse-propensity re-weighting **barely moves** the over-representation ratios (Silent 1.96→1.91×, Gen Z 0.09→0.09× on the 314,974-voter match). The age skew is a real property of *who gives*, not of who the matcher can find. (One residual bias it can't observe — donors who moved between giving and now — runs the *same* direction, so the raw skew is an upper bound; the named name-commonness mechanism explains almost none of it.) The voter file still has **no income and no race**, so the literature's richer/whiter claims remain proxied, not tested.
- **First analysis — DONE** ([cross-state-fec-money.md](#) §F2): the matcher-bias inverse-propensity re-weighting above; skews reported raw *and* re-weighted; income/race labeled untestable. Concentration on the pooled 314,974 match: top-1% **46.6%**, top-10% **79.3%**, Gini 0.857; 61.4% of dollars from two Seattle ZIP3s. Superseded per the panel note above — federal panel top-1% **41.2%** [38.6–43.4], Gini **0.815** [0.806–0.822], two-ZIP3 share **63.5%**; state panel top-1% **43.5%** [38.7–48.9], Gini **0.821** [0.806–0.838]. (*This line previously read 42.4% [40.2–44.9] for the federal panel — the all-tier value, contradicting the 41.2% in the panel note directly above it. The bootstrap CIs are the per-panel re-runs in [cross-state-fec-money.md](#) §F4.*)
- **Party-resolved — DONE for NY + ID** ([donor-class-and-the-electorate.md](#)): with party of record the demographic claim sharpens into a partisan one, and the striking result is that it holds **in both directions of the spectrum** — the donor class over-represents registered Democrats relative to the electorate in **deep-blue NY (+16.1 pts federal, +9.6 state)** and **deep-red Idaho (+8.6 federal, +9.8 state)**, under-represents the unaffiliated in both, and the age skew replicates (65+: NY federal **49.9%**, ID federal **66.8%**, ID state **51.3%**). Crossover: Democrats are near-monolithic donors (94% NY / **94.6%** ID → own party) and unaffiliated donors lean Democratic (~2:1 NY, nearly **4:1** ID). So the Democratic tilt is a property of who donates, not of a state's majority party. (The ID crossover and 51.3% figures are the state-money layer; see that paper's caveats.) **These are the primary (full-name-key) specification.** It discards 11–19% of matched donors who are younger and less Democratic than those retained, so part of the sharpened skew is selection rather than precision — the all-tier figures, which are the more conservative ones, are reported alongside in the donor-class paper.
- **Key literature.** Demos (*Whose Voice, Whose Choice*); Bonica (DIME); and the same-state result that tempers the benign reading — **Grumbach, Sahn & Staszak (APSR): Seattle's democracy vouchers did not diversify the donor pool.**

6. Money marks strength; it does not appear to move margin

Insight 42 · failure-contribution 14 · null: weak-against · V-Dem: egalitarian, electoral, deliberative · trajectory: indeterminate

- **Defensible claim.** Finance behaves as a marker of pre-existing candidate strength, not an independent vote-mover: raw fundraising log2(D/R) correlates **+0.55** with overperformance, but spend *allocation* has cross-cycle holdout R² **~0.00**, and the forecast model zeroes the fundraising term post-redistricting because the baseline already absorbs it. The honest reading — "money follows the scoreboard" — is, on the *vote-buying* axis, consistent with the healthy null, while leaving the **access/agenda-setting** channel untested.
- **Strongest objection.** The whole thing is a correlation with no exogenous variation — +0.55 is exactly what a true causal effect *would also* produce. The allocation-R²-zero null is underpowered and tests spend *mix*, not *level*. So the data can neither confirm nor refute vote-buying; it can only show money is endogenous to candidate quality.
- **First analysis — DONE** ([scripts/diag_ie_vs_margin.py](#)). For FEC-attributed races (Schedule E carries support/oppose + district), it regresses the *fundamentals-net residual* (actual – model-predicted Dem %, **not** the raw margin) on the net pro-Dem IE advantage. The run surfaced a harder boundary than the design assumed: **directional IE on disk exists for a single cycle (2024 FEC Schedule-E, 7 WA U.S. House races)**, and the \$70.6M of PDC state-legislative IE carries a NULL support/oppose flag — so it cannot enter a directional test at all. One cross-section of 7 races cannot bear the bootstrap CI, the early-vs-late split, or the next-cycle placebo, so **inference is withheld** (the script says so rather than reporting a coin-flip slope). What the descriptive cross-section *shows* still cuts toward the null: the association is, if anything, **negative** (–0.39 pp per \$1M net pro-Dem IE, Pearson r –0.39, n=7) — money flowing

toward the side that is *behind*, the textbook endogeneity signature — and the single most IE-saturated House race in the country, **WA-03 2024 (\$40.1M total IE, +\$16.2M net pro-Dem), finished +0.06 pp off its fundamentals: dead-on.** The citable Finding-6 result for WA today is therefore the *data ceiling itself*: the public machine-readable IE record is too thin to move the verdict off "cannot confirm or refute." Unlock = multi-cycle FEC Schedule-E backfill (`load_fec_independent_expenditures(cycle=...)` , rate-limited, federal House only); even then n stays small and uninstrumented.

- **Key literature.** Jacobson (spending endogeneity); Kalla & Broockman ($\approx$ zero average general-election persuasion across 49 experiments); Bonica (money as information, not purchase).

Boundary of inference — what this data cannot (yet) support

These questions are *significant* but the warehouse cannot currently carry them. They are reported as limits — and they map the data we would need to say more.

- **Elections financed by non-constituents (insight 31) — NO LONGER BLOCKED; the test ran.** This was reported as un-runnable because only donor-residence *outflow* was loaded and NY/TX contributions were absent. Both are now fixed: NY+TX+ID FEC contributions were loaded, and a **recipient-anchored inflow** dataset (5.50M gifts / \$1.21B, all-state donors → WA/NY/TX/ID federal candidates) was built, so the nationalization test *did* run (`cross-state-fec-money.md` §E-I). The result is **not** null-supporting at the inflow level: **27% (WA) / 41% (NY) / 34% (TX) / 53% (ID)** of these states' U.S.-House+Senate candidate money is out-of-region, and on the **outflow** side each state sends the *majority* of its candidate money out of region (NY 62%, ID 68% — the most — TX 43%) — Georgia Senate races alone drew ~\$68M from residents of the four states. The earlier "93.6% in-state" was indeed the ingestion artifact this item warned about. (State-level money — WA PDC, NY BOE, ID Sunshine, TX TEC — is now also loaded, though not yet folded into these federal cross-state cuts.)
- **Straight-ticket lock-in / candidate accountability (insight 28, data-sufficiency 3).** Individual ticket-splitting is **unmeasurable** (cross-party matchback fires for $\sim$ 0 voters), and the precinct ticket-split tables can't yield a 19-election trend (epoch-versioned precincts; 2022 renumbering → 0 common precincts). The only defensible cut is an *ecological* county-level cross-office consistency trend.
- **Is there a persuadable middle? (insight 28, the one "supports-null" finding).** `voter_party_choice` is Pierce-only, 2024-only; **exactly 0 voters** appear in two party-primary cycles. WA has no party registration. The honest result is a limits-and-null finding that *converges with* the gold-standard literature (Abramowitz & Webster; Kuriwaki's cast-vote-record work), not against it.
- **Are outcomes pre-determined? (insight 38).** The notarized backtest (MAE $\sim$ 6.4, $\sim$ 94% directional) bounds how much aggregate campaigning moves results — but predictability is **not** causal inertness (Gelman & King: predictability is the *footprint* of effective deliberation), the headline is safe-seat-inflated, and the data has no counterfactual.

The pattern that matters: of the three pieces of evidence that would most strengthen a failure case, **one is now in hand** — the *true cross-state money-flow test* ran (`cross-state-fec-money.md` §E-I) and shows pervasive nationalization. The remaining two — the *partisan consequence* of the turnout skew and *individual* cross-party behavior — still trace to a single missing asset: **individual party-of-record**, which Washington does not publish. That asset is the sole remaining gate on a stronger verdict.

The verdict-in-waiting, and what gates it

On present evidence: **electoral stress, not failure.** Turnout is steeply age-skewed; the donor class is narrow, old, and top-heavy; most seats are uncontested in the general; money tracks candidate strength. Each is real and quantifiable here. But each also has a strong benign reading that survives, the trajectories are mostly borrowed from the national literature rather than established in-data, and the accuracy gate holds the whole back from a stronger verdict.

To move past "weak-against-null," the program needed to close two gaps. **Both are now closed:**

1. **~~Party-of-record (statewide)~~ — DONE.** The NY (NYSVOTER FOIL, 13.54M) and ID (SoS statewide file, 1.03M) voter files were received and loaded, each with individual party of record. The partisan consequence of turnout inequality and the individual donor crossover are now tested party-resolved in a deep-blue and a deep-red state (`who-decides-new-york.md` , `who-decides-idaho.md` , `donor-class-and-the-electorate.md`). WA itself still lacks party of record before a possible 2028 PRA.

2. ~~A genuine multi-state money panel~~ — **DONE**. NY/TX contributions loaded, the FEC ingest direction corrected with a recipient-anchored inflow build, and the nationalization test run ([cross-state-fec-money.md](#) §E-I).

With both data gaps closed, the remaining work is human-owned: independent verification of the headline numbers and posting to SSRN/SocArXiv ([publication-checklist.md](#)).

Appendix A — Diagnostic ledger (all 10 questions)

Question	Insight	Failure	Null	Impair	Traj	Entrench	Counter [↓]	Converge
Who decides (turnout skew)	62	22	weak	5	4	7	8	6
Safe-seat democracy	52	34	weak	6	5	8	8	7
Whale vs small-dollar	48	22	weak	4	4	6	8	6
Giving reinforces voting	47	22	weak	5	3	7	7	7
Donor class ≠ electorate	42	22	weak	5	4	7	7	6
Money moves margin?	42	14	weak	3	4	5	8	6
Outcomes pre-determined?	38	14	weak	3	5	6	9	5
Non-constituent money	31	11	weak	4	3	6	7	4
Straight-ticket lock-in	28	17	weak	5	7	8	8	6
Persuadable middle	28	5	supports-null	2	2	4	9	5

Counter[↓] = counter-thesis strength (reverse): high values discount the failure signal. Accuracy-weighted failure signal across all findings = 22/100.

Appendix B — Suggested publication sequence (status as of 2026-06-28)

1. **Lead paper**: "Who Decides Washington State?" (Finding 1) — **DRAFTED** ([who-decides-washington.md](#)).
2. **Companion**: "Safe-Seat Washington / the Four-State Map" (Finding 2) — **DRAFTED** ([safe-seat-washington.md](#)); observed margins, complete four-state lower-chamber map, primary-turnout cut. (Uncontested-filer join folded in as the "no major-party choice" bucket.) Party-resolved companions now drafted: "Who Decides New York?" ([who-decides-new-york.md](#) , deep blue) and "Who Decides Idaho?" ([who-decides-idaho.md](#) , deep red).
3. **Money series**: Findings 3 + 5 — **DRAFTED** as [donor-class-and-the-electorate.md](#) (WA + NY + ID, party-resolved donor class, whale concentration, crossover), building on [cross-state-fec-money.md](#) (§F donor-class + matcher-bias correction; §A/E whale-vs-small-dollar + concentration). Conduit/earmark attribution verified (§E).
4. **Methods/curiosity piece**: Finding 6 — **DRAFTED** as [does-money-move-votes.md](#) (2026-07-26). The story is the honest near-null *and* the data ceiling: money is the strongest single correlate of overperformance (+0.55) yet leaves no causal fingerprint — allocation holdout $R^2 \approx 0$, the forecast model discards the term against a known baseline, and the directional IE cross-section runs *negative* (n=7). Directional IE is one cycle deep and the \$70.6M of PDC legislative IE has no support/oppose flag, so the verdict stays "cannot confirm or refute," and the data ceiling is the citable result.
5. **Party-of-record boundary questions** — **DONE** (NY + ID voter files loaded); the longitudinal/causal version of Finding 4 remains for a future extension.