

Electoral Health in Four States

Evidence on participation, contestation and political voice, from precinct results, campaign finance and the individual voter record — Washington, New York, Texas and Idaho

Synthesis of the electoral-health series. DRAFT — pending human/editorial sign-off. `scripts/verify_whitepaper.py` scrapes this document and asserts its figures, including against the companion papers that own them; that gate is automated and is not the sign-off.

This document was a research prospectus until 2026-08-16 — an outline of analyses to run, derived from a 2026-06-27 idea-scoring exercise, carrying a scalar "accuracy-weighted failure signal" of 22/100 and a ten-question scorecard. The analyses have since been run and written up as eight companion papers, and an external referee's judgement was that the document could no longer credibly be both a prospectus and a synthesis: every improvement in a companion created a propagation problem here, and the scorecard's model-authored scores were doing work as evidence that elicited judgement cannot do. It is now a synthesis. The gauntlet, the ten questions, the eleven scoring dimensions and the 22/100 are preserved in Appendix A as the programme's methodological history — which is what they are — and no longer appear in the argument.

What this paper claims, and what it does not

This is a synthesis of eight companion papers. It does not re-derive their results: where a figure belongs to a companion it is **scraped from that companion**, so that improving the companion moves this document or fails its gate loudly. That mechanism exists because the opposite happened — until 2026-08-15 this document carried a pooled donor panel, a retired match key and a headline slope that its companions had all superseded.

The claim. The four-state evidence documents persistent inequalities in who participates, who finances campaigns, and how often general elections are genuinely close. It does **not** support a scalar diagnosis of democratic failure, and it does **not** identify the causal electoral effect of campaign spending.

Scope and integrity statement (read first)

This program studies the **electoral** dimension of democracy — participation, representation, political equality, and contestation. It is **silent** on rule of law, civil liberties, press freedom, and executive constraint, and it makes no claim about them.

Three commitments govern every section below:

1. **Pre-specified working null.** The working hypothesis is that "Washington's (and the comparison states') electoral democracy is functioning." A finding moves against that null only when the evidence earns it. *This was called a "pre-registered null" until 2026-08-16. It is not one: there was no frozen, time-stamped registration of hypotheses, estimands, analysis rules and decision criteria before the analyses ran. A general prior position is a working null, and calling it pre-registration claims a discipline this programme did not exercise.*
2. **Accuracy gate.** Each finding contributes to any larger conclusion *only in proportion to its data sufficiency × inferential strength*. A striking number built on data the warehouse cannot actually carry is reported as a limit, not a result.
3. **A bounded conclusion, not a scalar verdict.** This document assembles the *components* of a democratic-health assessment. Its conclusion is that **the analyses document several measurable disparities and constraints in electoral participation, contestation and political voice, and do not establish system-level democratic failure**. What it does not do is reduce that to a number.

This commitment read "No verdict ... deliberately does not declare a failure of democracy" until 2026-08-15, while the closing section was titled "The verdict-in-waiting" and opened "electoral stress, not failure", and the paragraph below assigned a 22/100 failure signal. An external referee named the contradiction. Claiming to reach no verdict while publishing a scalar failure score is the version that could not be defended; the bounded interpretive conclusion above is stated plainly instead.

Data, and what each state can actually support

Data provenance. Every count below is a **live read that drifts** as the 2026 cycle accrues, and each is scoped to the state or

states named — the previous version of this paragraph was not, and an external referee read its contribution figures as program-wide when they were Washington's alone.

- **Washington:** ~5.1M precinct-result rows across 22 elections; **8.6M individual contributions, \$1.04B** (FEC \$646.2M + state PDC \$394.6M, 2018–2026) with employer/occupation/ZIP; ~16.1M voter-score rows (one per voter per district scope, ~5.5M distinct voters); ~27.1M individual VRDB vote records.
- **All four states, federal layer only: 28.9M contribution rows / \$4.73B** of resident-donor outflow (WA 5.58M/\$646.2M · NY 9.98M/\$2,066.6M · TX 12.56M/\$1,942.1M · ID 0.77M/\$76.2M). Adding each state's own finance system takes it to **36.1M rows / \$6.06B**.
- **The rare asset:** voters matched to their own donations at the person level. Quote these **per panel**, never pooled — WA federal 147,745 / state 217,114; NY federal 269,218 / state 378,383; ID federal 23,303 / state 23,613. The pooled 314,974 WA figure this document used to headline is a **retired estimand**; see the panel note in Finding 5.

Legal basis for each voter file, since a synthesis is where a reader looks for it. Washington under **RCW 29A.08.720**; New York under **NYSVOTER**, obtained by FOIL and used under the elections-purpose certification of **N.Y. Elec. Law § 3-103(5)**; Idaho under **Idaho Code § 34-437A**. None is redistributed. Texas publishes no voter file, which is why it appears in the money and contestation rows below and nowhere else.

The four states are not four symmetric assessments, and the design should say so first rather than let a reader discover it. Texas participates fully in the money and contestation evidence and has no voter file, so no person-level donor analysis exists for it. New York and Idaho publish individual party of record; Washington does not. Idaho's roll is a current extract with a survivorship problem unlike Washington's.

domain	Washington	New York	Texas	Idaho
Participation — age composition of the electorate	measured (27.1M vote records, birthdate)	measured (13.5M roll, year of birth)	unavailable (no voter file)	measured , with a roll-survivorship caveat
Participation — <i>party-resolved</i>	unavailable (no party of record)	measured	unavailable	measured
Contestation — observed margins and ballot choice	measured	measured	measured	measured
Political voice — aggregate money	measured (federal + state)	measured (federal + state)	measured (federal + state)	measured (federal + state)
Political voice — <i>person-level</i> voter↔donor linkage	measured	measured	unavailable (no voter file)	measured
Campaign effects — directional IE against margin	measured (34 federal + 129 state cells)	unavailable	unavailable	unavailable

Read every cross-state statement against this table. "Four states" is true of the money and contestation evidence; the person-level evidence is three states, and the campaign-effects evidence is one.

Method

Each research question was scored 0–10 on eleven dimensions by a **language model** (`idea-gauntlet/democracy_insight.workflow.js`) prompted to act as a quantitative methodologist and democratic-theory reviewer, to search the web, and to be adversarial toward its own conclusions:

- **Insight layer (research quality):** data sufficiency, inferential strength, novelty (vs. the existing political-science literature), systemic significance, usefulness, robustness.
- **Diagnostic layer (anatomy of a failure claim):** function impairment (severity vs. the democratic ideal), trajectory (is it worsening?), entrenchment (does it resist self-correction?), counter-thesis strength (how good is the "this is actually healthy" reading — *reverse-scored*), convergence (does it triangulate with the other findings?).

Function impairment is **conceptually mapped** to the V-Dem component indices (electoral, liberal, participatory, deliberative, egalitarian) so severity is legible in a familiar vocabulary.

This paragraph said "anchored to" until 2026-08-15. An external referee objected that "anchored" claims a calibration that does not exist, and he is right: these are not V-Dem component-index observations, no reproducible transformation maps V-Dem indicators onto these 0–10 scores, and the scores were produced by a language-model workflow. The mapping is a labelling convention borrowed from V-Dem's conceptual scheme, and nothing about the scores inherits V-Dem's authority.

The four evidence domains, and which of them share data

The evidence below is grouped into four domains rather than the ten scored questions of the original gauntlet (Appendix A). The regrouping is not cosmetic. **The ten questions were not ten independent studies**, and counting them as though they were inflated an implicit evidence count — "convergence" was itself one of the scored dimensions.

domain	what it measures	companion papers	shares data with
1. Participation	who actually votes, by age and party	who-decides- {washington,new-york,idaho}, who-returns-ballot	shares the voter files with Domain 3's linkage
2. Contestation	whether general elections are close, and whether they offer a major-party choice	safe-seat-washington	shares election outcomes with Domain 4
3. Political voice	who funds campaigns, person by person	donor-class-and-the-electorate, cross-state-fec-money	3a/3b/3c all draw on the same contribution systems; 3b/3c share one voter↔donor linkage
4. Campaign effects	whether money moves margin	does-money-move-votes	shares election outcomes and the forecast residual with Domain 2

So Domains 1 and 3 are not independent of each other, and 2 and 4 are not independent of each other; within Domain 3, the three results are three transformations of one data-generating process. Triangulating several estimands from the same data is legitimate and is what this programme does. Treating the count of results as a count of corroborating studies is not.

The evidence, in four domains

Domain 1 — Participation: who actually votes

- **Defensible claim.** The electorate that decides odd-year general elections in WA is roughly half the size of the presidential electorate and dramatically older. From **27.1M** VRDB vote records (~100% birthdate coverage): voters 65 and older were **36.7%, 40.2% and 40.3%** of the 2021, 2023 and 2025 odd-year electorates against **28.5%** in 2024, while voters 18–29 fell from **14.2%** in 2024 to about **7.6%** off-cycle. Individually, 18–29 participation falls from **58.4%** (2024) to about **16%** off-year, while 65+ slips only from **88.3%** to **~61%**.

Re-quoted at the owning paper's own precision on 2026-08-16. This bullet gave the off-year senior share as a rounded band ("~37–40%"), and gave two within-cohort rates — **15.8%** and **61.3%** — to a decimal place that [who-decides-washington.md](#) does not publish, because that paper reports them as "about 16%" and "~61%". A synthesis stating a companion's figure more precisely than the companion does is drift waiting to happen, and there is no basis on which the extra digit is more correct. All of these are now scraped from the companion. - **Strongest objection.** This is *voluntary* differential participation, not disenfranchisement. WA has all-mail, postage-paid, automatic/same-day registration — the gray electorate reflects who *chose* to vote, and the gap is mechanically expected from low salience. The fix is purely institutional (move local races on-cycle to even Novembers), which is itself evidence the system is responsive, not failing. - **First analysis — DONE** ([who-decides-washington.md](#)): `Join voting_history to voters.birthdate ; classify age-as-of-each-election into cohorts ; produce a cycle-type × cohort table of both (a) within-cohort turnout and (b) cohort share of the actual electorate, for every general 2021–2025. Report composition shares (denominator-free) alongside rates to neutralize the registration-churn objection. Stop short of any partisan-consequence claim` — that needs statewide party-of-record (see §2 of the follow-on). - **Key literature.** Hajnal & Trounstein (off-cycle timing skews the electorate older/whiter); Anzia (*Timing and Turnout*); Lucero et al. 2025, surveying cities that switched, which puts voters **over 45** at **58.4%** of the off-cycle electorate

against **49.7%** of the presidential-year one (an earlier version of this line gave the cohort, the magnitude and the design wrong, and its "28%" appears to have been Washington's own presidential 65+ share); and **Ornstein (2024)** on California's SB 415 across 236 local governments, which finds the expected turnout and diversity gains but **no** detectable effect on representation, candidacy, incumbency, housing policy or public-employee salaries — the disconfirming half of the same literature.

Domain 2 — Contestation: whether general elections are close

- **Defensible claim, on OBSERVED margins.** The large majority of seats are not close. On the most recent general for partisan legislative and congressional offices, the not-close share runs **WA 87.8% · NY 88.0% · TX 94.0% · ID 92.9%** of lower-chamber seats. This is a structural counting result on certified returns that does **not** depend on any blocked/weak signal or on any model.

Two changes here on 2026-08-15, both at an external referee's insistence, and both make the finding harder to attack. *First, the live forecast table is gone.* This bullet led with model-projected 2026 margins from the latest `forecast_predictions` snapshot per state — figures that were unpinned, unprobed, drifting, due to move wholesale at the pre-November re-lock, and disagreeing with the observed count by thirteen points in Texas. They are not reproduced here, because enumerating four unpinned live reads is the thing the note objects to. The observed result is both stronger and stable. *Second, "trajectory: worsening" is withdrawn.* Washington's observed not-close share runs **78.9–88.1%** across 2016–2024 with a dip in the 2018 wave; that is persistence, not a monotonic trend, and five cycles cannot carry one. *And the sentence "when the general is foregone, the operative decision moves to lower-turnout primaries" is removed as an inference the safe-seat companion has expressly withdrawn* — an ex-post margin cannot date the binding decision. *(A basis note here described how the four retired band cells were computed from `forecast_predictions` — including that the Texas district count differed between two snapshots four days apart, which is why they were never trustworthy as published figures. It went with the cells.)* - **Strongest objection.** Those are *model-projected* 2026 margins, sensitive to the 10-pt threshold; and "non-competitive" conflates a closed-shop seat with one that is lopsided because voters *genuinely* lean that way (self-sorting). A Solid-D Seattle seat and a Solid-R rural-ID seat each *represent* their electorates — the null's "outcomes reflect voter choice." - **First analysis — DONE** (`safe-seat-washington.md`): Ran **observed** margins, not the projection: per state, the most recent general for partisan legislative + congressional offices → margin per district → band counts. **The companion has since split this into two dimensions that must not be merged:** *candidate competition* (was the race close, on the top-two margin regardless of party) and *partisan availability* (did the ballot offer both a Democrat and a Republican). They overlap but are different questions — in WA 2024, 83.5% of seats were not close and 34.6% offered no D-v-R option, and of the sixteen same-party generals fifteen were also lopsided but **one was decided by six points**. Treating "no major-party choice" as automatically non-competitive, as this bullet originally did, is the conflation the rewrite removed. For WA, built the 2016–2024 seat-level trend (not-close share runs **78.9–88.1%** across the five cycles, dipping in the 2018 wave) and the primary-to-general turnout ratio in safe seats (~0.5) — the one place this data *adds* to Cook/Ballotpedia/Unite America by attaching real turnout (and the voter-donor join) to the "primary decides" claim. - **Key literature.** Cook PVI (swing districts 164 in 1997 → ~72 post-2016); Ballotpedia Competitiveness Index (38% of state-leg seats uncontested in 2024); Unite America (*The Primary Problem*).

Domain 3 — Political voice: who funds campaigns

The three results below share their contribution systems, and the second and third share a single voter↔donor linkage. They are one body of evidence seen three ways.

3a. Small transactions coexist with high dollar concentration

Titled "Whale-dominated money behind a small-dollar facade" until 2026-08-16. A low median gift and a high Gini coexist naturally in any heavy-tailed distribution; neither is a facade concealing the other, and "facade" asserted a concealment the measurement does not show. This is also, and only, an itemized-contribution finding: a small transaction in an FEC itemized file is not an observation of the universe of small donors, whom the disclosure floors keep out of the record entirely.

- **Defensible claim.** WA money is broad by headcount but concentrated by dollar. The **median itemized gift is \$25** in both money systems, while per-recipient-cycle concentration is substantial: median Gini **0.578** federal and **0.579** state, over recipient-cycles with ≥100 distinct donors (**834** federal, **1,997** state). The data can quantify, per race and cycle, the sub-\$200 retail share against the whale layer.

Corrected 2026-08-10. These figures were previously given as "n=2,821 ... single gifts reach \$2.5M ... Gini ~0.61", which was computed on the **pooled** FEC+PDC table — the exact pooling Finding 5's panel note says the series corrected everywhere else. Pooled, the count is 2,831 today; separated it is 834 federal and 1,997 state. The **\$2.5M maximum is a PDC state gift**; the federal maximum is \$929,600, so the old sentence paired a median from one money

system with a maximum from another. The 0.61 does not reproduce on any basis I can construct — both layers give 0.578 — and it is withdrawn rather than restated. **Basis, and a live-read warning (2026-08-10).** Population: recipient-cycles keyed (`fec_candidate_id`, `election_cycle`) over `individual_contributions`, positive amounts, split by money system on the PDC: id prefix, restricted to ≥ 100 distinct (name, zip5) donors. *Pooled is a separate grouping, not the sum of the two* — a recipient-cycle below the threshold in each system alone can clear it pooled — which is why $834 + 1,997 \neq 2,831$. **These three counts are unpinning live reads and they drift upward** as the 2026 PDC cycle accrues: they read 822 / 1,989 / 2,814 when this correction was first written, hours earlier the same day. All eight figures in this finding are now asserted by `scripts/verify_whitepaper.py`, which had asserted none of them — including the 0.578 that the series' withdrawn-claim register already recorded as guarded by it. - **Strongest objection.** The recipient key is `fec_candidate_id`, which holds **two identifier systems**: FEC committee ids and PDC: -prefixed state filer ids. On the federal side the highest-volume "recipients" are conduits (ActBlue, JFCs), so a naive per-recipient Gini measures a *conduit's pass-through book*, not a candidate's race; on the state side — which is 71% of the pooled recipient-cycles — that objection does not apply, and a different one does. A true per-race figure needs an earmark-attribution layer that does not yet exist and must be validated against double-counting. And a Gini that mixes \$25 and \$2.5M is mechanically high regardless of democratic health. - **First analysis.** Build the conduit-attribution layer *first* (partition direct-to-candidate vs. conduit-routed; verify whether ActBlue/WinRed rows carry attributable earmark memos). Then compute, per ultimate-recipient $\times$ cycle: Gini, top-1%/top-10% dollar share, and sub-\$200 vs. $\geq \$200$ split — reported **with and without** conduit pass-throughs, since that toggle determines the answer. - **Key literature.** Bonica & Rosenthal (wealth-elasticity of giving); Bouton & Cagé (*Small Campaign Donors*); Brennan Center (small-dollar grew, megadonors grew faster — "thin layer over a concentrated core").

3b. Donors are also high-frequency voters

- **Defensible claim (cross-sectional only).** Donors are a participation elite. On the measure the owning paper now treats as clean — registrants who existed before the first election in the window, so every retained person could have voted in all of them, then age-standardized — the donor/non-donor turnout gap runs **+22.9 to +26.3 points** across the four panels (NY federal +25.1, NY state +26.3, WA federal +26.0, WA state +22.9). Financial voice and electoral voice concentrate on the same individuals rather than offsetting.

Retired 2026-08-15: "87.6% are super-voters vs 50.9%" of non-donors, at an average turnout propensity 0.967 vs 0.749 and a ratio of 1.72 $\times$. Those figures came from the **pooled** 314,974-voter match, an estimand the donor paper has since replaced with per-source panels. Worse, the measure itself is defective for this comparison: `voter_scores.is_super_voter` is defined as *last voted on or after 1 Jan 2022 and registered at least eight years*, so registration tenure sits inside the outcome definition — every registrant of under eight years' standing is false by construction, the variable cannot carry a tenure adjustment, and it is not measure-comparable with New York's plain count of generals voted. The owning paper says all of this and publishes the eligible-for-all figures above instead. An external referee caught the synthesis still headlining the dramatic superseded metric, which is precisely what a synthesis must not do. - **Strongest objection.** The *causal/longitudinal* version ("giving *makes* people vote more; inequality *deepens*") is unsupportable here: donors are pre-selected for engagement (reverse causation is equally plausible), the match is biased toward older/stable-address/super-voting/uncommon-named people (inflating the gap), and `voter_donor_affiliation` collapses each voter to one row, so first-gift-then-vote sequencing can't be reconstructed on a shallow 2021–2026 history. The benign reading — donating as a *gateway* that broadens participation — is live. - **First analysis — DONE**, and superseded by `donor-class-and-the-electorate.md`, which is the citable owner of this result. Framed strictly as association: the giving $\leftrightarrow$ voting overlap is real, the *causal/longitudinal* version remains out of reach, and the eligible-for-all age-standardized gaps above are the figures to quote. - **Key literature.** Verba/Schlozman/Brady (*Voice and Equality; Unheavenly Chorus*) — money is the most income-skewed form of participation; the donor pool is a structural elite (the *constant* against which any worsening must be measured).

3c. The donor class is not the electorate

□ **Panel note (2026-07-26).** The WA figures in this finding were computed on a **pooled** voter $\leftrightarrow$ donor match, before it was discovered that WA's `individual_contributions` holds federal (FEC: , \$646.2M) and state (PDC: , \$394.6M) money and the matcher had no source filter. Pooling stacks one person's federal and state giving into a single donor total and inflates concentration. `donor-class-and-the-electorate.md` now reports two panels — **federal** 147,745 donors / top-1% **41.2%** / Gini **0.815**, **state** 217,114 / **43.5%** / **0.821** — which supersede the pooled numbers below. Every finding survives, and the age skew *strengthens* on the federal panel (Silent **2.67 $\times$** , Gen Z **0.04 $\times$**).

- **Defensible claim.** At the person level, donors are a narrow slice that does not mirror the electorate: **~4.0–5.7% of voters** — matched donors as a share of all registrants on the pooled panels, ID **4.0%**, NY **4.1%**, WA **5.7%** — skewed old (pooled WA: Silent **1.96 $\times$** , Boomer **1.71 $\times$** over-represented; Gen Z **0.09 $\times$** , Millennial **0.54 $\times$** under-represented — 2.67 $\times$ / 2.04 $\times$ / 0.04 $\times$ / 0.35 $\times$ on the federal panel), overwhelmingly super-voters, geographically concentrated (**61.4%**

of WA donor dollars from two Seattle-metro ZIP3s; 63.5% federal-only), and internally top-heavy (top 1% of matched donors supply **46.6%** of matched dollars pooled, **41.2%** federal; top 10% **79.3% / 74.2%**). RETIRED (\$154.0M, 23.8%) and NOT EMPLOYED (\$128.8M, 19.9%) are the two largest occupation blocs. *(These are contribution-level, not matched-donor, figures — occupation is on the gift, not the voter — and they are stated on the series' documented outflow basis: FEC rows, Washington-resident donors, \$646.2M. Until 2026-07-27 they read \$221.7M / 21.3% and \$147.4M / 14.1%, computed on the unfiltered pooled table — FEC plus state PDC plus non-resident donors, \$1,050.8M — i.e. the same two-money-system pooling this series corrected everywhere else, sitting inside a finding whose every other number is panel-scoped. On the federal basis both blocs are a LARGER share, so the point sharpens.)* The voter→donor join makes this *person-level*, not ecological — the genuine value-add over the standard FEC-aggregate literature.

- **Strongest objection — one mechanism tested and excluded, the objection not closed.** The objection is that the skew is a **matcher artifact**: the unique-key over-selects the older, rarer-named, stable-address residents it then "finds" over-represented. What has been tested is the **roll-side** half of that. On the current full-name key the probability a voter is uniquely matchable is nearly flat across age — **NY 94.5–95.4%** (0.9-pt spread), **WA 96.3–97.6%** (1.4-pt spread) — so inverse-propensity re-weighting does not move the distribution: NY's 65+ donor share goes 47.9% → **47.9%**, and every Washington generation multiplier is unchanged to two decimal places in both panels.

The defensible conclusion is therefore **"measured roll-side strict-key matchability does not explain the age skew"** — not "matcher bias has been rejected". The test says nothing about donor-side selection: how a contributor writes their own name on a filing, residential mobility between the filing and the extract, work or second addresses, stale addresses, or namesakes. The owning paper enumerates those and leaves them open. The voter file still has **no income and no race**, so the literature's richer/whiter claims remain proxied, not tested.

Corrected 2026-08-15. This bullet said the objection was "tested and rejected" and that the skew "is a real property of *who gives*, not of who the matcher can find" — a categorical claim the owning paper explicitly declines to make. It also quoted **69.1%–73.3%** matchability, which is the *retired* (last name, first initial, ZIP) key; the current full-name key runs 94.5–97.6%. An external referee caught both. - **First analysis — DONE** ([cross-state-fec-money.md](#) §F2): the matcher-bias inverse-propensity re-weighting above; skews reported raw *and* re-weighted; income/race labeled untestable. (*Retired estimand, kept only as the thing the per-panel figures replaced: on the pooled 314,974 match, top-1% 46.6%, top-10% 79.3%, Gini 0.857, and 61.4% of dollars from two Seattle ZIP3s.*) Per panel — federal panel top-1% **41.2%** [38.6–43.4], Gini **0.815** [0.806–0.822], two-ZIP3 share **63.5%**; state panel top-1% **43.5%** [38.7–48.9], Gini **0.821** [0.806–0.838]. (*This line previously read 42.4% [40.2–44.9] for the federal panel — the all-tier value, contradicting the 41.2% in the panel note directly above it. The bootstrap CIs are the per-panel re-runs in [cross-state-fec-money.md](#) §F4; the panels are full populations, so read the intervals as resampling sensitivity to donor composition, not sampling error.*) - **Party-resolved — DONE for NY + ID** ([donor-class-and-the-electorate.md](#)): with party of record the demographic claim sharpens into a partisan one, and the striking result is that it holds **in both directions of the spectrum** — the donor class over-represents registered Democrats relative to the electorate in **deep-blue NY (+16.1 pts federal, +9.6 state) and deep-red Idaho (+8.6 federal, +9.8 state)**, under-represents the unaffiliated in both, and the age skew replicates (65+: NY federal **49.9%**, ID federal **66.8%**, ID state **51.3%**). **Crossover, on the federal panels only:** currently enrolled Democrats are near-monolithic donors (**95%** NY federal), and unaffiliated donors' resolved giving leans Democratic — and both survive assigning the *entire* unresolved pool to the rival side (NY unaffiliated **52.1%** against 45.1%; ID unaffiliated **65.0%** against 34.1%, as shares of matched donors). So the Democratic tilt is a property of who donates, not of a state's majority party.

Panel corrected 2026-08-15. This sentence quoted the **state** panels — ID 94.6% and "nearly 4:1" — where recipient party is reconstructed rather than published and resolution is only 51.1% for Idaho and 37.7% for New York (39.0% for Idaho's unaffiliated row). The owning paper's bound analysis says in terms that the two patterns **do not survive** the extreme unresolved-pool assignment on the state panels, and **do** survive it on the federal panels, where resolution is 87.6–88.8%. The synthesis was resting a headline on the panel that fails the test while the panel that passes it sat unused. State-panel crossover is suggestive and is reported as such in the companion. **These are the primary (full-name-key) specification.** It discards 11–19% of matched donors who are younger and less Democratic than those retained, so part of the sharpened skew is selection rather than precision — the all-tier figures, which are the more conservative ones, are reported alongside in the donor-class paper. - **Key literature.** Demos (*Whose Voice, Whose Choice*); Bonica (DIME); and the same-state result that tempers the benign reading — Yorgason (*APSR 2024/25*), **"Campaign Finance Vouchers Do Not Expand the Diversity of Donors: Evidence from Seattle": the voucher recipients were the same wealthier, whiter, older, more civically engaged people who dominated before the reform.** (An earlier version attributed this to Grumbach, Sahn & Staszak, who wrote a different paper on a different question.)

Domain 4 — Campaign effects: money and margin

- **Defensible claim.** Fundraising is the single strongest correlate of overperformance in the panel — raw log₂(D/R)

receipts correlate **+0.60** — while spend *allocation* carries a cross-cycle holdout R^2 of **0.028**, and the forecast model zeroes the fundraising term post-redistricting because the baseline already absorbs it. **What none of that establishes is direction.** The association is equally consistent with money buying votes and with money chasing candidates who were already going to win, and no design available here separates them. The citable result is a **non-identification**: the public record can measure the association precisely and cannot sign the causal effect.

This finding was titled "Money marks strength; it does not appear to move margin" until 2026-08-15, and an external referee was right that the title contradicted the bullet beneath it — which already conceded that the correlation "is exactly what a true causal effect *would also* produce". A heading cannot assert non-effect while the text concedes non-identification. Two figures were stale at the same time: the correlation read **+0.58** and the holdout R^2 **0.02**, against the owning paper's current **+0.60** and **0.028** — and the retired R^2 also rounded to a different first decimal than the current one, so the old cell carried a stale rounding on top of a stale figure. Both are scraped from [does-money-move-votes.md](#) rather than restated; the scrapes had silently broken when that paper's tables were rebuilt, which is how the figures drifted unnoticed. - **Strongest objection — and it is sustained.** The whole thing is a correlation with no exogenous variation. The allocation- R^2 null is underpowered and tests spend *mix*, not *level*. So the data can neither confirm nor refute vote-buying; it can only show money is endogenous to candidate quality. - **First analysis — DONE** (`scripts/diag_ie_vs_margin.py`). For FEC-attributed races (Schedule E carries support/oppose + district), it regresses the *fundamentals-net residual* (actual – model-predicted Dem %, **not** the raw margin) on the net pro-Dem IE advantage. **Directional IE on disk spans five cycles (2018–2026 FEC Schedule-E, 34 scorable WA U.S. House races)**, and the **\$51.7M of direction-coded PDC state-legislative IE** (form C-6 section C6.3, ingested 2026-08-09) adds **129 scorable district-cycles** on the same design — see `scripts/diag_pdc_ie_vs_margin.py` . That extension does not settle the question: every interval still spans zero and the **sign becomes specification-dependent** (–3.836 to +4.871 across four specifications), which is evidence the constraint is structural rather than a shortage of cells. The regression **now runs**: the slope is **+0.515 pp per \$1M net pro-Dem IE (Pearson r +0.186, n=34)**, with a bootstrap interval of –0.600 to +2.821 that spans zero.

That positive sign rests on one observation, and the owning paper says so. Its leave-one-out sweep over the same 34 races runs from **–0.035** (dropping WA-08 2018) to **+0.832** (dropping WA-03 2024): deleting a single race erases the positive slope entirely. WA-08 2018 pairs the panel's largest net pro-Democratic IE (**+\$8.26M**) with a **+8.78**-point residual and carries a Cook's distance of **0.69**; dropping every WA-08 observation gives **–0.065**. A district-clustered bootstrap, which respects the fact that the panel measures ten districts repeatedly, widens the interval to **–1.595 to +1.268**.

What this bullet claimed until 2026-08-15, and why it is withdrawn. It said "the public record can *bound* the persuasion effect but not sign it". A bootstrap interval around an observational coefficient under endogenous treatment bounds the *association*, not the persuasion effect; calling it the latter smuggles back the identification the same sentence denies. It also headlined +0.515 without the leverage result — which the owning paper had already derived, published and led its own abstract with. That is the propagation failure this document was carrying in miniature: the companion did the careful work and the synthesis kept the dramatic version. The most heavily funded race in the panel, **WA-03 2024 (\$18.61M total IE, +\$6.09M net pro-Dem), finished +0.06 pp off its fundamentals: dead-on. Two corrections against the earlier version of this bullet, both material.** It reported a single cycle and 7 races; three further cycles were simply not loaded, and the backfill took about five minutes. And it reported a *negative* slope (–0.39, r –0.39, n=7) read as the endogeneity signature; the sign reversed on the fuller panel, which is what a coin-flip estimate does. It also called WA-03 the most IE-saturated House race in the country at \$40.1M — an artifact of a notice/periodic double-count that inflated every IE total roughly twofold. At the true \$18.61M it ranks 22nd of 387. - **Key literature.** Jacobson (spending endogeneity); Kalla & Broockman ($\approx$ zero average general-election persuasion across 49 experiments); Bonica (money as information, not purchase).

What the evidence does not support

These questions are *significant* but the warehouse cannot currently carry them. They are reported as limits — and they map the data we would need to say more.

- **Elections financed by non-constituents (insight 31) — NO LONGER BLOCKED; the test ran.** This was reported as un-runnable because only donor-residence *outflow* was loaded and NY/TX contributions were absent. Both are now fixed: NY+TX+ID FEC contributions were loaded, and a **recipient-anchored inflow** dataset (5.50M gifts / \$1.21B, all-state donors → WA/NY/TX/ID federal candidates) was built, so the nationalization test *did* run ([cross-state-fec-money.md](#) §E-I). The result is **not** null-supporting at the inflow level: **27% (WA) / 41% (NY) / 34% (TX) / 53% (ID)** of these states' U.S.-House+Senate candidate money is out-of-region, and on the **outflow** side three of the four send half or more of their candidate money out of region (ID 68% — the most — NY 62%, WA 50.8%), while **Texas sends 43%** — Georgia Senate races alone drew ~\$68M from residents of the four states. The earlier "93.6% in-state" was indeed

the ingestion artifact this item warned about. (State-level money — WA PDC, NY BOE, ID Sunshine, TX TEC — is now also loaded, though not yet folded into these federal cross-state cuts.)

- **Straight-ticket lock-in / candidate accountability (insight 28, data-sufficiency 3).** Individual ticket-splitting is **unmeasurable** (cross-party matchback fires for ~0 voters), and the precinct ticket-split tables can't yield a 22-election trend (epoch-versioned precincts; 2022 renumbering → 0 common precincts). The only defensible cut is an *ecological* county-level cross-office consistency trend.
- **Is there a persuadable middle? (insight 28, the one "supports-null" finding).** `voter_party_choice` is Pierce-only and holds **one** party-primary cycle (2024), so no voter can appear in two — that is a property of the load, not a measured null, and an earlier version reported it as "exactly 0 voters" as though it were evidence. WA has no party registration. The honest result is a limits-and-null finding that *converges with* the gold-standard literature (Abramowitz & Webster; Kuriwaki's cast-vote-record work), not against it.
- **Are outcomes pre-determined? (insight 38).** The backtest (MAE 6.41, 94% directional on 163 filtered cells; **10.12 / 85% on the unfiltered 170-cell grid**) bounds how much aggregate campaigning moves results. It is a **retrospective fit**, not a notarized out-of-sample result — what is notarized is the locked 2026 forward prediction set, which scores in November — and the model's coefficients were tuned against a residual matrix, so the figure is not pre-registered. Further — predictability is **not** causal inertness (Gelman & King: predictability is the *footprint* of effective deliberation), the headline is safe-seat-inflated, and the data has no counterfactual.

The pattern that matters: of the three pieces of evidence that would most strengthen a failure case, **one is now in hand** — the *true cross-state money-flow test* ran (`cross-state-fec-money.md` §E-I) and shows pervasive nationalization. The remaining two — the *partisan consequence* of the turnout skew and *individual* cross-party behavior — still trace to a single missing asset for Washington: **individual party-of-record**. Washington publishes it only for presidential-primary declarants and only under **RCW 29A.56.050**, which directs the Secretary of State to prescribe rules for providing the declarations, or a list of the voters who participated, to the state and county committees of that party. The statute sets no public-disclosure duration; the "~60–90 day window" this project has worked from is **unverified** (`wa-pp-party-of-record-pra-scope.md` flags it as needing confirmation, twice) and is not asserted here. What is certain is that only Pierce County was obtained in time for 2024, which is why the holding is Pierce-County-only; it is a collection-timing constraint, not a disclosure-regime one. New York and Idaho do publish party of record, and both voter files are now loaded, so the two questions are answerable there and are answered in the companions — what remains gated is the *Washington* version.

Conclusion

The four-state evidence documents persistent inequalities in who participates, who finances campaigns, and how often general elections are genuinely close. It does not support a scalar diagnosis of democratic failure, and it does not identify the causal electoral effect of campaign spending.

Concretely, and in the order of the domains above: turnout is steeply age-skewed, and sharply more so as salience falls; the donor class is narrow, old and top-heavy, and over-represents the same party at both partisan poles; general elections are mostly not close, while a smaller share offer no major-party choice at all; and money is strongly associated with candidate performance by a design that cannot say which way the association runs. Each is real and quantifiable here.

Each also has a benign reading that survives the evidence assembled. Differential participation is voluntary in states with all-mail, postage-paid, same-day registration. A lopsided seat may be lopsided because voters genuinely lean that way. A heavy-tailed contribution distribution is what any voluntary funding system produces. And a correlation between money and performance is exactly what a world in which money follows expected winners would look like. **The evidence here does not adjudicate between the concerning and the benign reading of any of the four**, which is why the conclusion is bounded rather than scored.

Two things left this section on 2026-08-15 and 2026-08-16. It opened "On present evidence: electoral stress, not failure" and closed by saying "the accuracy gate holds the whole back from a stronger verdict" — both of which read the gauntlet's scores as evidence about electoral health, which is a use elicited judgement cannot support; those scores are now Appendix A. And it said "most seats are uncontested in the general", which is false under the series' own definitions: in Washington 2024, 83.5% of seats were not close and 34.6% offered no D-v-R option. "Not close" is not "uncontested", and the safe-seat paper split the two precisely because merging them was the original conflation. The synthesis had reintroduced it in its own closing verdict — outside every coverage span, while the finding that corrects it was gated.

On the data itself, two acquisition gaps that once bounded this programme are closed:

1. **~~Party-of-record (statewide)~~ — DONE.** The NY (NYSVOTER FOIL, 13.54M) and ID (SoS statewide file, 1.03M)

voter files were received and loaded, each with individual party of record. The partisan consequence of turnout inequality and the individual donor crossover are now tested party-resolved in a deep-blue and a deep-red state ([who-decides-new-york.md](#) , [who-decides-idaho.md](#) , [donor-class-and-the-electorate.md](#)). WA itself still lacks party of record before a possible 2028 PRA.

2. ~A genuine multi-state money panel~ — **DONE**. NY/TX contributions loaded, the FEC ingest direction corrected with a recipient-anchored inflow build, and the nationalization test run ([cross-state-fec-money.md](#) §E-I).

The major acquisition gaps for the present analyses are largely closed. What remains is not only human sign-off, and this paragraph said it was until 2026-08-15. Still open and not delegable to a signature: Washington still lacks individual party of record (see *Boundary of inference*); the four states do not have equivalent coverage, so the cross-state claims are not symmetric (Texas has no voter file and therefore no person-level donor analysis; NY and ID publish party registration and WA does not; Idaho's roll carries a survivorship problem unlike Washington's); and the model-generated diagnostic scores in Appendix A are stale against the data as it now stands. The genuinely human-owned items are independent verification of the headline numbers and posting the papers not yet posted — the Washington companion is already on SSRN (7149263) and owes a revision, not a first upload ([electoral-health-audit-log.md](#)).

Appendix A — Methodological history: the 2026-06-27 idea gauntlet

None of what follows is evidence about electoral health, and it is preserved because it is the honest record of how this programme chose what to study.

The programme began with a scoring exercise. Ten candidate research questions were each scored 0–10 on eleven dimensions by a **language model** ([idea-gauntlet/democracy_insight.workflow.js](#)) prompted to act as a quantitative methodologist and democratic-theory reviewer, to search the web, and to be adversarial toward its own conclusions:

- **Insight layer (research quality):** data sufficiency, inferential strength, novelty, systemic significance, usefulness, robustness.
- **Diagnostic layer (anatomy of a failure claim):** function impairment, trajectory, entrenchment, counter-thesis strength (*reverse-scored*), convergence.

Of the ten, nine scored "weak-against-null" and one "supports-null"; none reached "moderate" or "strong" against the null. The exercise reported an **accuracy-weighted failure signal of 22/100**.

Why that number is not in the paper. It is an **elicited-judgement index, not a measurement**: the mean of ten model-authored 0–100 severity scores, weighted by twenty more model-authored 0–10 scores, of which only the first ten were ever published — so it **cannot be recomputed from the table below**, which gives an unweighted mean of 18.3. The distribution it summarises is partly prescribed by the prompt, which instructed the scorer to "*default to 'data-insufficient' or 'weak-against-null' unless the evidence is genuinely strong*", so reading "nine weak-against-null" as a finding is partly circular. And it was computed on 2026-06-27 against a data description that has since moved — 19 elections, ~\$1.3B in contributions and 9.2M voter-score rows, against 22, \$1.04B (Washington) and 16.1M today — and never re-run, although Domain 4's sign has since reversed and Domain 2's Washington figure has moved.

An external referee's verdict was that this is "a pseudo-quantitative statistic that adds false precision to what is properly a qualitative synthesis", and that no additional disclaimer cures it. That is right, and the same applies to the impairment, trajectory, entrenchment, counter-thesis and convergence columns below: they are a useful internal research-prioritisation rubric and they are not evidence about a state's electoral health.

The ledger, as scored on 2026-06-27:

Question	Insight	Failure	Null	Impair	Traj	Entrench	Counter↓	Converge
Who decides (turnout skew)	62	22	weak	5	4	7	8	6
Safe-seat democracy	52	34	weak	6	5	8	8	7
Whale vs small-dollar	48	22	weak	4	4	6	8	6
Donors are also high-frequency voters	47	22	weak	5	3	7	7	7
Donor class ≠ electorate	42	22	weak	5	4	7	7	6

Question	Insight	Failure	Null	Impair	Traj	Entrench	Counter↓	Converge
Money moves margin?	42	14	weak	3	4	5	8	6
Outcomes pre-determined?	38	14	weak	3	5	6	9	5
Non-constituent money	31	11	weak	4	3	6	7	4
Straight-ticket lock-in	28	17	weak	5	7	8	8	6
Persuadable middle	28	5	supports-null	2	2	4	9	5

Counter↓ = counter-thesis strength (reverse): high values discount the failure signal. Accuracy-weighted failure signal across all findings = 22/100.

Appendix B — The companion papers, and their status

1. **Lead paper:** "Who Decides Washington State?" (Finding 1) — **DRAFTED** ([who-decides-washington.md](#)).
2. **Companion:** "Safe-Seat Washington / the Four-State Map" (Finding 2) — **DRAFTED** ([safe-seat-washington.md](#)); observed margins, complete four-state lower-chamber map, primary-turnout cut. (Uncontested-filer join folded in as the "no major-party choice" bucket.) Party-resolved companions now drafted: "Who Decides New York?" ([who-decides-new-york.md](#) , deep blue) and "Who Decides Idaho?" ([who-decides-idaho.md](#) , deep red).
3. **Money series:** Findings 3 + 5 — **DRAFTED** as [donor-class-and-the-electorate.md](#) (WA + NY + ID, party-resolved donor class, whale concentration, crossover), building on [cross-state-fec-money.md](#) (\$F donor-class + matcher-bias correction; \$A/E whale-vs-small-dollar + concentration). *(This line read "Conduit/earmark attribution verified (\$E)" until 2026-08-15, which contradicted Finding 3's own objection bullet — that a true per-race figure "needs an earmark-attribution layer that does not yet exist". What \$E verified is that FEC records earmarked gifts under the candidate committee as transaction type 15E, so the conduit-side 24T rows are correctly excluded and nothing is double-counted. That is an ingest correctness check, not the ultimate-recipient attribution Finding 3 says it lacks. An external referee found the two sentences arguing with each other.)*
4. **Methods/curiosity piece:** Finding 6 — **DRAFTED** as [does-money-move-votes.md](#) (2026-07-26). The story is the honest non-identification: money is the strongest single correlate of overperformance (+0.60) yet leaves no causal fingerprint that this design can read — allocation holdout R^2 0.028, the forecast model discards the term against a known baseline, and the directional test, run across 34 federal district-cycles and a further 129 state-legislative ones, returns intervals that span zero on every specification and a headline slope that one deleted race reverses. The verdict stays "cannot confirm or refute," and the citable result is that the limit is the **design** — outside money concentrates in a few races and is targeted at expected closeness — not the disclosure record.
5. **Party-of-record boundary questions** — **DONE** (NY + ID voter files loaded); the longitudinal/causal version of Finding 4 remains for a future extension.